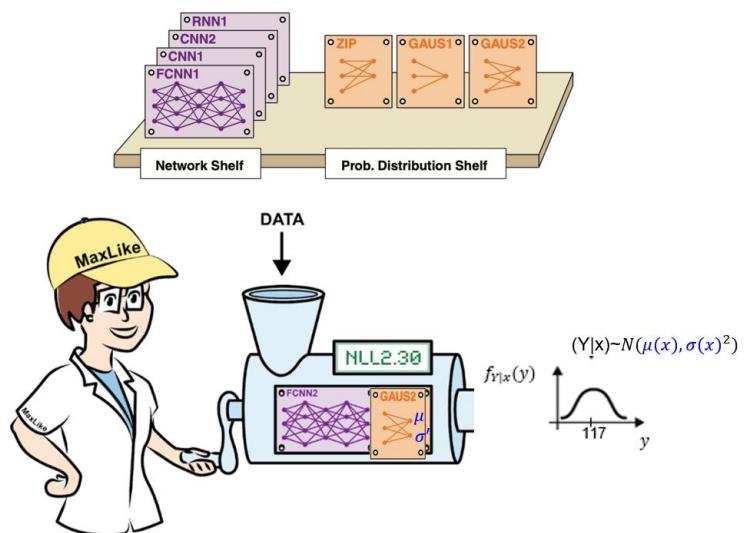


WBL Probabilistic Deep Learning:: Half day 1

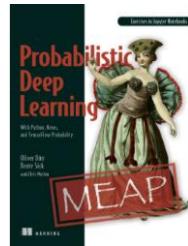
Beate Sick, Oliver Dürr

Lecture 1: Introduction to probabilistic deep learning



Literature

- Course website
 - https://tensorchiefs.github.io/dlwl_eth25/
- Probabilistic Deep Learning: Our probabilistic take
 - <https://www.manning.com/books/probabilistic-deep-learning>
- Deep Learning with Python, Second Edition
 - <https://www.manning.com/books/deep-learning-with-python-second-edition>
- Keras Documentation:
 - <https://keras.io/>
- Other Courses
 - Convolutional Neural Networks for Visual Recognition <http://cs231n.stanford.edu>

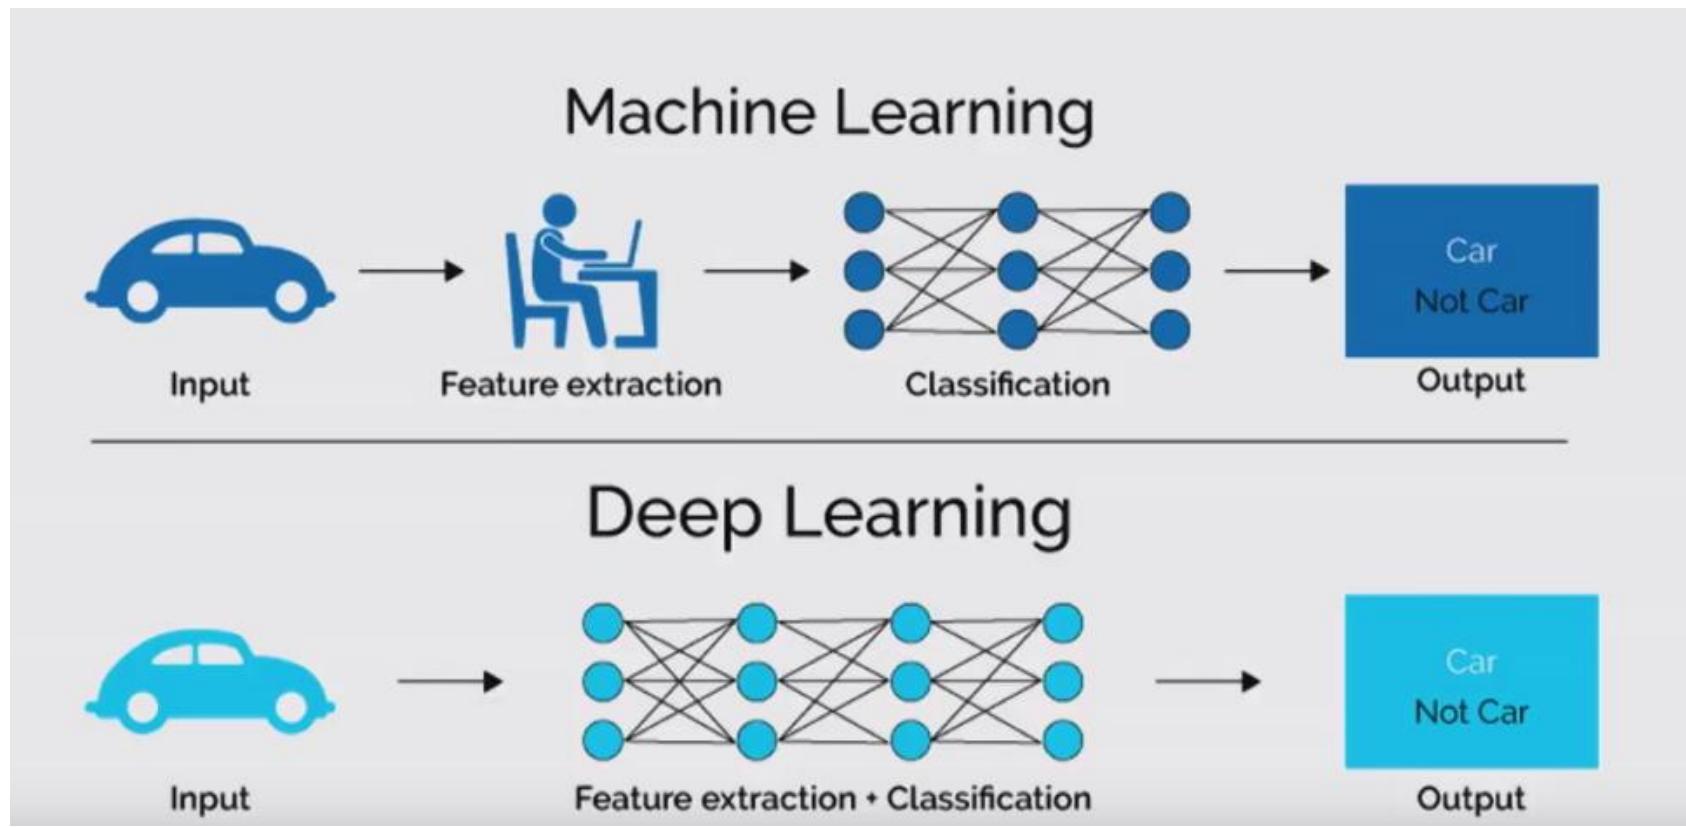


Topics (may change)

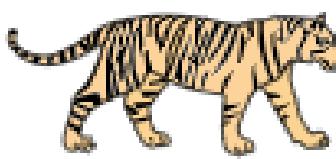
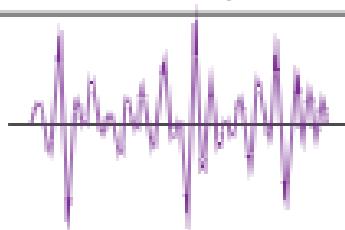
- Half day 1
 - Introduction to probabilistic DL
 - Introduction to Keras with pytorch backend
- Day 2
 - Convolutional Neural Networks (CNN) for image data
 - Transformer Architectures
 - Foundation Models = Pretrained models on many data
 - Transfer learning
 - Projects
- Day 3
 - Model evaluation
 - NN-based logistic regression with tabular and image data
 - Poisson regression with NNs
 - Projects
- Half day 4
 - Project presentations
 - tba

Date
03.02.2025
10.02.2025
17.02.2025
24.02.2025

Deep Learning vs. Machine Learning



Use cases of deep learning

Input x to DL model	Output y of DL model	Application
Images 	Label "Tiger"	Image classification
Audio 	Sequence / Text "see you tomorrow"	Voice recognition
Sequence of tokens representing e.g. "Hallo, wie gehts?"	Next token :	Translation text generation, QA
Sequence of tokens representing e.g. This movie was rather good	Label (Sentiment) positive	Sentiment analysis

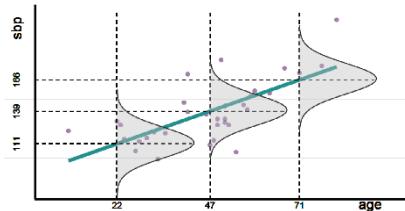
Deep learning models can „see“, „hear“, „read“, „write“ .

Last breakthrough 2022: Large Language Models (LLMs) like ChatGPT

Course topic is part of our research focus

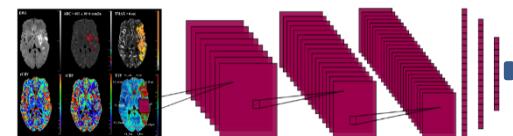
Statistics

- + transparent & interpretable
- + valid uncertainty measures
- needs structured data
- not robust vs distribution changes



Deep Learning

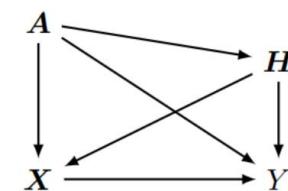
- + can handle structured & unstructured data
- + high prediction performance
- black box,
- not robust vs distribution changes



Interpretable, robust
(causally inspired)
& probabilistic DL

Causality

- + robust vs distribution changes
- + allows intervention planning
- + transparent & interpretable
- needs stronger assumptions



Focus in this DL course:

Probabilistic Viewpoint

Tasks in supervised DL

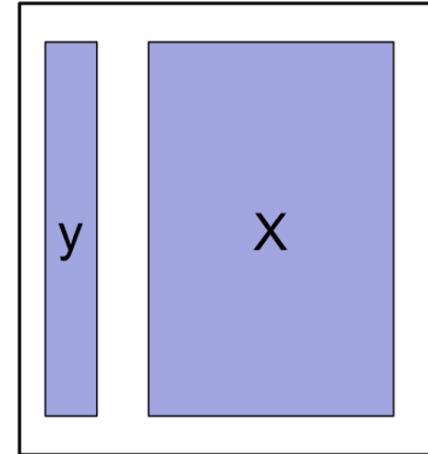
- Typical supervised DL task: predict y given x

- **Classification**

- Point prediction: Predict a class label
 - Probabilistic prediction: predict a discrete probability distribution over the class labels
 - First, we focus on probabilistic binary classification where $Y \in \{0,1\}$

- **Regression**

- Point prediction: Predict a number
 - Probabilistic prediction: predict a continuous distributions



Supervised Learning

Probabilistic vs deterministic models

Deterministic

“Classification”



“Regression”

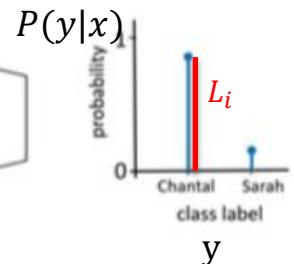


Probabilistic

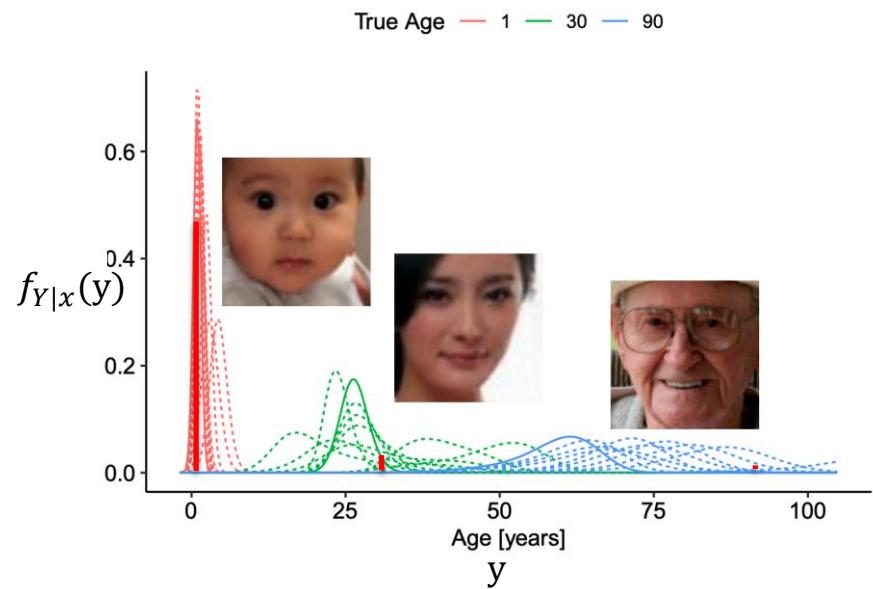


$$(x_i, y_i) = (\text{Bild}, \text{Chantal})$$

Likelihood L_i of the observed outcome y_i under the predicted conditional outcome distribution $P(Y|x_i)$

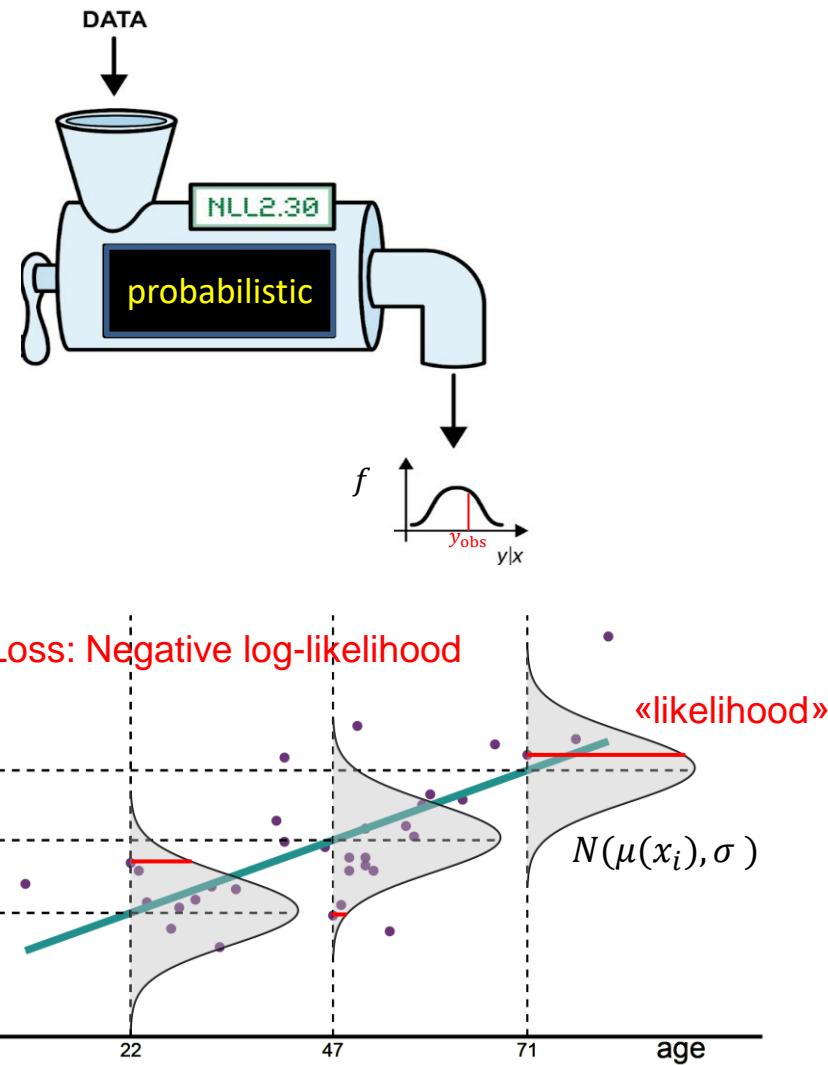
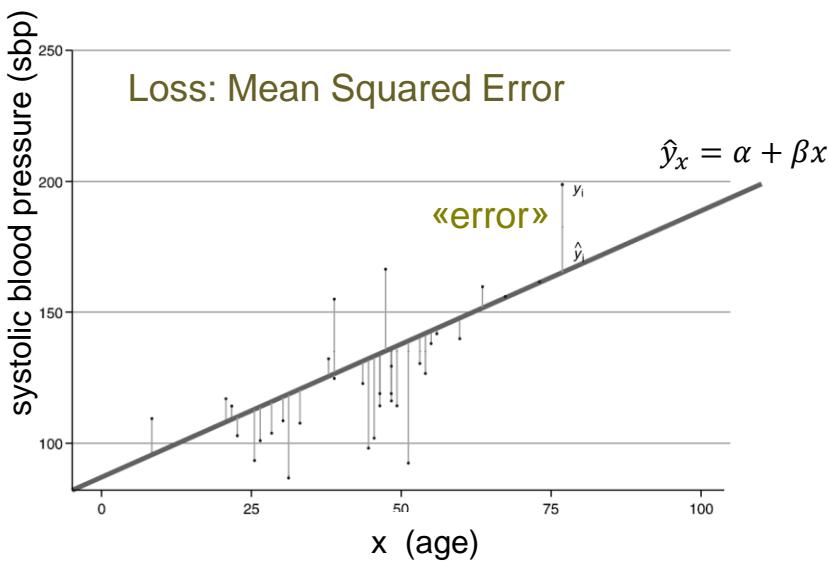
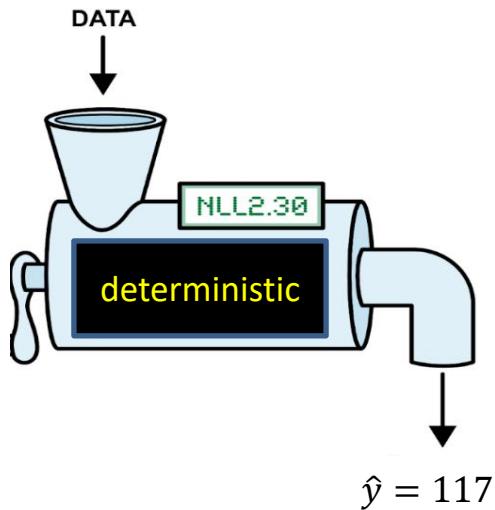


74



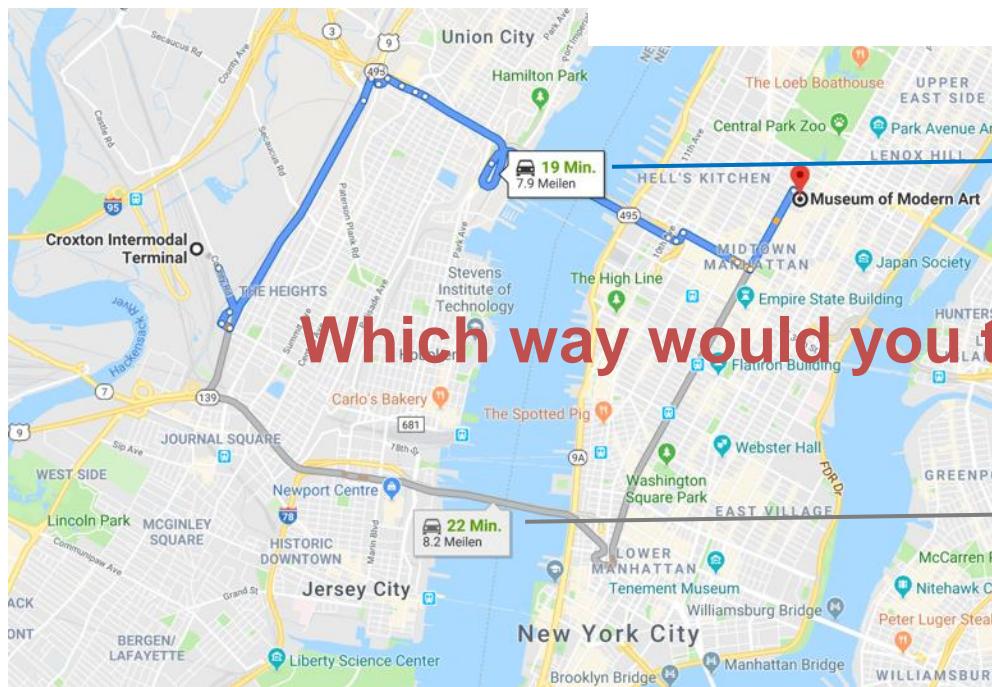
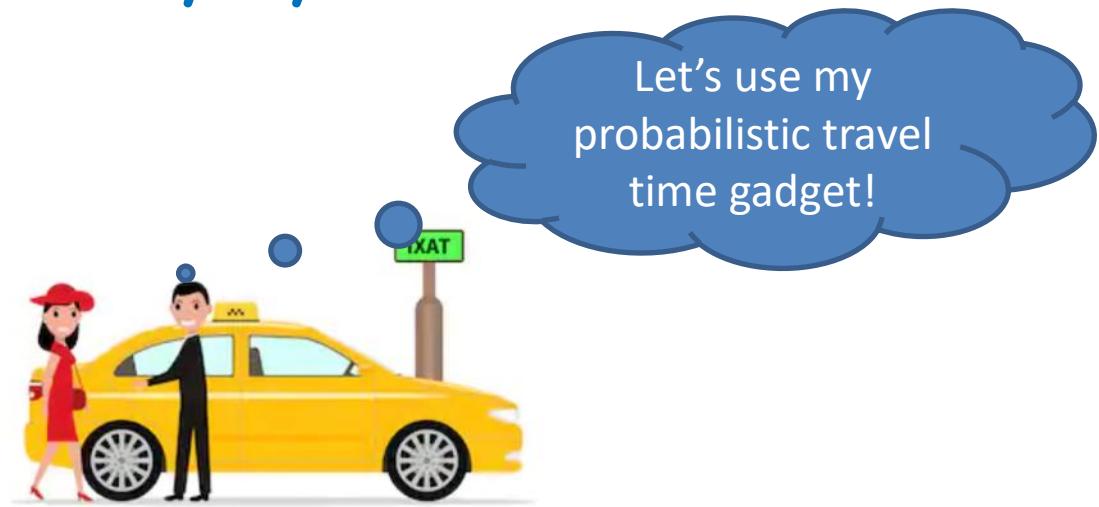
Conditional probability distribution $P(y|x)$ for continuous y aka $f_{Y|x}(y)$

Non-probabilistic versus probabilistic regression DL models



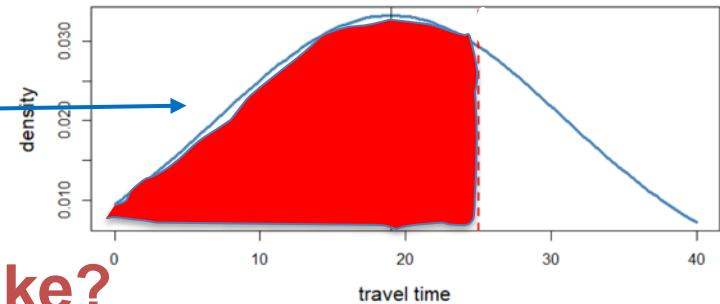
Probabilistic models in everyday tasks

You'll get 500\$ tip if I arrive at MOMA within 25 minutes!

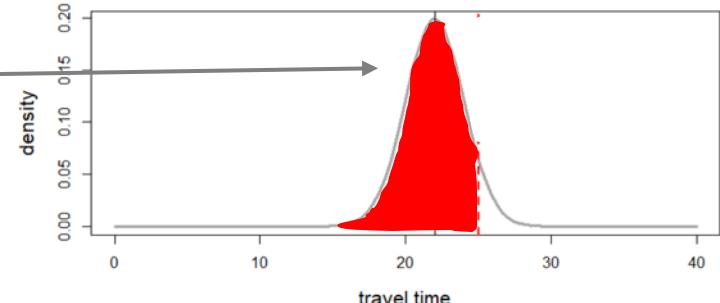


Which way would you take?

Chance to get tip: 69%



Chance to get tip: 93%

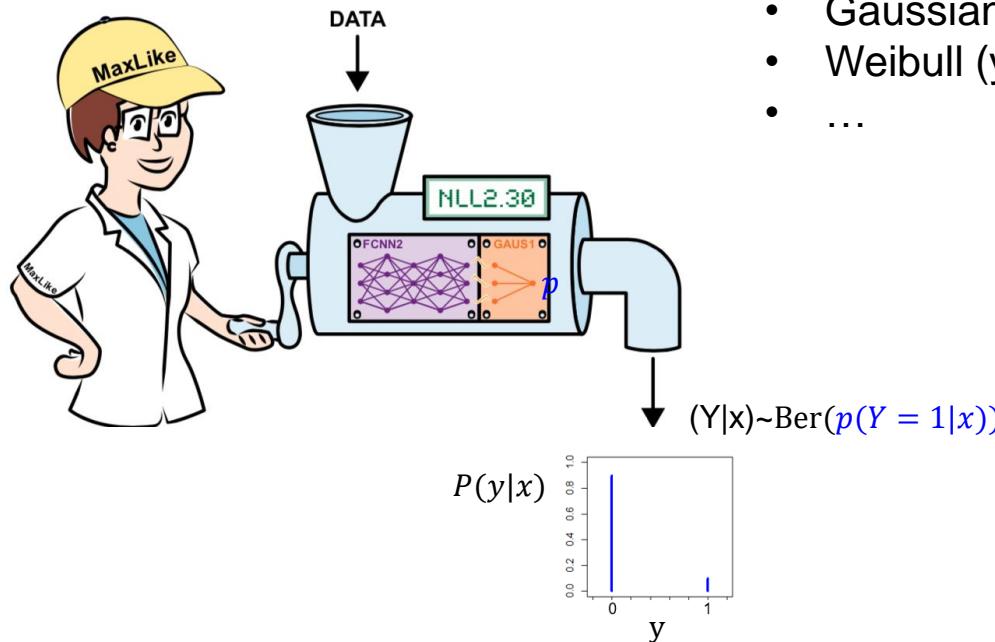
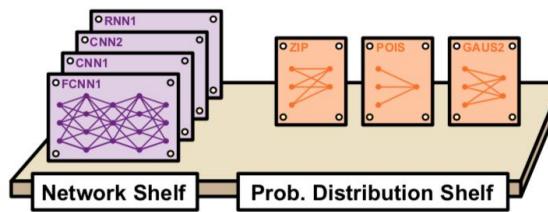


Guiding Theme of the course

- We treat DL as probabilistic models to model a conditional outcome distribution $P(y|x)$
- The models are fitted to training data with maximum likelihood (or Bayes)

Input X determines choice of the NN architectures:

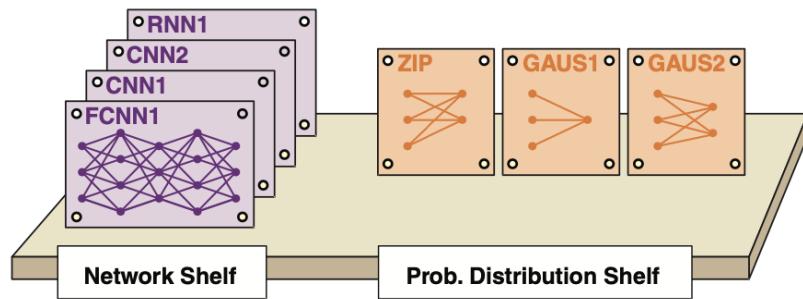
- Tabular data \rightarrow fcNN, tabPFN
- Images \rightarrow CNN, ViT
- Text \rightarrow LLMs



Family of outcome distribution, parameters are controlled by NN:

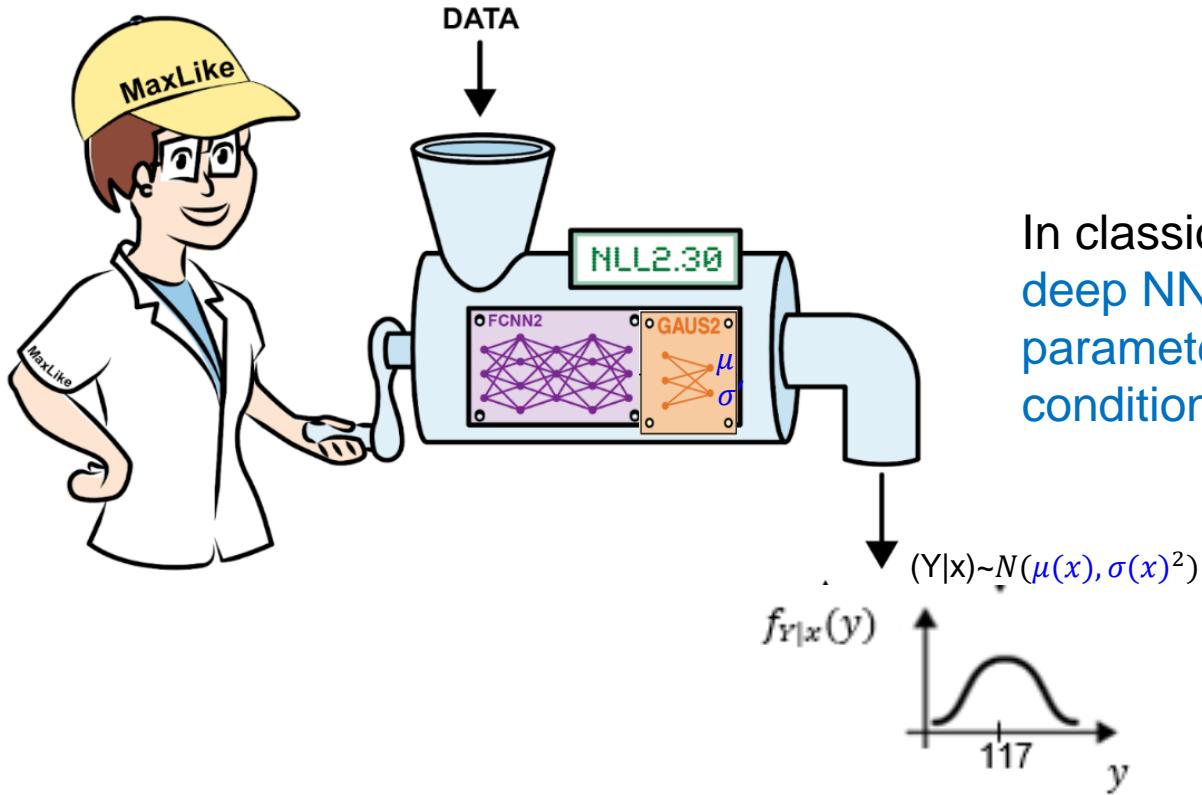
- Bernoulli (y : binary)
- Multi-Nomial (y : classes)
- Poisson (y : count)
- Negative Binomial (y : count)
- Gaussian (y : real number)
- Weibull (y : real number)
- ...

Probabilistic Deep Learning



Family of outcome distribution

- Gaussian
- Bernoulli
- Poisson
- ...



In classical probabilistic deep learning
deep NNs are used to estimate the
parameters of a pre-defined
conditional outcome distribution.

How to use a NN to do
Logistic Regression?

Recall logistic regression from statistics

$$Y \in \{0,1\}$$

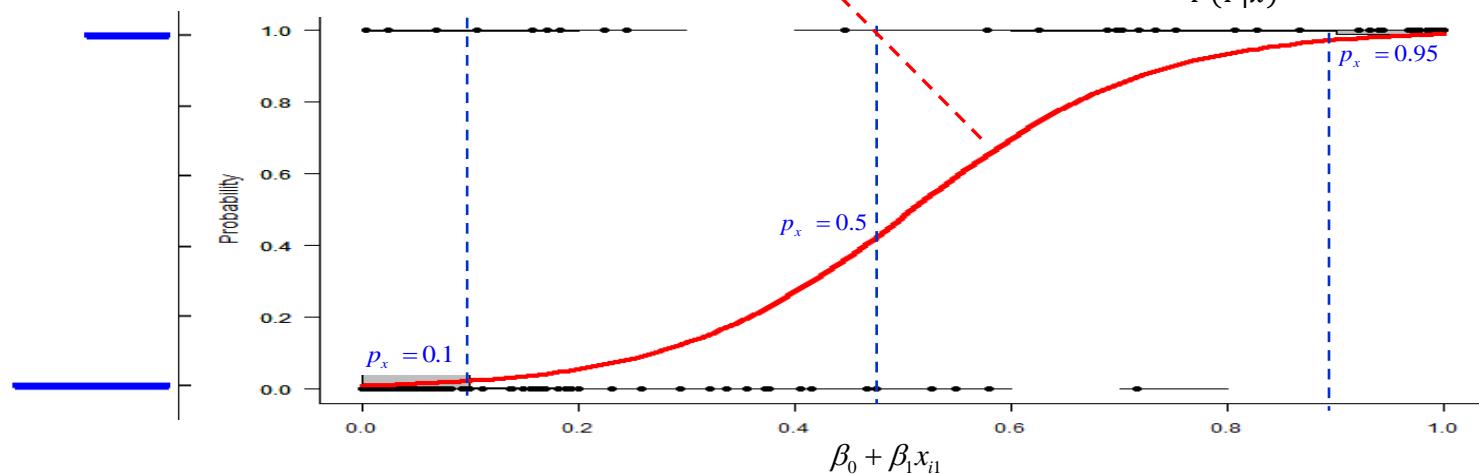
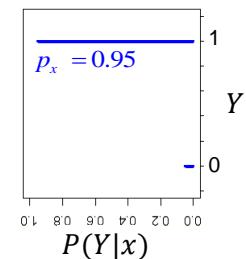
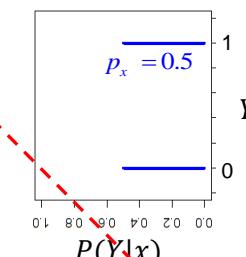
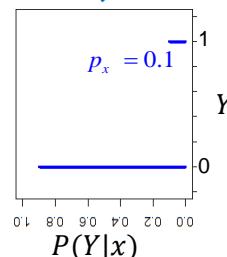
$$(Y|X_i) \sim \text{Ber}(p_{x_i})$$

$$\log\left(\frac{p_{x_i}}{1-p_{x_i}}\right) = \beta_0 + \beta_1 x_{i1}$$

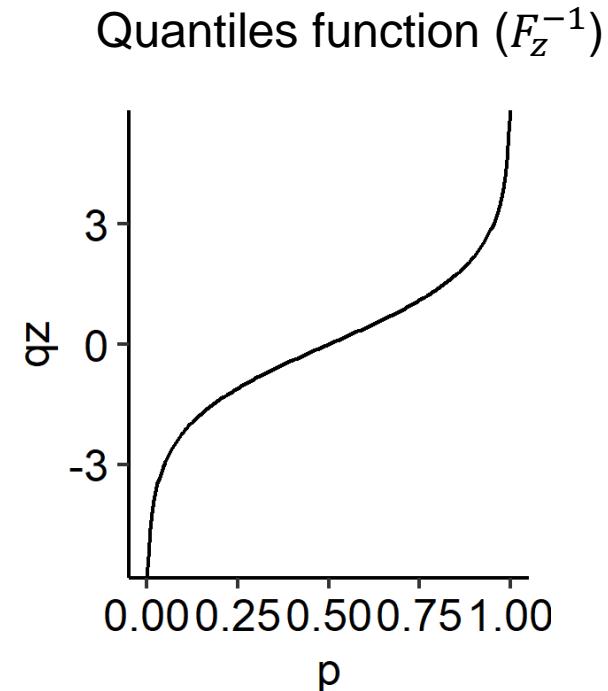
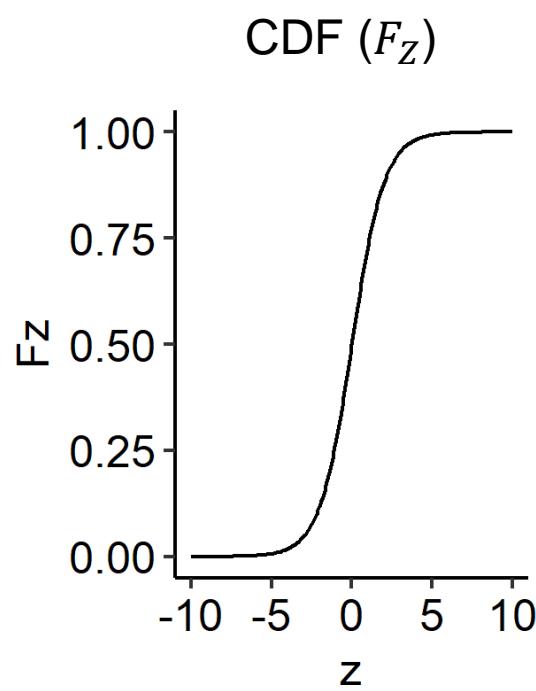
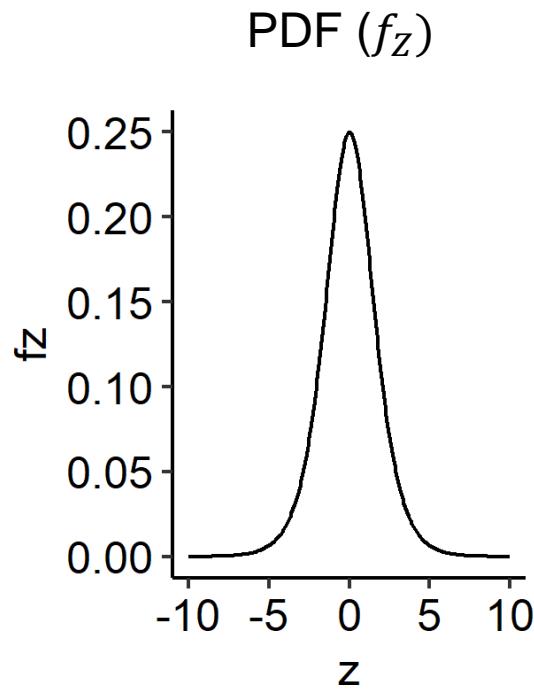
$$P(Y=1|x_i) = p_{x_i} = \sigma(\beta_0 + \beta_1 x_{i1}) = \frac{1}{1+e^{-(\beta_0 + \beta_1 x_{i1})}}$$

Logistic regression predicts a conditional Bernoulli $P(Y|x_i)$

$$P(Y|X=x) = \begin{cases} p_x & , y=1 \\ 1-p_x & , y=0 \end{cases}$$



Recall the Standard Logistic Distribution



$$f_z(z) = \frac{e^{-z}}{(1 + e^{-z})^2}$$

$$F_z(z) = \frac{1}{1 + e^{-z}}$$

$$F_z(z) = \text{expit}(z)$$

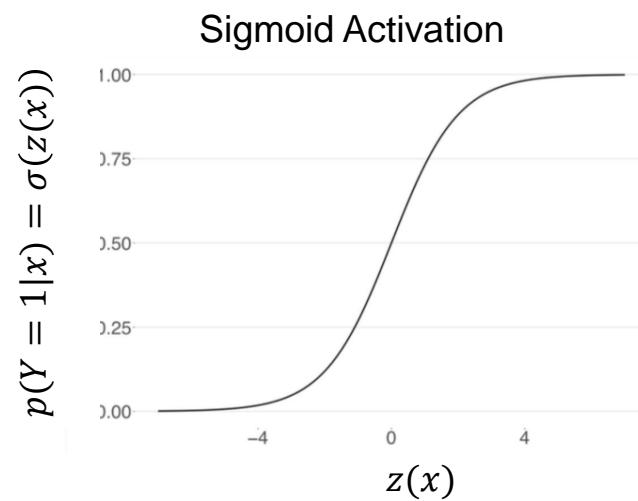
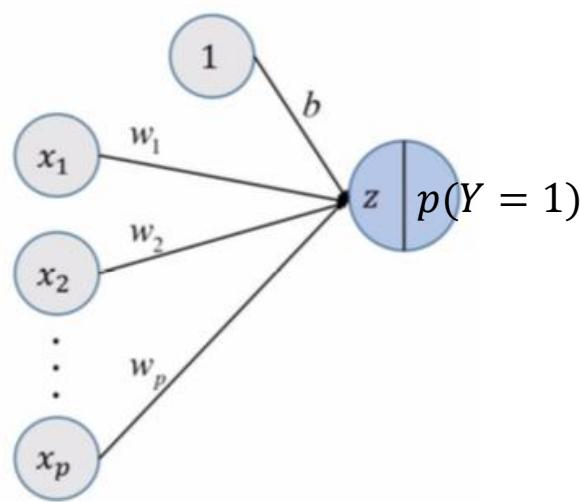
aka “sigmoid”

$$F_z^{-1}(p) = \log \frac{p}{1-p}$$

$$F_z^{-1}(p) = \log(\text{odds}) \\ = \text{logit}(p)$$

Logistic regression as NN

Activation in the last layer: “sigmoid”



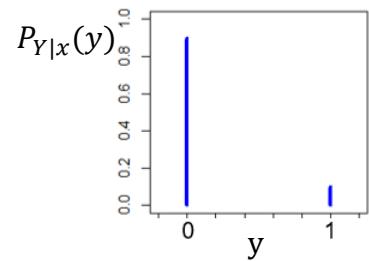
$$z(x) = b + x_1 \cdot w_1 + x_2 \cdot w_2 + \cdots + x_p \cdot w_p$$

$$p(Y = 1|x) = \sigma(z) = \sigma(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_{ip} x_{ip}) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_{ip} x_{ip})}}$$

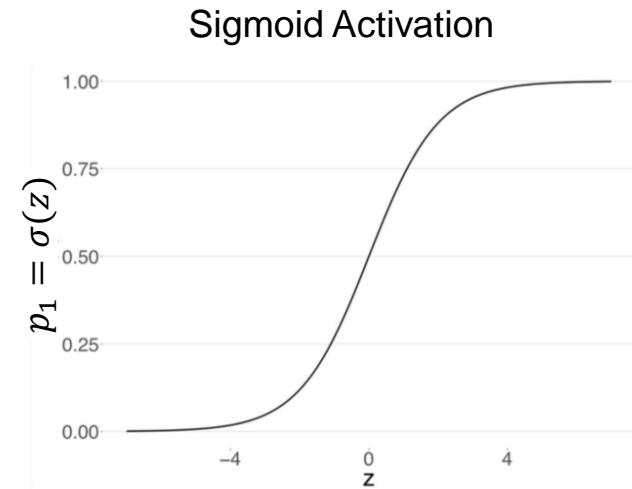
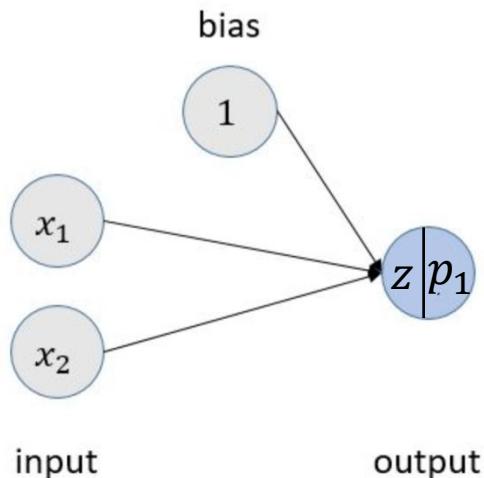
The output after the sigmoid activation:

$$P(y_i = 1|x_i) = \sigma(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_{ip} x_{ip})$$

→ A NN without hidden layers and
with one output node and a sigmoid-activation is a logistic regression model!



Exercise: Part 1



Model: The above network models the **probability** p_1 that a given banknote is false.

TASK (with pen and paper)

The weights (determined by a training procedure later) are given by

$$w_1 = 0.3, w_2 = 0.1, \text{ and } b = 1.0$$

What is the probability that a banknote, that is characterized by $x_1=1$ and $x_2 = 2.2$, is a faked banknote?

Pause bis 15:05 15:15

Colab Demo [Demo Time]

- Explain Colab based on
 - https://github.com/tensorchiefs/dlwl_eth25/blob/master/notebooks/00_colab_intro_forR.ipynb
 - Try to use AI (copilot) to translate R code to python
 - Try to use AI to generate code to solve your task
 - Show how to store Colab notebooks

Exercise run your first network



Notebook: [01_fcnn_with_banknotes](#)

Play with the code until and solve the task

🔧 Calculate the probability that a banknote is fake

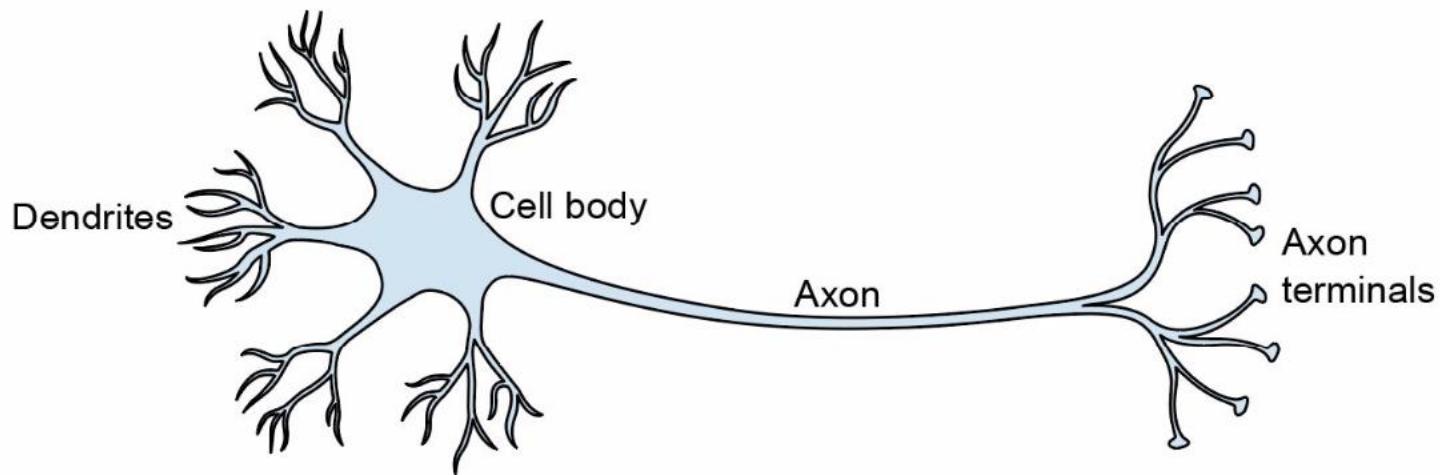
The banknote is described by two features $x_1 = 1$ and $x_2 = 2.2$ and the weights of the trained NN are $w_1 = 0.3$ and $w_2 = 0.2$ and the bias is $b = 1$. Use the following python function to calculate the probability that the banknote is fake.

```
[ ] 1 def predict_prob(x1, x2, w1, w2, b):  
2     return 1/(1+np.exp(-(w1*x1 + w2*x2 + b)))
```

🔑 Solution Code

▶ [Code anzeigen](#)

Neuron in a neural network: Biological Motivation



Sketch shows a model of single neuron in a brain.
The Dendrites collect signal from other neurons and accumulate it.
If signal exceeds a certain value the neuron “fires”, meaning a signal is sent to other neurons that coupled via the axon terminals

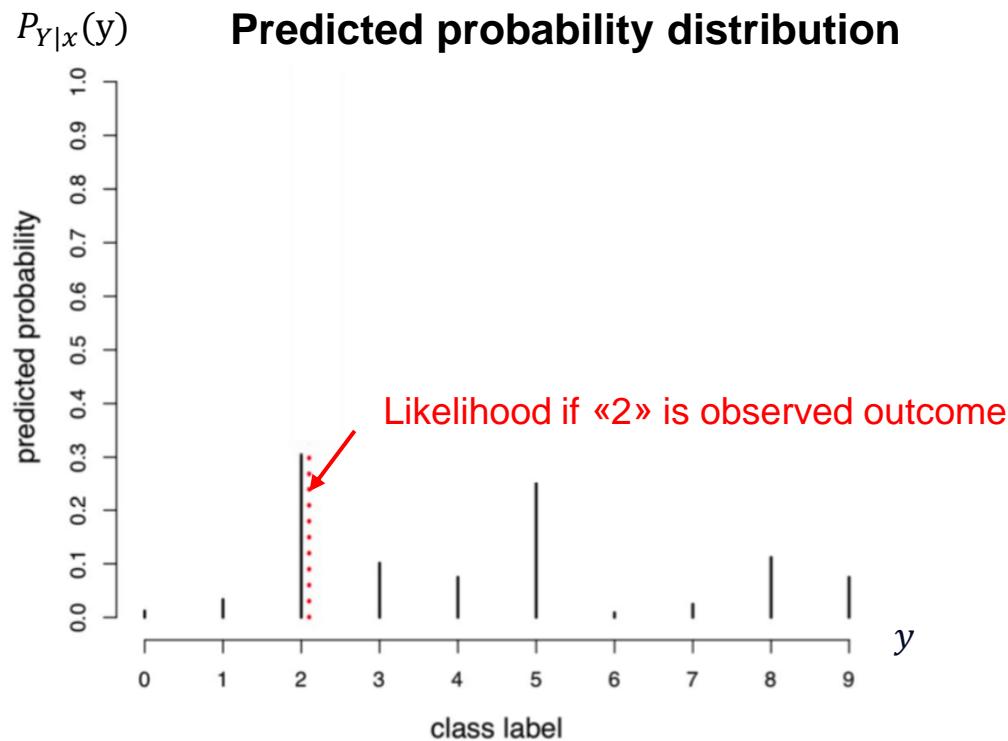
Neural networks are **loosely** inspired by how the brain works

How to use a NN to do
Multinomial Regression?

Multinomial regression aka multi-class regression

Multinomial regression is used when the outcome variable Y is categorial nominal and has more than two categories, c_1, c_2, \dots, c_k , that do not have an order.

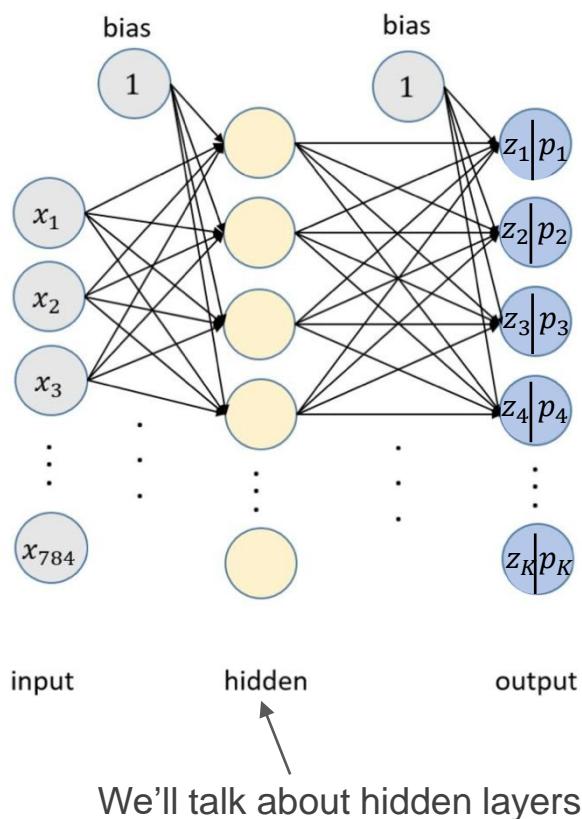
$$Y \in \{c1, c2, \dots, cK\}, P_{Y|x}(Y = c_k) = P(Y = c_k|x) = p_k(x), \sum_{k=1}^K p_k(x) = 1$$



Remark: If Y is not nominal but ordinal, often a proportional odds model is fitted which does assume more than just $\sum_k p_k = 1$.

Multi-class regression as NN

Activation in the last layer: “softmax”



$p_0, p_1 \dots p_9$ are probabilities for the classes c_1 to c_K .

Incoming to last layer $z_i \ i = 1 \dots K$

$$p_i = \frac{e^{z_i}}{\sum_{j=0}^K e^{z_j}}$$

Makes outcome positive

Ensures that p_i 's sum up to one

This activation is called *softmax*

NN with “softmax” activation in last layer outputs the probabilities for the classes.

How to use a NN to do
Linear Regression?

Recall linear regression from statistics

$$(Y|X = x_i) \sim N(\mu(x_i), \sigma^2)$$

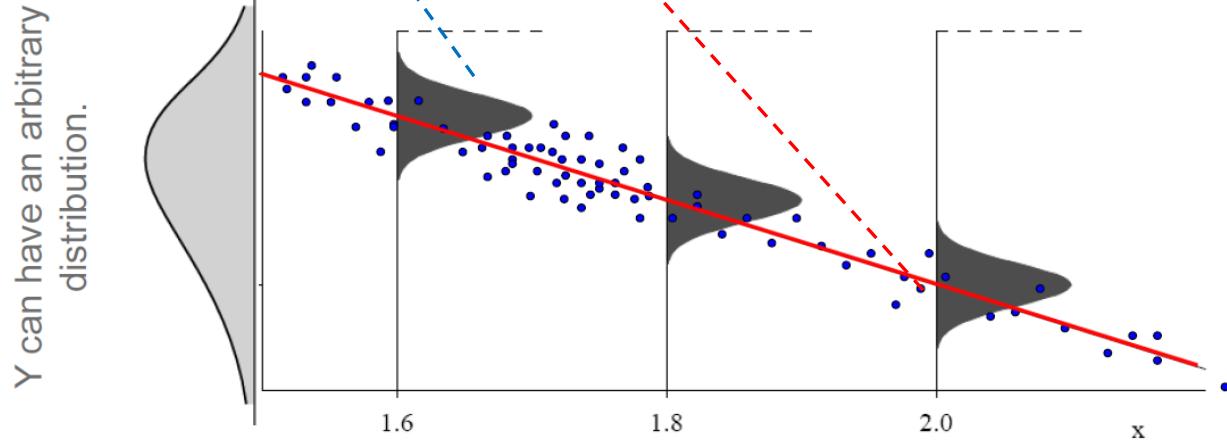
$$Y \in \mathbb{R}, \quad \mu_x \in \mathbb{R}$$

Probabilistic linear regression predicts for each input x_i a Gaussian conditional probability distribution for the output $P(Y|x_i)$ that assigns each possible value of Y a likelihood.

point prediction
probabilistic prediction

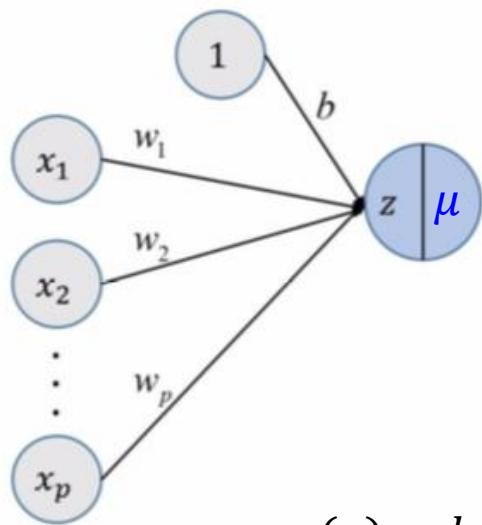
$$\begin{aligned}y_i &= \mu(x_i) + \varepsilon_i = \beta_0 + \beta_1 \cdot x_{i1} + \varepsilon_i \\ \hat{E}(Y_{X_i}) &= \hat{\mu}_{x_i} = \hat{\beta}_0 + \hat{\beta}_1 \cdot x_{i1} \\ \text{Var}(Y_{X_i}) &= \text{Var}(\varepsilon_i) = \sigma^2 \\ \varepsilon_i &\sim N(0, \sigma^2)\end{aligned}$$

Variance is assumed to be constant



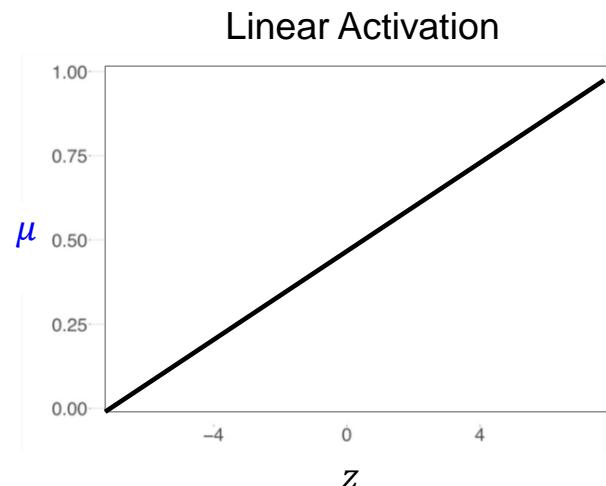
Linear Regression as NN

Activation in the last layer: “linear”



$$z(x) = b + x_1 \cdot w_1 + x_2 \cdot w_2 + \dots + x_p \cdot w_p$$

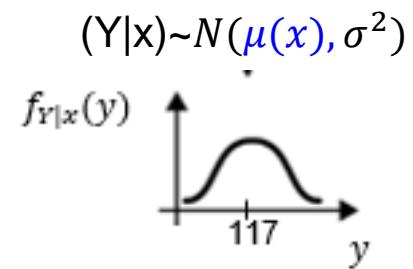
$$\mu(x) = z(x)$$



The output after the linear activation can be interpreted as estimated expected value of the Gauss $\mu(x) = E(Y|x)$

→ A fcNN without hidden layers and with one output node and a sigmoid-activation is a logistic regression model!

Remark: after fitting the linear predictor $\hat{\mu}(x)$ the constant variance is estimated from the residuals



How to use a NN to do
Non-linear Regression?

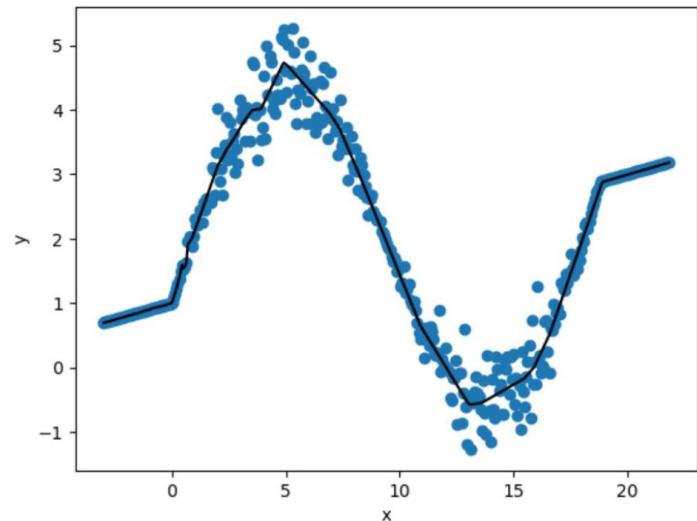
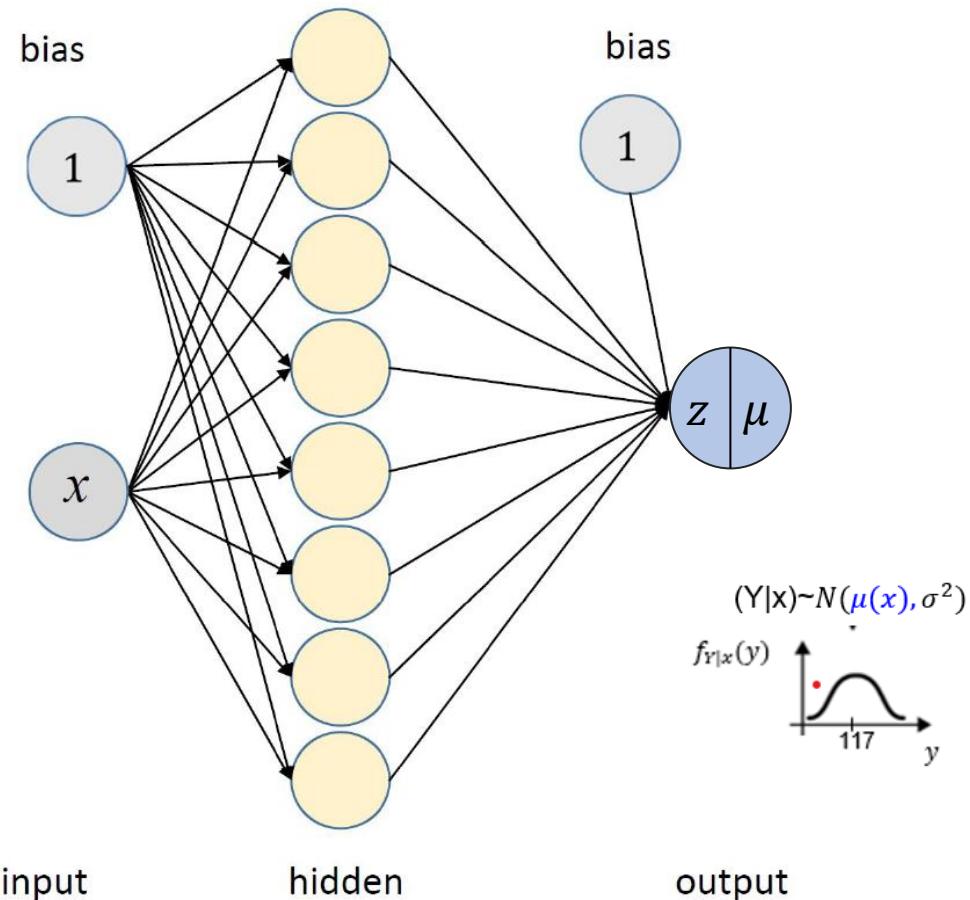
A close-up shot from the movie Inception. Two men in dark suits are facing each other. The man on the left has his eyes closed and a slight smile. The man on the right has his eyes open and is looking directly at the other man. The lighting is warm and dramatic.

WE NEED TO GO

DEEPER

Non-linear Regression as NN

Activation in hidden layer non-linear, in last layer linear



Here, the conditional outcome distribution is a Gaussian with a μ that depends non-linearly on x 😊 a variance σ^2 that is constant ⚡

Remark: It has been shown, that a NN with one hidden layers can approximate any function, if the the hidden layer holds enough neurons.

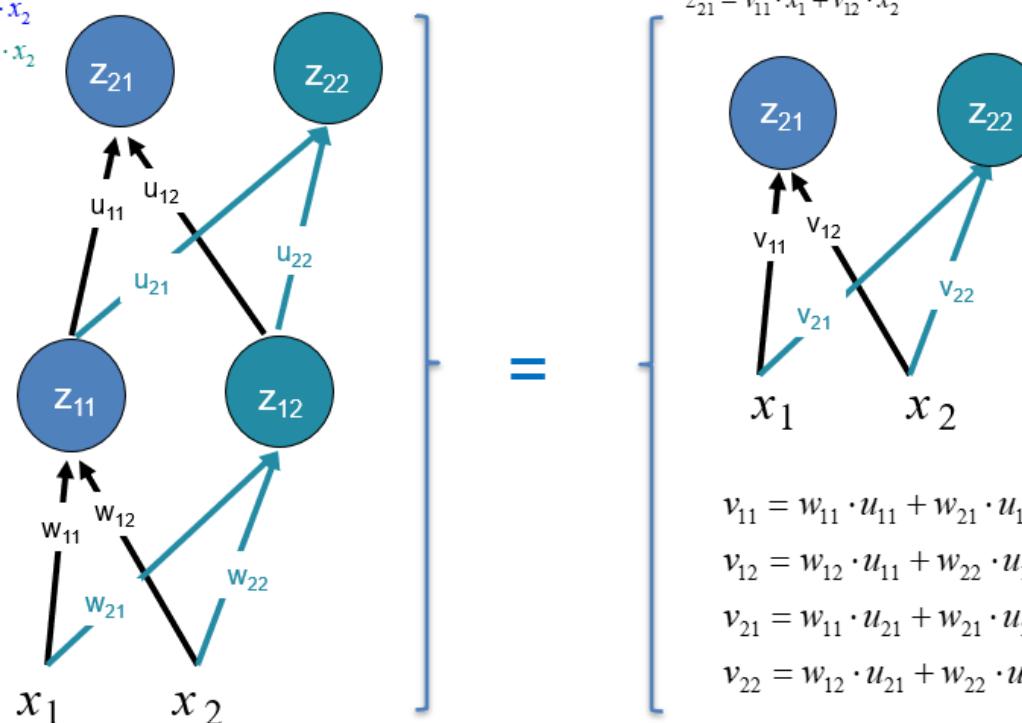
Hidden layer need non-linear activation functions

2 linear layers can be replaced by 1 linear layer -> can't go deep with linear layers!

$$\begin{aligned}z_{21} &= z_{11} \cdot u_{11} + z_{12} \cdot u_{12} = (w_{11} \cdot x_1 + w_{12} \cdot x_2) \cdot u_{11} + (w_{21} \cdot x_1 + w_{22} \cdot x_2) \cdot u_{12} \\&= x_1 \cdot (w_{11} \cdot u_{11} + w_{12} \cdot u_{12}) + x_2 \cdot (w_{21} \cdot u_{11} + w_{22} \cdot u_{12})\end{aligned}$$

$$z_{11} = w_{11} \cdot x_1 + w_{12} \cdot x_2$$

$$z_{12} = w_{21} \cdot x_1 + w_{22} \cdot x_2$$



$$z_{21} = v_{11} \cdot x_1 + v_{12} \cdot x_2$$

$$v_{11} = w_{11} \cdot u_{11} + w_{21} \cdot u_{12}$$

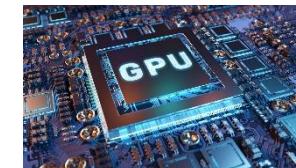
$$v_{12} = w_{12} \cdot u_{11} + w_{22} \cdot u_{12}$$

$$v_{21} = w_{11} \cdot u_{21} + w_{21} \cdot u_{22}$$

$$v_{22} = w_{12} \cdot u_{21} + w_{22} \cdot u_{22}$$

Computation NNs with linear activation are matrix multiplications → GPUs

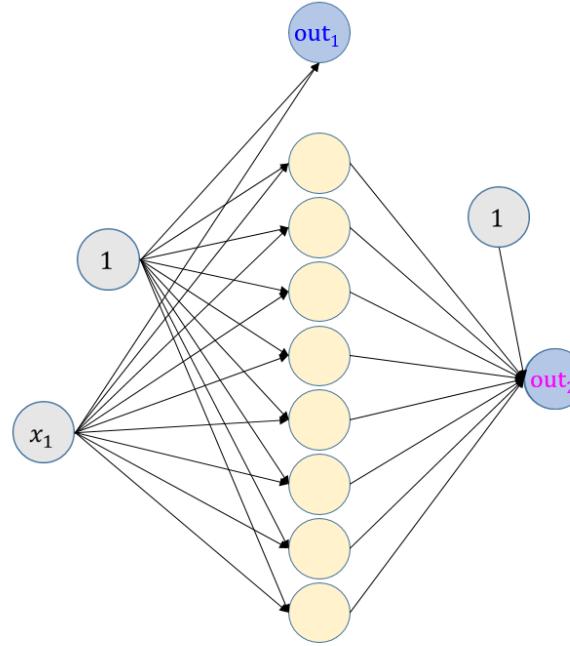
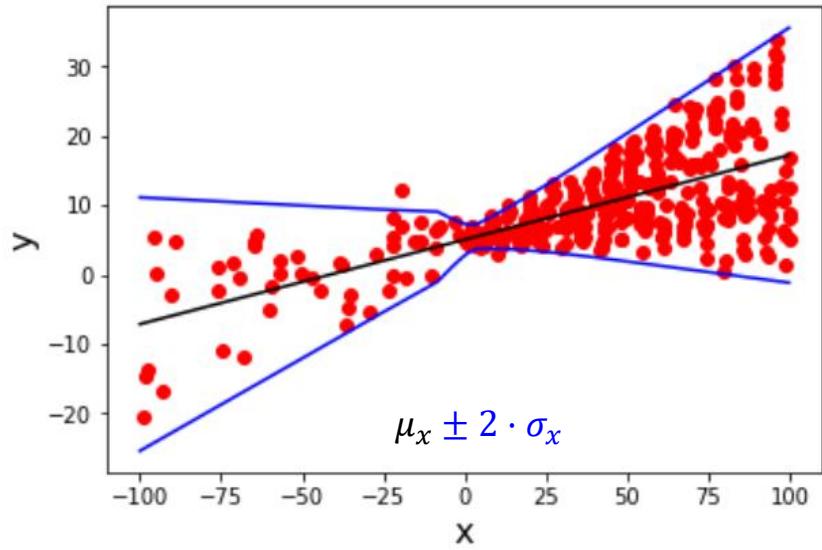
$$z = (x \cdot W) \cdot U = x \cdot (W \cdot U) = x \cdot V$$



Remark: biases are ignored here, but do not change fact

NN for linear regression with not-constant variance

$$Y_{X_i} = (\mathbf{Y} | \mathbf{X}_i) \sim N(\mu_{x_i}, \sigma_x^2)$$



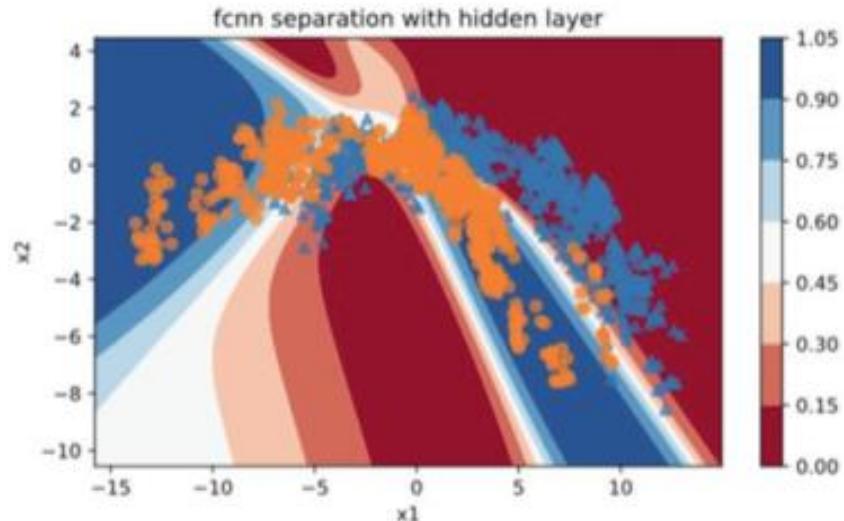
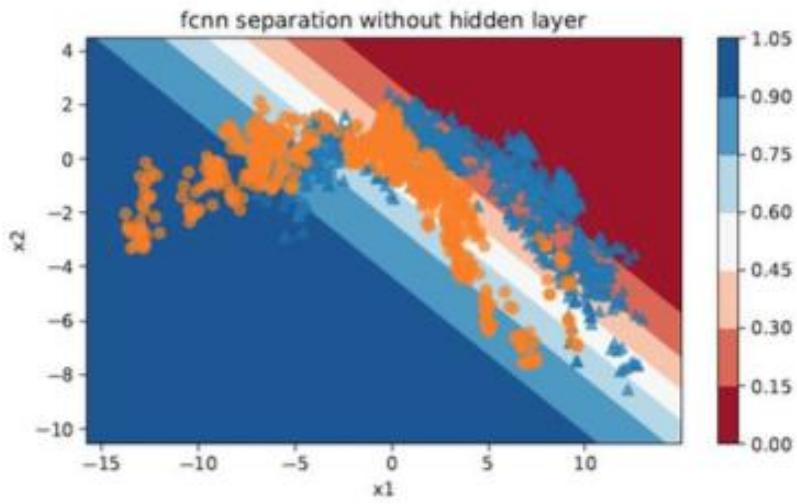
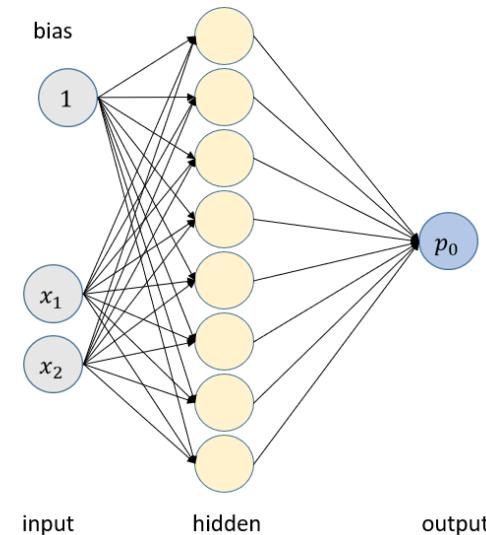
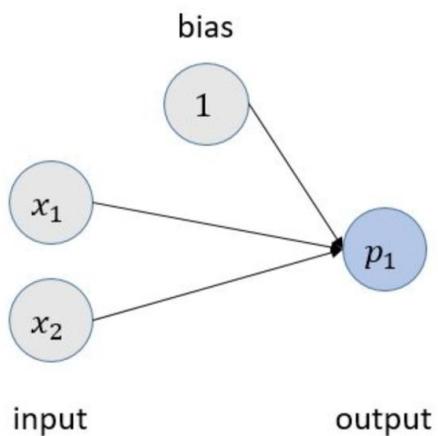
$$\mu_{x_i} = out_{1_i}$$

$$\sigma_{x_i} = e^{out_{2_i}}$$

Note: You will see soon that the maximum likelihood principle allows to learn the variance σ_x of the conditional Gaussian without requiring a ground truth for σ_x .

Non-linear decision boundary in binary classification

Activation in hidden layer “ReLU” in last layer: “linear”



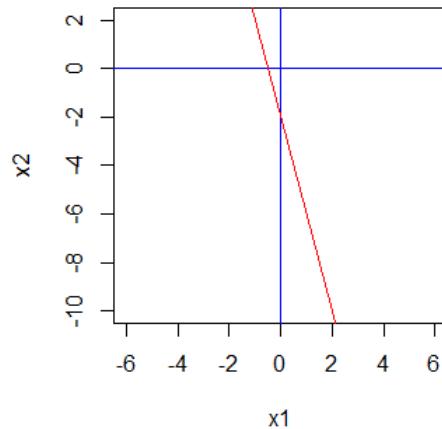
Recall: Logistic regression yield linear/planar decision curves

Logistic regression model:

$$\ln\left(\frac{p}{1-p}\right) = 1 + 2x_1 + 0.5x_2$$

Determine the **decision curve between Y=1 and Y=0** in the feature room which is spanned by x_1 and x_2 and draw it in the following plot x_2 and x_1 .

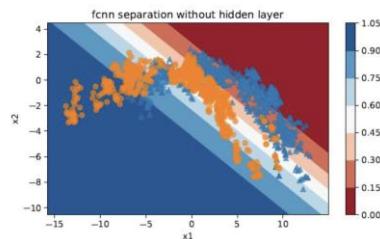
Hint: on the separation curve should hold: $P(Y = 1|x) = 0.5$
-> plug in 0.5 for p and solve for x_2 .



$$\ln\left(\frac{0.5}{1-0.5}\right) = 0 = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

$$x_2 = -\frac{\beta_0}{\beta_2} - \frac{\beta_1}{\beta_2} \cdot x_1$$

$$x_2 = -2 - 4 \cdot x_1$$



Hence a NN with 1 output neuron with sigmoid-activation w/o hidden layer have linear decision boundary

Deep Learning Libraries

Deep-learners talk about tensors - What is a tensor?

For Deep Learning: A tensor is an array with several indices (like in numpy). Order (a.k.a rank) is the number of indices and shape is the range.

```
In [1]: import numpy as np

In [2]: T1 = np.asarray([1,2,3]) #Tensor of order 1 aka Vector
T1

Out[2]: array([1, 2, 3])

In [3]: T2 = np.asarray([[1,2,3],[4,5,6]]) #Tensor of order 2 aka Matrix
T2

Out[3]: array([[1, 2, 3],
   [4, 5, 6]])

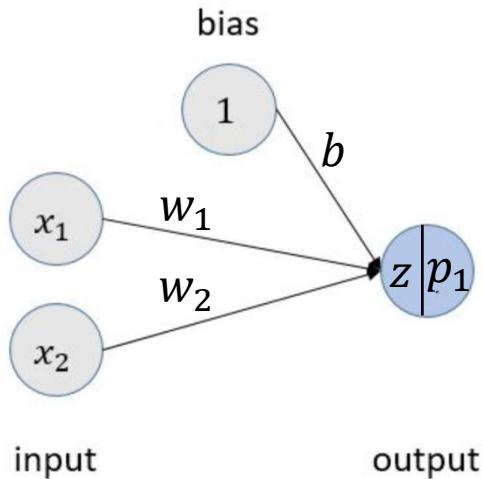
In [4]: T3 = np.zeros((10,2,3)) #Tensor of order 3 (Volume like objects)

In [6]: print(T1.shape)
print(T2.shape)
print(T3.shape)

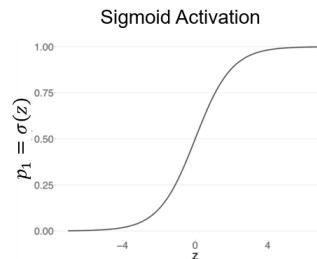
(3,)
(2, 3)
(10, 2, 3)
```

Tensors (as a generalization of matrix and vectors) are the basic quantities in DL.

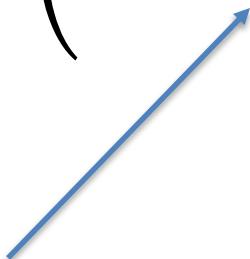
GPUs love Vectors



$F^{\mu\nu}$

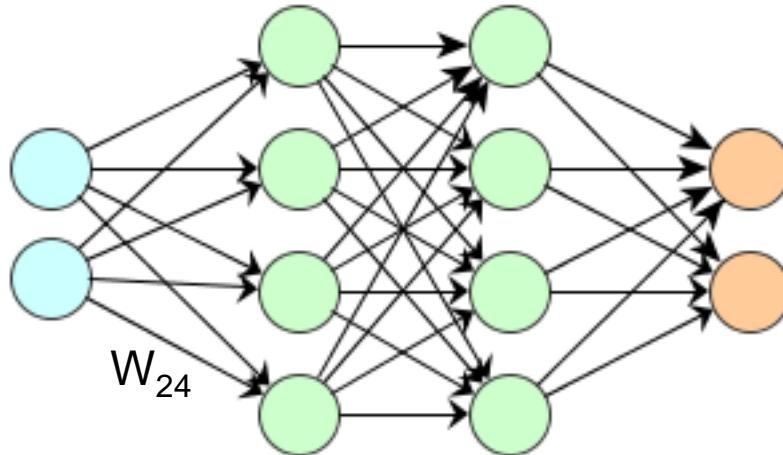


$$p_1 = \text{sigmoid} \left((x_1 \quad x_2) \cdot \begin{pmatrix} w_1 \\ w_2 \end{pmatrix} + b \right)$$



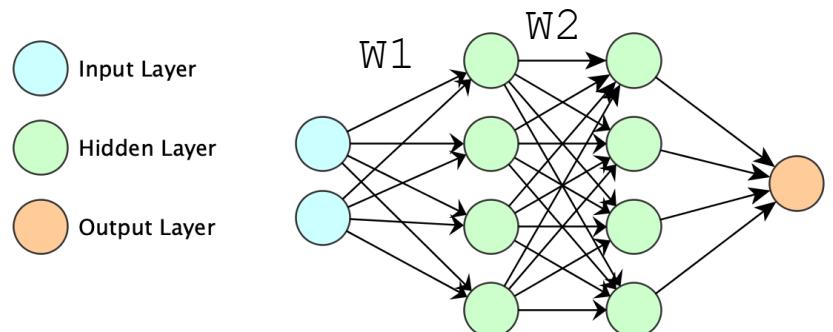
DL: better to have column vectors

Typical Tensors in Deep Learning



- The input can be understood as a vector
- A mini-batch of size 64 of input vectors can be understood as tensor of order 2
 - (index in batch, x_j)
- The weights going from e.g. Layer L_1 to Layer L_2 can be written as a matrix (often called W)
- A mini-batch of size 64 images with 256,256 pixels and 3 color-channels can be understood as a tensor of order 4.

Structure of the network



In code:

```
## Solution 2 hidden layers
def predict_hidden_2(X):
    hidden_1 = F.sigmoid(X @ W1 + b1) # @ is Matrix Mult.
    hidden_2 = F.sigmoid(hidden_1 @ W2 + b2)
    return F.sigmoid(hidden_2 @ W3 + b3)
```

In math ($f = \text{sigmoid}$) and $b1=b2=b3=0$

$$p = f(f(f(x W^1)W^2))$$

Defining the NN with code is done with so-called low-level libraries (here pytorch)
Looks a bit like onions, matryoshka (Russian Dolls) or lego bricks...

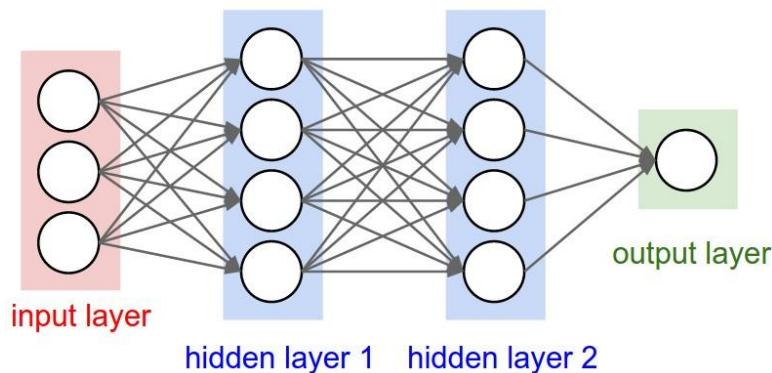
Question: Dimensions of $W1$, $W2$, and $W2$

Keras as High-Level library



We use Keras as high-level library

Libraries make use of the Lego like block structure of networks



Under the “hood” of Keras runs the engine;

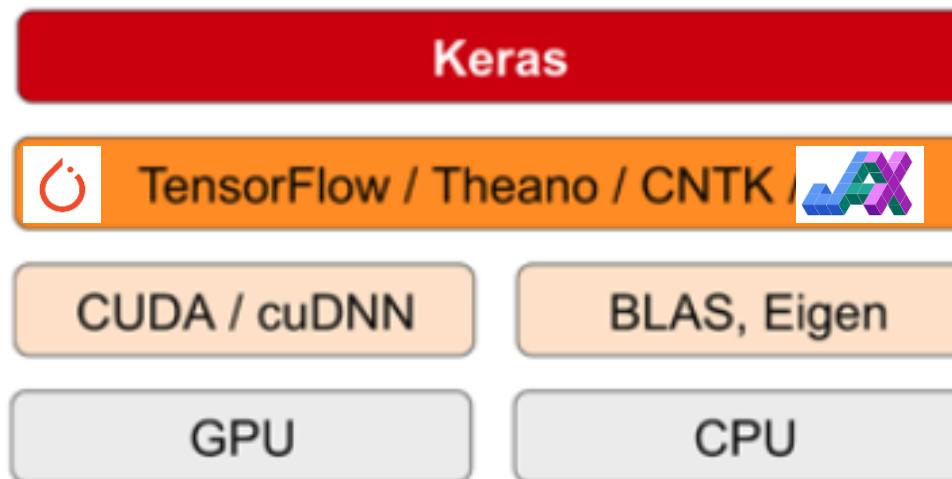


Figure: From Deep Learning with Python, Francois Chollett

Load packages

- Select “Engine” (before `import keras`)

 PyTorch TensorFlow

```
import os
os.environ["KERAS_BACKEND"] = "torch"      # or 'jax' or 'tensorflow'
import keras
```

Load building blocks

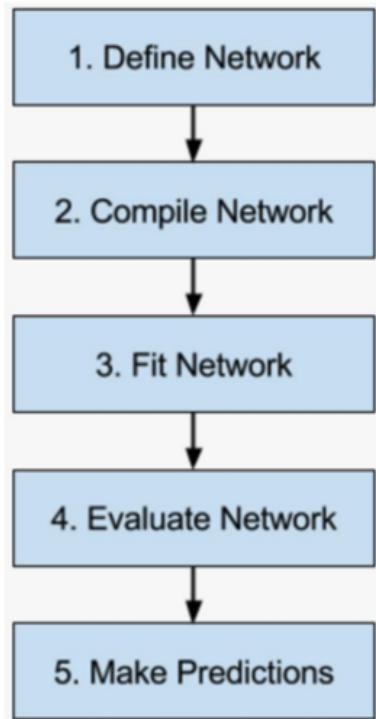
- Layers, functions etc.

```
from keras.models import Sequential
from keras.layers import Dense,
from keras.optimizers import SGD
```



Define the network

Sequential API, layers output are the input for the next layer, Alternative Functional API



```
# Define fcNN
model = Sequential()
model.add(Dense(8,
               input_shape=(2,),
               activation='sigmoid'))
model.add(Dense(1,
               activation='sigmoid'))
```

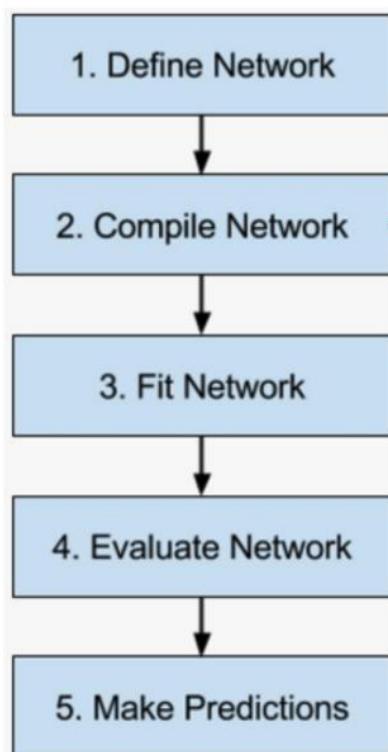
Hidden layer with 8 Neurons,
Activation function: sigmoid

Last layer with one output neuron for
binary classification

Input shape needs to be defined only at the beginning.

Alternative: `input_dim=2`, Functional API or Sequential API

Compile the network



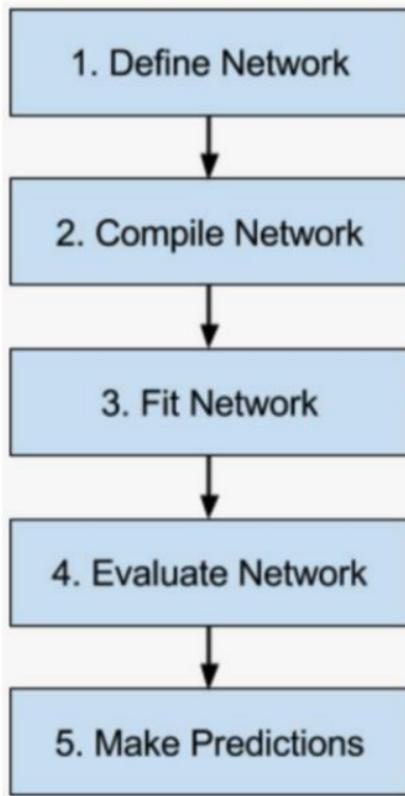
```
model.compile(optimizer=SGD(learning_rate=0.01),  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

Loss Function to optimize
Here: Binary CE

Which optimizer should be used?
Here Stochastic Gradient Descent

Which metrics do we want to track,
Here: Accuracy

Fit the network



Data Tensor X for training

Label Tensor Y

Validation at
End of each epoch

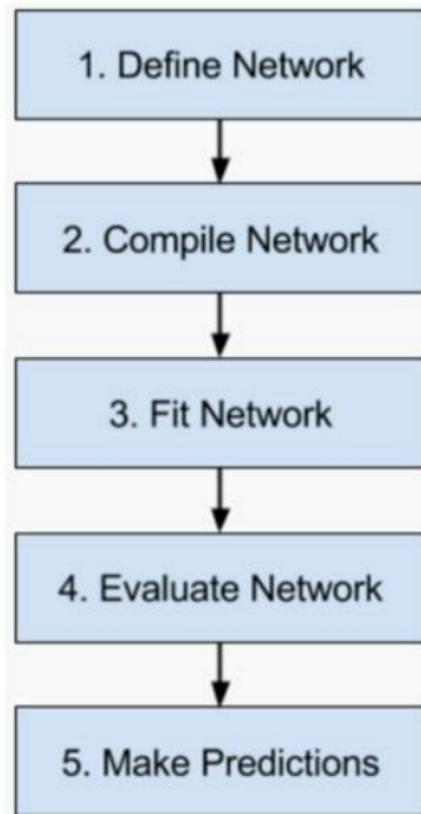
```
history = model.fit(x_train,  
                     y_train,  
                     validation_data=(x_val, y_val),  
                     epochs=1000,  
                     batch_size=10,  
                     verbose=1)
```

How many time do
we feed the whole
dataset to the
model

Should we see the whole
output? If no verbose = 0

How many samples per
batch, 1 batch = 1 iteration
of weight updates

Evaluate the network



Unseen (to the model) Test Data X with labels Y

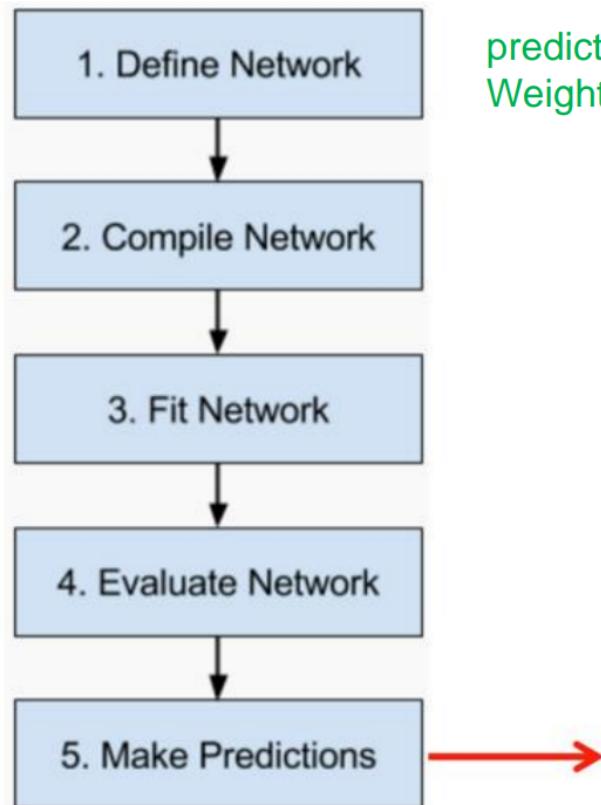
```
model.evaluate(x_test, y_test, verbose=0)
```

Test Loss: 0.33591118454933167, Test Accuracy: 0.8581818342208862

Remember from the
.compile

```
loss='binary_crossentropy',  
metrics=['accuracy'])
```

Make Predictions



predict with the model,
Weights are fixed now

First 10 rows
(observations) from
Testdata

All features, here: 2

```
model.predict(x_test[0:10,:])
```

```
1/1 [=====] - 0s 148ms/step
array([[5.9251189e-01],
       [8.8978779e-01],
       [7.8563684e-01],
       [9.1934395e-01],
       [6.1078304e-01],
       [8.2838058e-01],
       [2.5862646e-01],
       [3.1314052e-05],
       [2.6938611e-01],
       [3.7005499e-02]], dtype=float32)
```

Each of the observations
Gets a probability to
Belong to class 1

Building a network (API Styles)

Three API styles

- The Sequential Model
 - Dead simple
 - Only for single-input, single-output, sequential layer stacks
 - Good for 70% of use cases
- The functional API
 - Like playing with Lego bricks
 - Multi-input, multi-output, arbitrary static graph topologies
 - Good for 95% of use cases
- Model subclassing
 - Maximum flexibility
 - Larger potential error surface

Sequential API

The Sequential API

```
import keras
from keras import layers

model = keras.Sequential()
model.add(layers.Dense(20, activation='relu', input_shape=(10,)))
model.add(layers.Dense(20, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

model.fit(x, y, epochs=10, batch_size=32)
```

Functional (do you spot the error?)

The functional API

```
import keras
from keras import layers

inputs = keras.Input(shape=(10,))
x = layers.Dense(20, activation='relu')(x)
x = layers.Dense(20, activation='relu')(x)
outputs = layers.Dense(10, activation='softmax')(x)

model = keras.Model(inputs, outputs)
model.fit(x, y, epochs=10, batch_size=32)
```

Functional (do you spot the error?)

The functional API

```
import keras
from keras import layers

inputs = keras.Input(shape=(10,))
x = layers.Dense(20, activation='relu')(x) ←'inputs'
x = layers.Dense(20, activation='relu')(x)
outputs = layers.Dense(10, activation='softmax')(x)

model = keras.Model(inputs, outputs)
model.fit(x, y, epochs=10, batch_size=32)
```

Projects

Overview of projects

- ~3 Persons
- Instead of a test you do a short presentation and a poster to be presented at the last day
- Topics free but should have to do with probabilistic deep learning
- Data and Ideas can be found e.g. <https://www.kaggle.com/>
- 2 Example Projects
- Approach us if you want to do your own

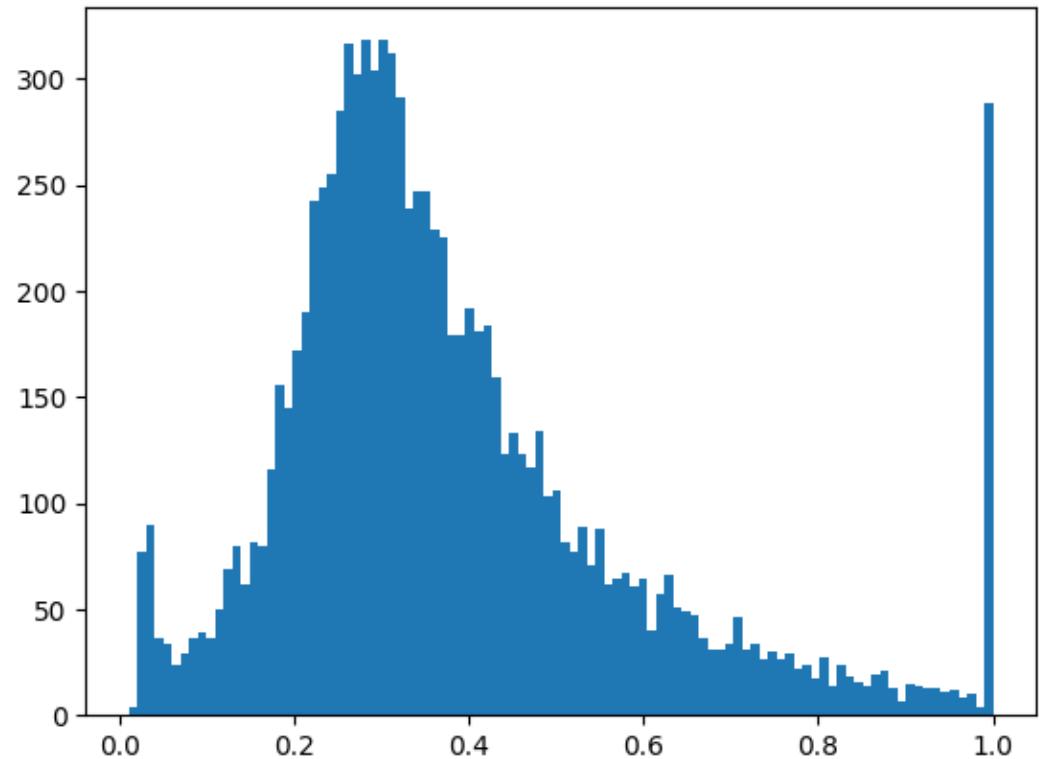
Animal Cuteness

PetFinder.my - Pawpularity Contest

~10000 Photos of Animals



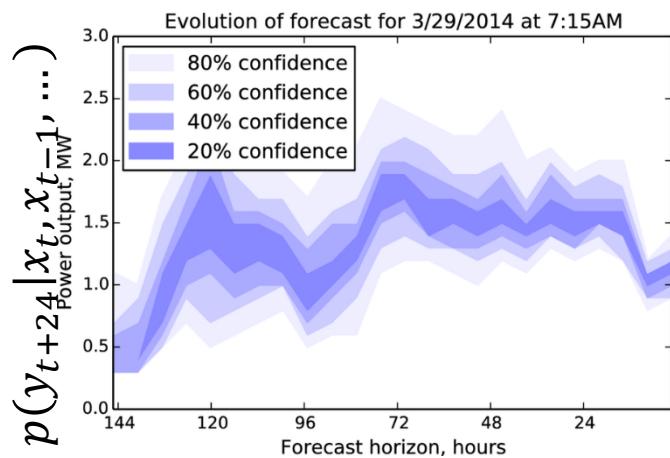
Predict cuteness score (0-100%)



Weather Data

Aim of the project

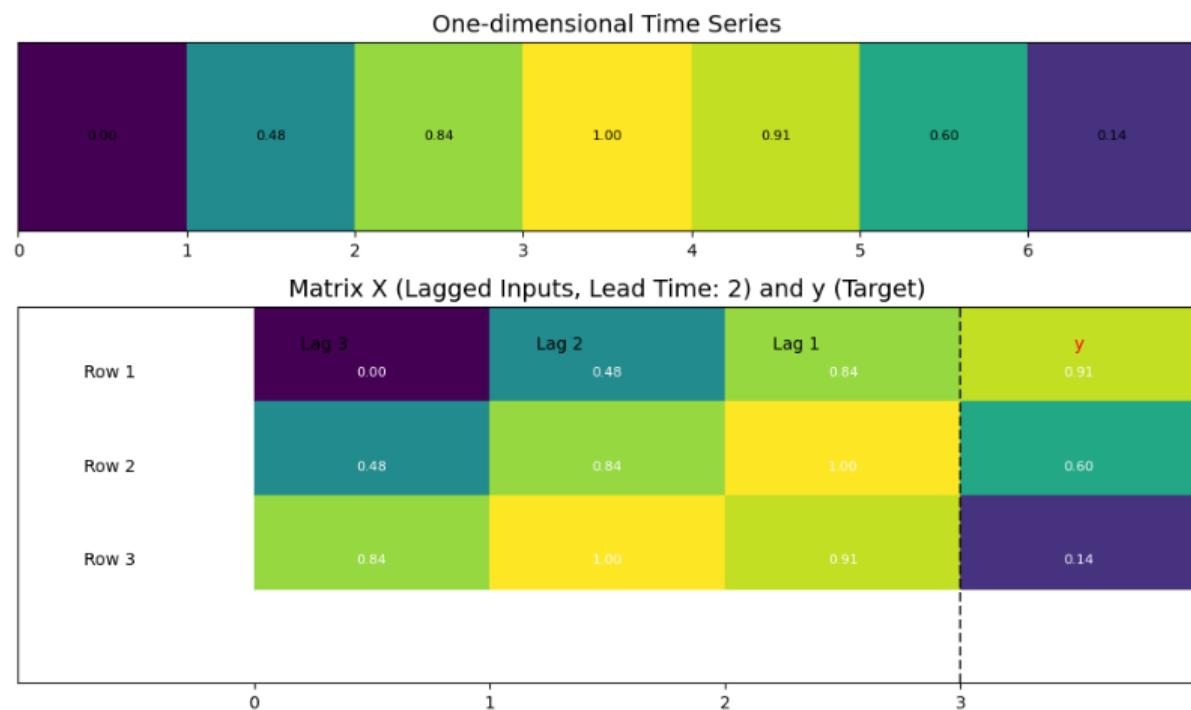
- Provide probabilistic forecast for a weather variable (e.g. temperature)
 - E.g. $p(y_{t+24}|x_t, x_{t-1}, \dots)$



- Data (possibilities)
 - Historic observations from DWD stations (starter set contains Konstanz)
 - Historic observations using Historical Weather API (from openmeteo)
 - Historic observations and historic numerical weather forecasts
 - For Konstanz data is available from Nov 2024-
 - For other places openmeteo provides data

Using historical Observations

Training Data from historical observations



- Simple Demo NN (non-probabilistic)
- Starter set
https://github.com/tensorchiefs/dlwbl_eth25/blob/master/project_starters/weather/weather_pred.ipynb

Pause bis 17:15

Exercise: use keras for a NN with hidden layers



Notebook: [01_fcnn_with_banknotes](#)

Play around with the Keras code for building and training NNs!

We discuss afterwards the loss function and how fitting works.

