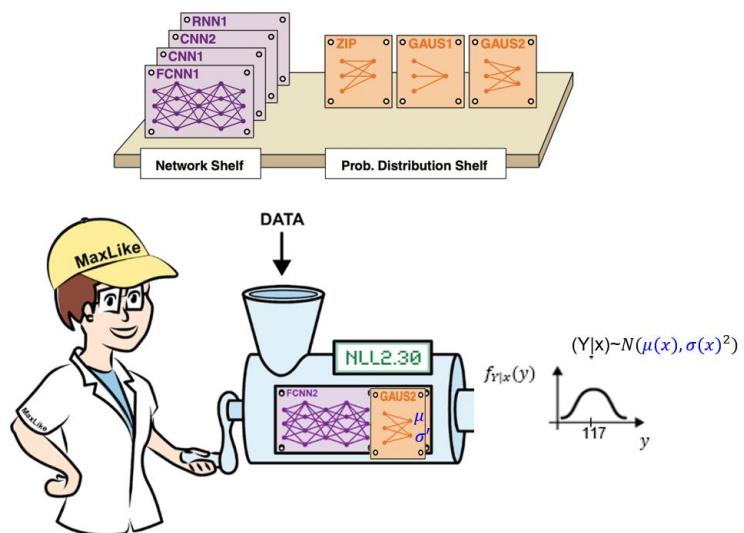


# WBL Probabilistic Deep Learning:: Day 2

Preliminary version

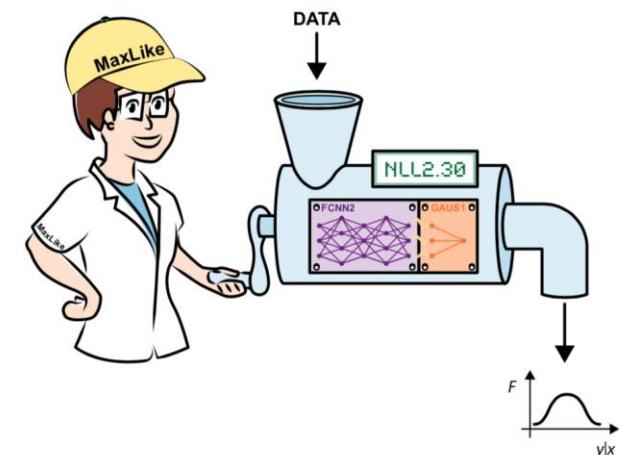
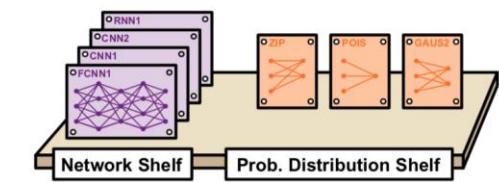
*Beate Sick, Oliver Dürr*

Day 2: NN training and different NN architectures



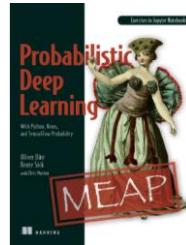
# Topics of today

- Training of NN
  - Negative Log Likelihood (NLL) as loss function
  - Optimization via stochastic gradient descent (SGD) or a variant (e.g. adam)
- NN architectures
  - Fully connected NN for tabular data
  - Convolutional Neural Networks (CNN) for images
  - Next time: Transformer
- Tricks of the trade
  - Early stopping
  - Input data normalization
  - Batchnorm
  - Dropout
  - Data augmentation
- Working with pretrained (foundation) models
  - Feature extraction followed by a RF or fcNN or ....
  - Finetuning, transfer learning



# Literature

- Course website
  - [https://tensorchiefs.github.io/dlwl\\_eth25/](https://tensorchiefs.github.io/dlwl_eth25/)
- Probabilistic Deep Learning: Our probabilistic take
  - <https://www.manning.com/books/probabilistic-deep-learning>
- Deep Learning with Python, Second Edition
  - <https://www.manning.com/books/deep-learning-with-python-second-edition>
- Keras Documentation:
  - <https://keras.io/>
- Other Courses
  - Convolutional Neural Networks for Visual Recognition <http://cs231n.stanford.edu>

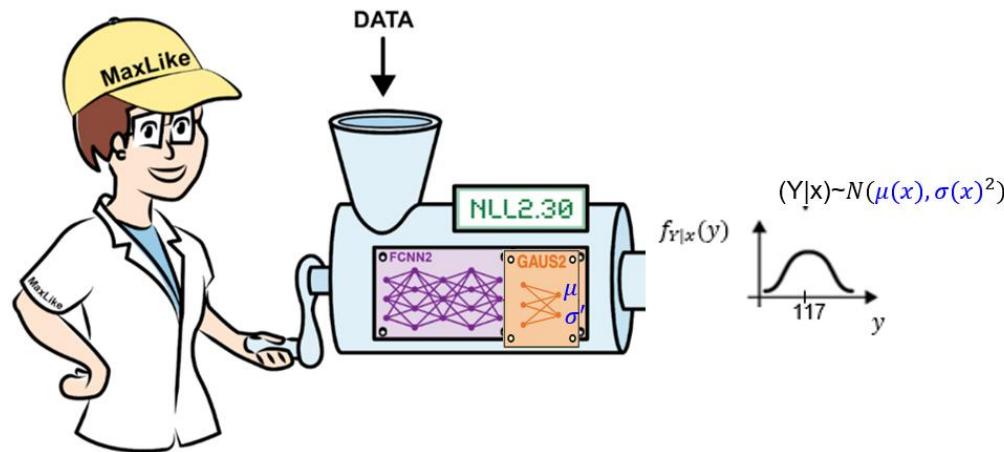
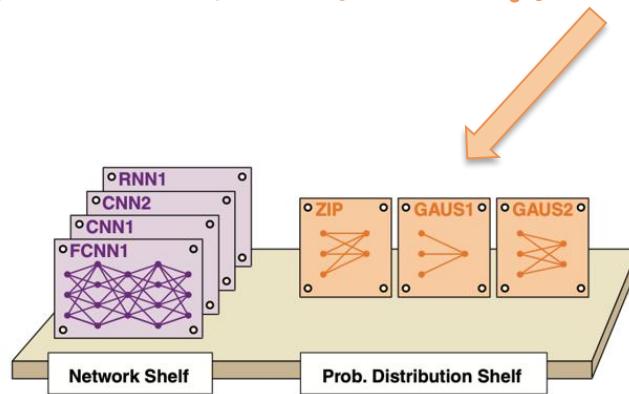


# Topics (may change)

- Half day 1
  - Introduction to probabilistic DL
  - Introduction to Keras with pytorch backend
- Day 2
  - Loss function and optimization
  - Convolutional Neural Networks (CNN) for image data
  - Working with pretrained models
  - Transfer learning
  - Projects
- Day 3
  - Transformer Architectures
  - Model evaluation
  - NN-based logistic regression with tabular and image data
  - Poisson regression with NNs
  - Projects
- Half day 4
  - Project presentations
  - tba

Date
03.02.2025
10.02.2025
17.02.2025
24.02.2025

NLL is used as loss function which depends on outcome distribution

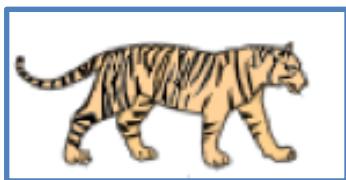


# Training: Tune the weights to minimize the loss

Input  $x^{(i)}$

True class  $y^{(i)}$

predicted class  
(class with highest predicted probability)



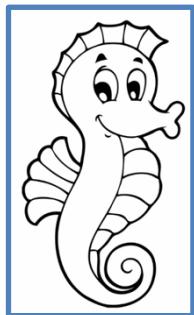
Tiger

Lion



Tiger

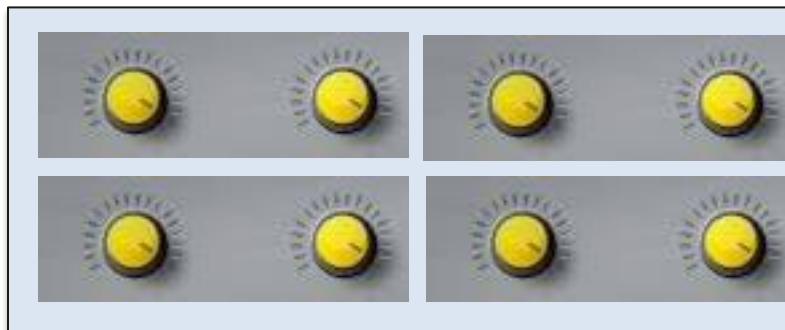
Tiger



Seehorse

Seehorse

Neural network with many weights  $w$



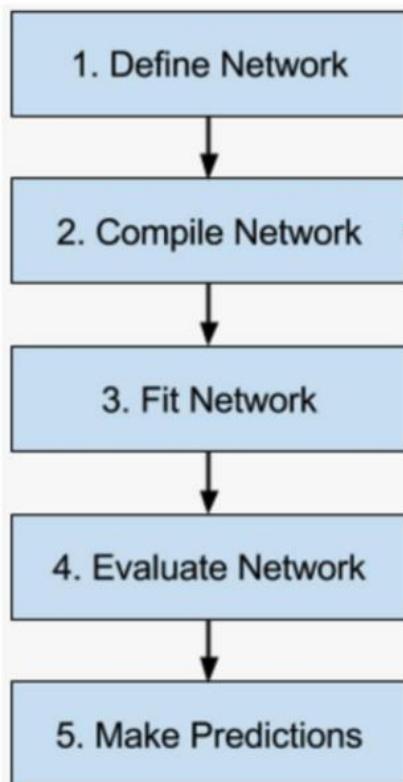
Trainingsprinciple:

Weights are tuned so a loss functions gets minimized.

$$\hat{w} = \underset{w}{\operatorname{argmin}} \text{ loss}(\{y^{(i)}, x^{(i)}\}, w)$$

...

# Recall: in Keras we choose a loss function



```
model.compile(optimizer=SGD(learning_rate=0.01),  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

Loss Function to optimize  
Here: Binary CE

# Losses in Keras

The purpose of loss functions is to compute the quantity that a model should seek to minimize during training.

## Available losses

Note that all losses are available both via a class handle and via a function handle. The class handles enable you to pass configuration arguments to the constructor (e.g. `loss_fn = CategoricalCrossentropy(from_logits=True)`), and they perform reduction by default when used in a standalone way (see details below).

▶ [Keras 3 API documentation / Losses / Probabilistic losses](#)

### Probabilistic losses

- `BinaryCrossentropy class`
- `BinaryFocalCrossentropy class`
- `CategoricalCrossentropy class`
- `CategoricalFocalCrossentropy class`
- `SparseCategoricalCrossentropy class`
- `Poisson class`
- `CTC class`
- `KLDivergence class`
- `binary_crossentropy function`
- `categorical_crossentropy function`
- `sparse_categorical_crossentropy function`
- `poisson function`
- `ctc function`
- `kl_divergence function`

### Regression losses

- `MeanSquaredError class`
- `MeanAbsoluteError class`
- `MeanAbsolutePercentageError class`
- `MeanSquaredLogarithmicError class`
- `CosineSimilarity class`

▶ [Keras 3 API documentation / Losses / Regression losses](#)

## Probabilistic losses

### BinaryCrossentropy class

[so]

```
keras.losses.BinaryCrossentropy(  
    from_logits=False,  
    label_smoothing=0.0,  
    axis=-1,  
    reduction="sum_over_batch_size",  
    name="binary_crossentropy",  
    dtype=None,  
)
```

## Regression losses

### MeanSquaredError class

[source]

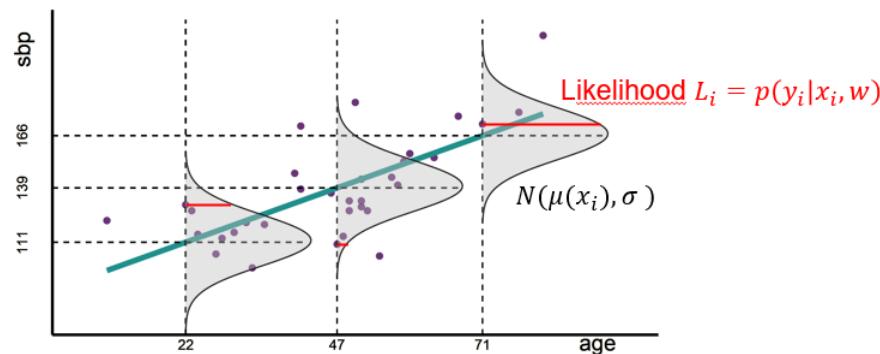
```
keras.losses.MeanSquaredError(  
    reduction="sum_over_batch_size", name="mean_squared_error", dtype=None  
)
```

# Which loss to use in probabilistic DL?

## Negative Log Likelihood!



# Recap: Maximum Likelihood (one of the most beautiful ideas in statistics)



Ronald Fisher in 1913  
Also used before by  
Gauss, Laplace

Tune the parameters weights of the network such, that observed data (training data) is most likely under the predicted outcome distribution.

We assume iid training data  $\rightarrow$  multiplication of likelihood over all data points:

$$\hat{w} = \underset{w}{\operatorname{argmax}} \prod_{i=1}^N L_i = \underset{w}{\operatorname{argmax}} \prod_{i=1}^N p(y_i|x_i, w)$$

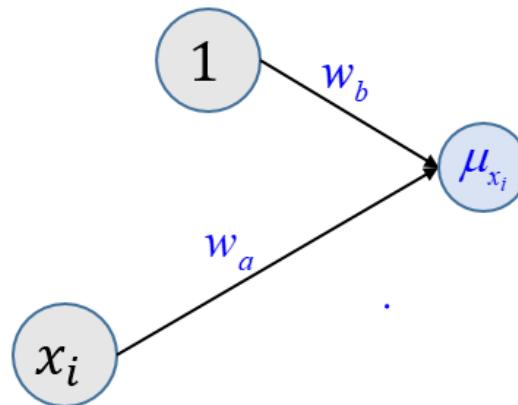
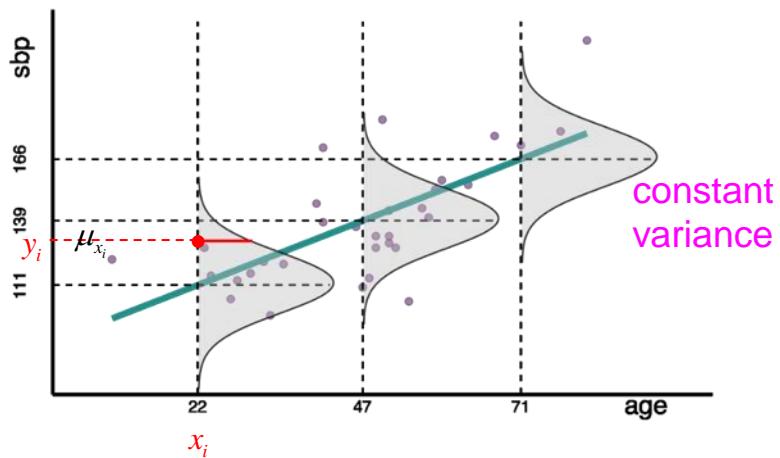
Practically: Use Negative Log-Likelihood NLL as loss and minimize NLL

take log; minimize after multiplication with -1, note that NLL in DL is usually the mean-negative log-likelihodd

$$\hat{w} = \underset{w}{\operatorname{argmin}} \text{NLL}(y^{(i)}, x^{(i)}, w) = \underset{w}{\operatorname{argmin}} \frac{1}{N} \sum_{i=1}^N -\log(p(y_i|x_i, w))$$

# Fit “linear regression NN” via Maximum likelihood principle

$$Y_{X_i} \sim N(\mu_{x_i} = w_a \cdot x_i + w_b, \sigma^2)$$



ML-principle:

$$\begin{aligned} \mathbf{w}_{\text{ML}} &= \underset{w}{\operatorname{argmax}} \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(y_i - \mu_{x_i})^2}{2\sigma^2}} \\ &= \underset{w}{\operatorname{argmin}} \sum_{i=1}^n -\log \left( \frac{1}{\sqrt{2\pi\sigma^2}} \right) + \frac{(y_i - \mu_{x_i})^2}{2\sigma^2} \end{aligned}$$

Negative Log-Likelihood (NLL)

$$(\hat{w}_a, \hat{w}_b)_{\text{ML}} = \underset{w_a, w_b}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n (y_i - (w_a \cdot x_i + w_b))^2$$

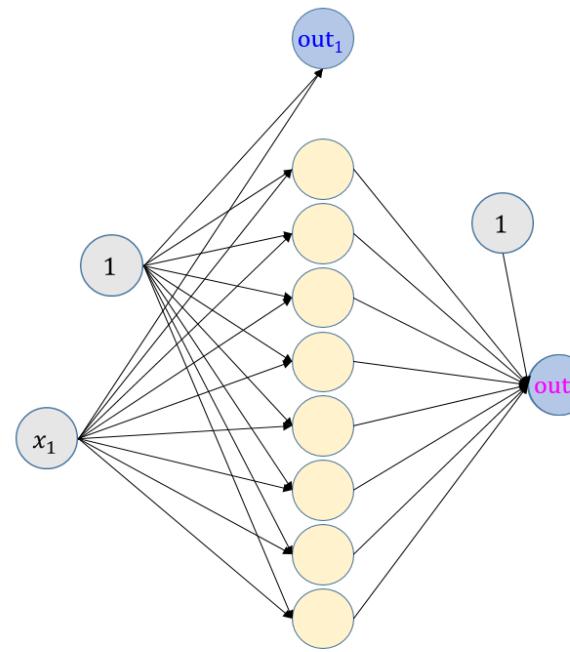
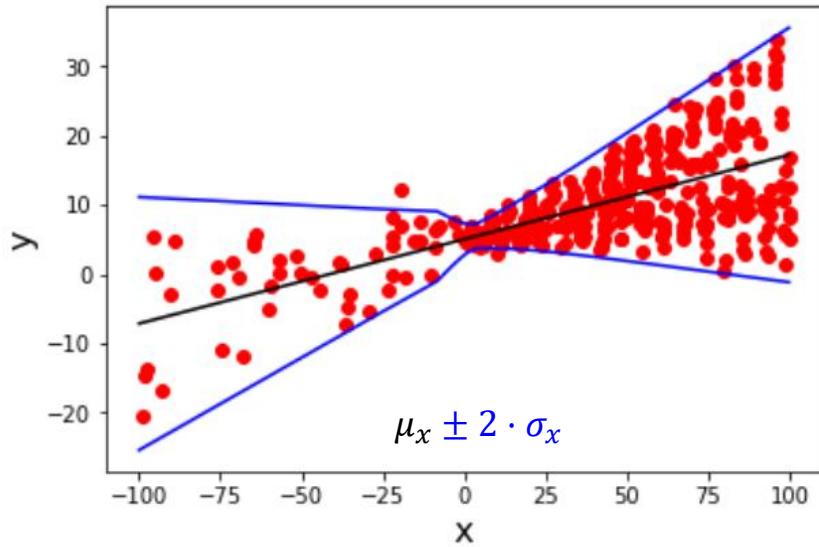
Minimize MSE loss

$\hat{w}_a$        $\hat{w}_b$

Minimizing NLL loss  $\stackrel{\sigma \text{ const}}{=} \text{Minimizing MSE loss}$

# Linear regression with flexible non-constant variance

$$Y_{X_i} = (Y|X_i) \sim N(\mu_{x_i}, \sigma_x^2)$$



Minimize the mean negative log-likelihood (NLL) on train data:

$$NLL(w) = \frac{1}{n} \sum_{train-data} -\log(p(y_i|x_i, w))$$

gradient descent with NLL loss

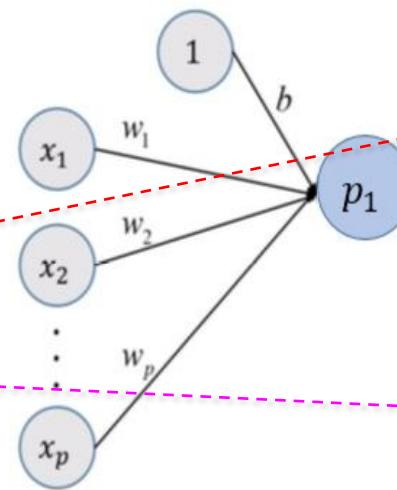
$$\hat{\mathbf{w}}_{ML} = \operatorname{argmin}_w \sum_{i=1}^n -\log\left(\frac{1}{\sqrt{2\pi\sigma_{x_i}^2}}\right) + \frac{(y_i - \mu_{x_i})^2}{2\sigma_{x_i}^2}$$

Note: we do not need to know the “ground truth for  $\sigma$ ” – the likelihood does the job!

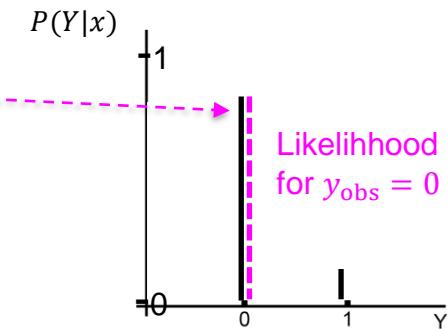
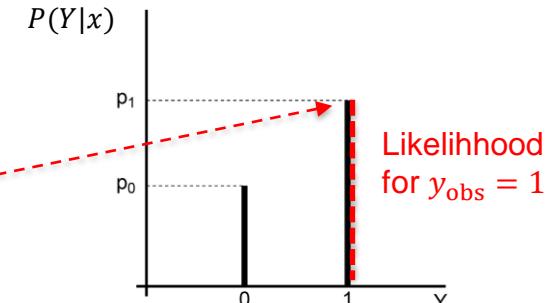
# Maximum likelihood principle in logistic regression

Training data set

<b>x1</b>	...	<b>xp</b>	<b>y</b>
0.4	...	-1.3	0
1.1	...	0.2	1
:	:	:	:
0.6	...	-0.9	0



probabilistic prediction



Given (=conditioned on) the input features  $x$  of an observation  $i$ , a well trained NN

- should predict large  $p_1 = P(Y = 1|x)$  if the observed class is  $y_i = 1$
  - should predict small  $p_1$  hence large  $p_0 = P(Y = 0|x)$  if observed is  $y_i = 0$
- The likelihood (for the observed outcome) or LogLikelihood should be large

$$\text{LogLikelihood} = \sum_{i=1}^n [y_i \log(p_{1i}) + (1 - y_i) \log(1 - p_{1i})]$$

Notation trick in statistics – it selects correct log-probability since  $y_i \in \{0,1\}$

# Fitting a “logistic regression NN”

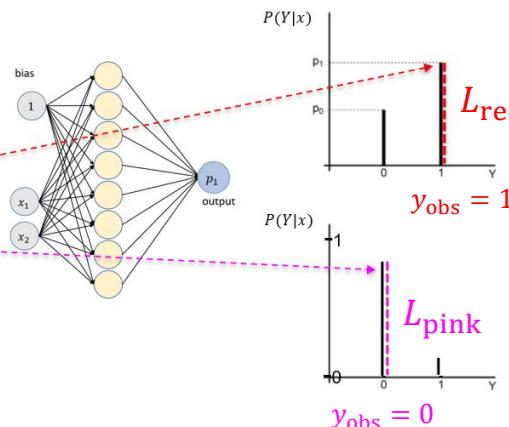
The Likelihood of an observation  $i$  is the likelihood (probability), that the predicted probability distribution  $P(Y|x_i)$  assigns to the observed outcome  $y_i$ .

Note: the predicted  $P(Y|x_i)$  and the corresponding likelihood for the observed  $y_i$  depends on the data-point  $(x_i, y_i)$  and the model parameter values

→ The higher the likelihood, the better is the model prediction  $P(Y|x)$

Training data set

x1	...	xp	y
0.4	...	-1.3	0
1.1	...	0.2	1
:	:	:	:
0.6	...	-0.9	0



$L_{\text{red}}$   $L_{\text{red}}$  is likelihood of red observation with  $y_{\text{obs}} = 1$

$L_{\text{pink}}$   $L_{\text{pink}}$  is likelihood of pink observation with  $y_{\text{obs}} = 0$

## Maximum likelihood principle:

Statistical models are fit to maximize the average LogLikelihood

$$L = \frac{1}{N} \sum L_i = \frac{1}{N} \sum [y_i \log(p_{1i}) + (1 - y_i) \log(1 - p_{1i})]$$

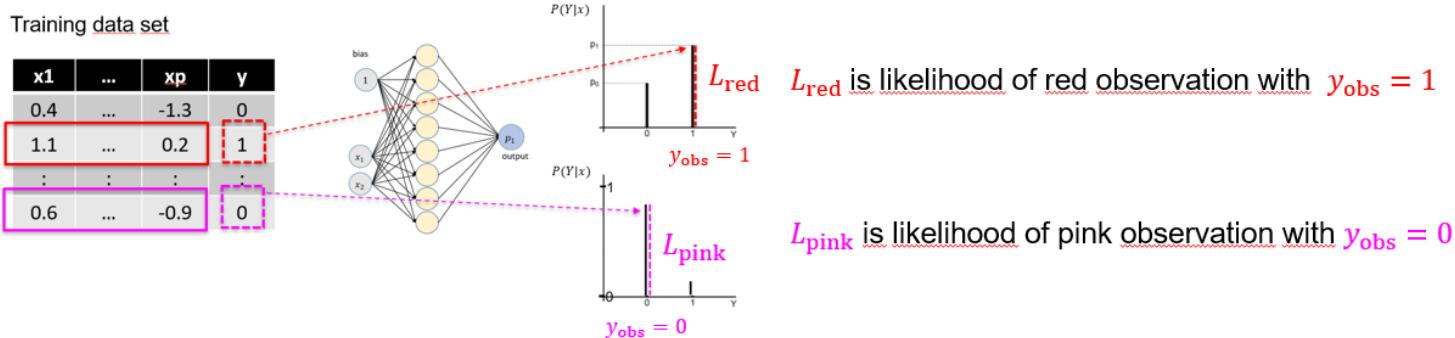
we often use the simplified notation average logLik :  $L = \frac{1}{N} \sum \log(p_i)$

Predicted probability for  
the observed outcome  $y_{\text{obs}}$

# NLL Loss for probabilistic binary classification

In DL we aim to minimize a loss function  $L(\text{data}, w)$  which depends on the weights  
→ Instead of maximizing the average LogLikelihood  
we minimize the averaged Negative LogLikelihood (NLL):

$$\text{loss} = \text{NLL} = -\frac{1}{N} \sum \log(L_i) = -\frac{1}{N} \sum \log(p_i)$$



The best possible value of the NLL contribution of an observation  $i$  is  $-\log(L_i) = -\log(1) = 0$

The worst possible value of the NLL contribution of an observation  $i$   $-\log(L_i) = -\log(0) = \infty$

Note: In Keras the NLL for binary outcome is called '`binary_crossentropy`',  
if we do probabilistic binary classification with **1 output node with sigmoid activation**.  
If we would use **2 output nodes and softmax** than the NLL is called '`categorical_crossentropy`'

# Define the NLL loss for a probabilistic model

- Some keras losses correspond to NLL, e.g.
  - Linear Regression with constant variance
    - Mean Squared Error (keras.losses.mse)
  - Multiclass classification
    - Categorical\_crossentropy (keras.losses.categorical\_crossentropy)
- For a given outcome distribution family, you can build the NLL loss
  1. Define a model with as many outputs as the number of parameters of the outcome distribution family
  2. Link the outputs of the NN to the parameters of the outcome distribution
  3. Use negative log likelihood as loss function

## 1. Define a model with as many outputs as the number of parameters of the distribution



```
# Define the model
inputs = Input(shape=(1,))
hidden = layers.Dense(10, activation="relu")(inputs)
# ... more hidden layers ...
outputs = layers.Dense(2)(hidden) # <--- Outputs mean and log(sd)
model = Model(inputs=inputs, outputs=outputs)
```

[8] ✓ 0.0s

Python

## 2. Link the outputs of the model to the parameters of the distribution

```
# Wrapper function to convert model output to a PyTorch Normal distribution
@staticmethod
def output_to_gaussian_distribution(output):
    mean = output[:, 0:1] # Gets the first column while keeping dimensions (like output[, 1, drop=FALSE] in R)
    log_sd = output[:, 1:2] # Gets the second column while keeping dimensions
    scale = torch.exp(log_sd) # Inverse link function to ensure positive scale
    return Normal(loc=mean, scale=scale)
```

[7] ✓ 0.0s

Python

## 3. Define NLL

```
# Custom Negative Log-Likelihood Loss
def negative_log_likelihood(y_true, output):
    dist = output_to_gaussian_distribution(output)
    return -dist.log_prob(torch.tensor(y_true)).mean()

# Compile the model
model.compile(optimizer="adam", loss=negative_log_likelihood)
```

[10] ✓ 0.0s

Python

# Exercise Define a custom loss function



Notebook: 02

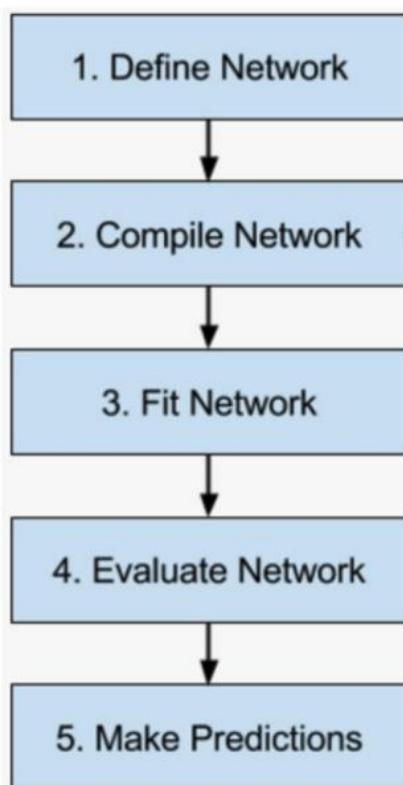
[https://github.com/tensorchiefs/dlbl\\_eth25/blob/master/notebooks/02\\_custom\\_loss.ipynb](https://github.com/tensorchiefs/dlbl_eth25/blob/master/notebooks/02_custom_loss.ipynb)

Fit a NN via  
**Stochastic Gradient Descent**

# Recall: in Keras we choose a loss function



Which optimizer should be used?  
Here Stochastic Gradient Descent

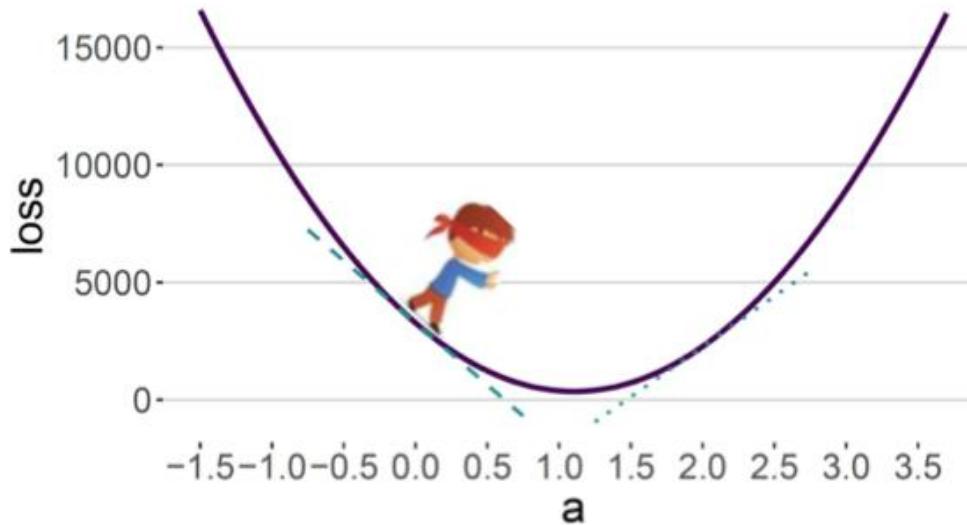


```
model.compile(optimizer=SGD(learning_rate=0.01),  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

Loss Function to optimize  
Here: Binary CE

# Idea of gradient descent

- Shown loss function for a single parameter  $a$

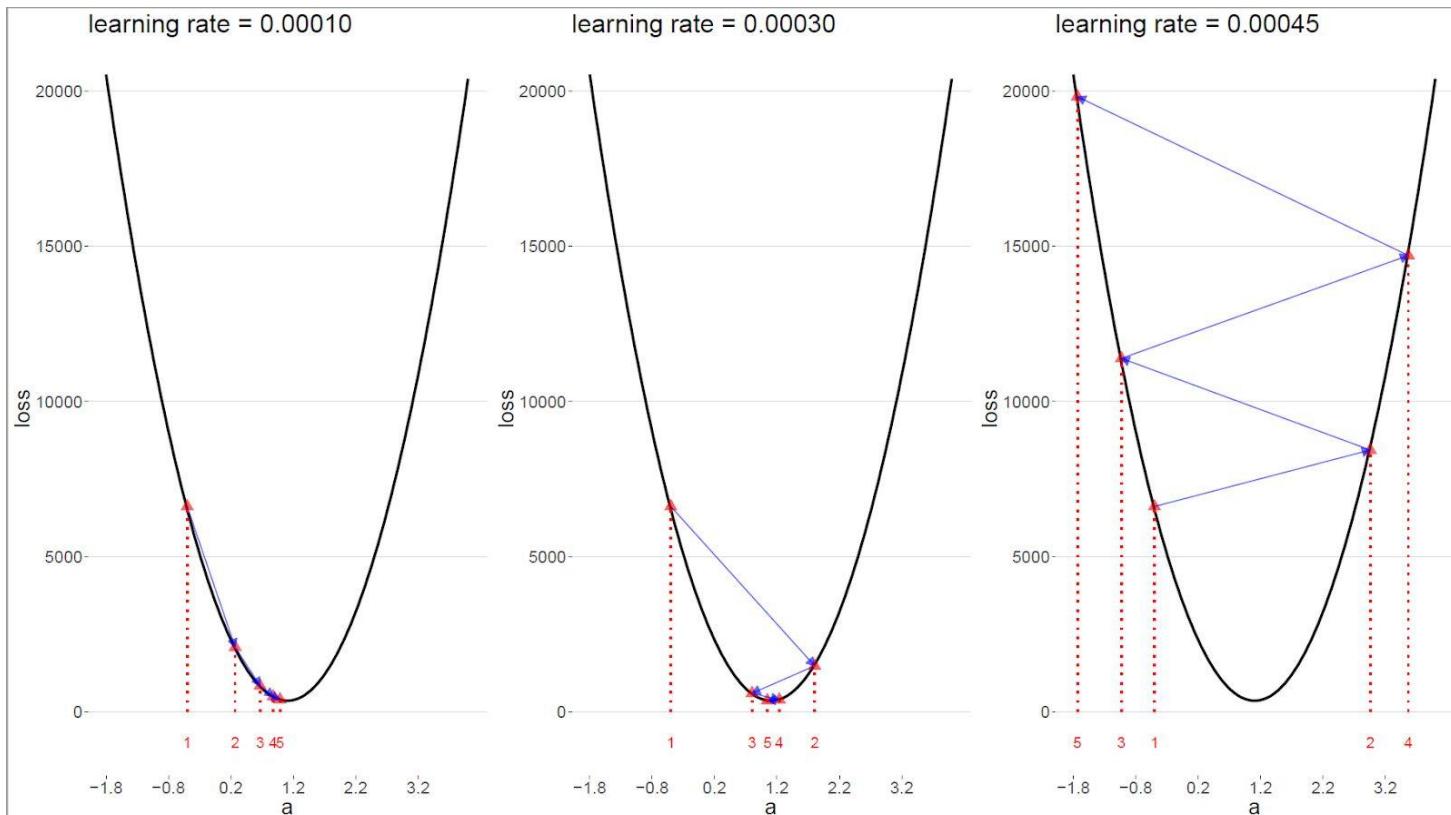


- Take a large step if slope is steep (you are away from minimum)
- Slope of loss function is given by gradient of the loss w.r.t.  $a$
- Iterative update of the parameter  $a$

$$a^{(t)} = a^{(t-1)} - \epsilon \frac{\partial L(a)}{\partial a} \Big|_{a=a^{(t-1)}}$$

learning rate

# The learning rate is a very important parameter for DL



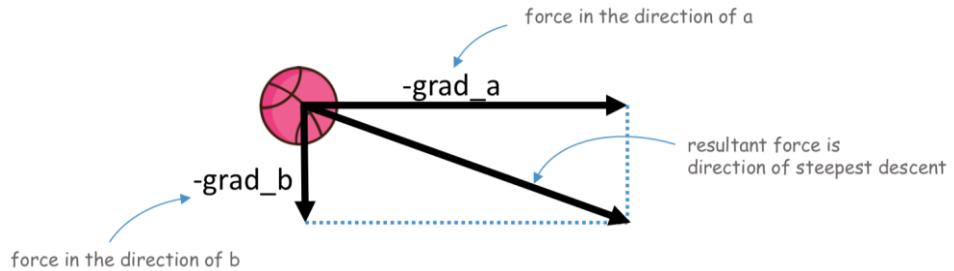
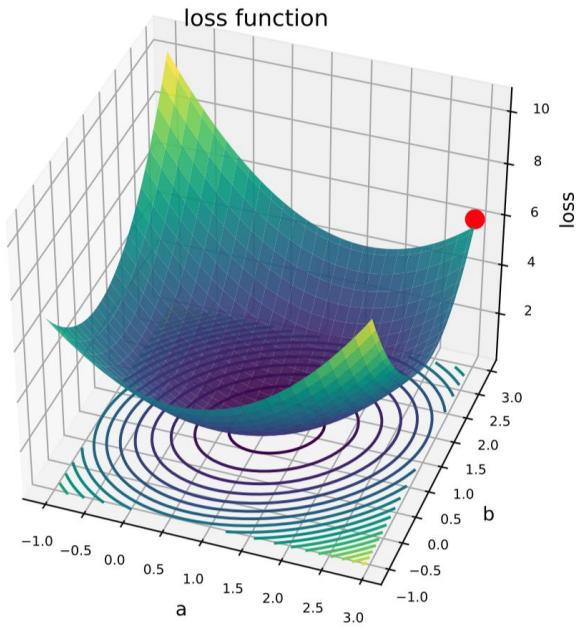
good – little slow

good

learning rate too high

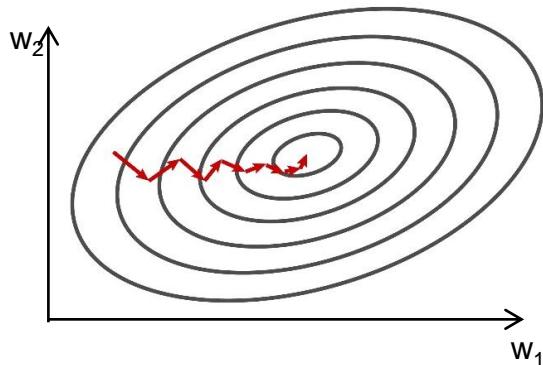
If the loss diverges to infinity: Don't panic, lower the learning rate!

# In two dimensions



Gradient is perpendicular to contour lines

$$w_i^{(t)} = w_i^{(t-1)} - \epsilon \frac{\partial L(\mathbf{w})}{\partial w_i} \Big|_{w_i=w_i^{(t-1)}}$$



Remark:

DL frameworks can do the differentiation see [notebook](#).

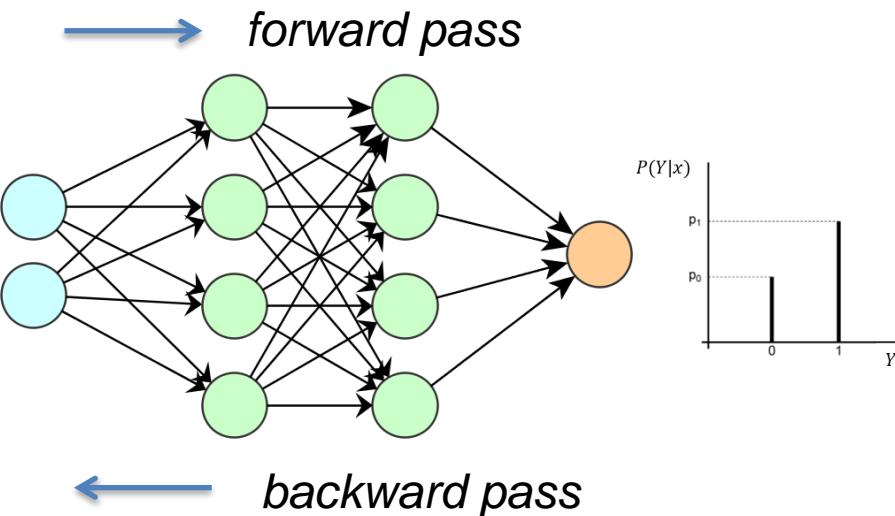
In DL the Loss is computed as mean NLL over all examples in a mini-batch, then weight update is performed

→ Since minibatches hold random part of all data, we talk about Stochastic Gradient Descent (SGD).

→ 'SG'D is the basic of most optimizers, but usually 'adam', an improved updating procedure is used in DL. See <https://keras.io/api/optimizers/> for an overview.

# Backpropagation

- We efficiently train the weights in a NN via forward and backward pass
  - Forward Pass propagate training example through network
    - Predicts as output  $P(Y|x_i)$  for each input  $x_i$  in the batch given the NN weights  $w$   
→ With  $P(Y|x_i)$  and the observed  $y_i$  we compute the loss  $L = \left( \text{NLL} = -\frac{1}{N} \sum \log(L_i) \right)$
  - Backward pass propagate gradients through network
    - Via chain rule all gradients  $\frac{\partial L(\mathbf{w})}{\partial w_k}$  are determined  
→ update the weights  $w_i^{(t)} = w_i^{(t-1)} - \epsilon \frac{\partial L(\mathbf{w})}{\partial w_i} \Big|_{w_i=w_i^{(t-1)}}$



# Typical Training Curve: ReLU is used for fast training

Loss curves for NN with ReLu or sigmoid

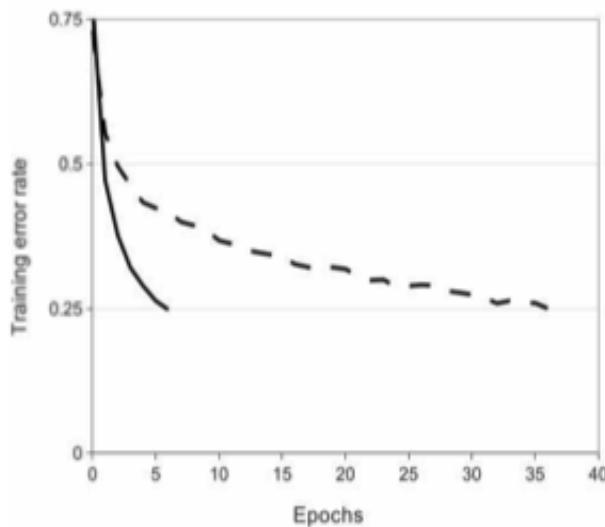
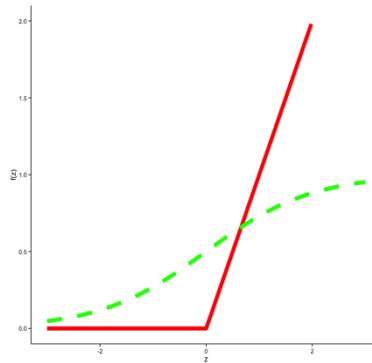
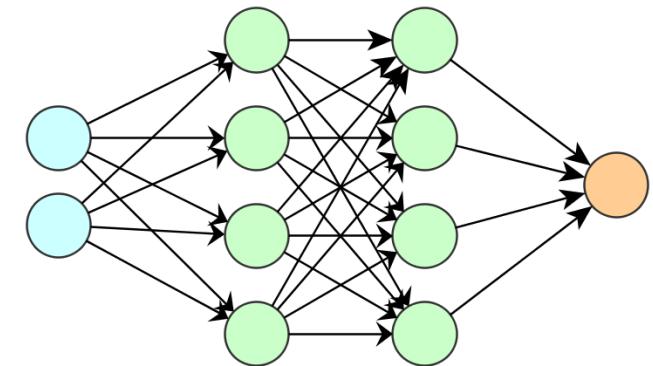
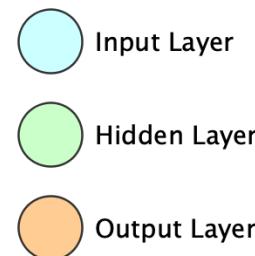


Figure 1: A four-layer convolutional neural network with ReLUs (solid line) reaches a 25% training error rate on CIFAR-10 six times faster than an equivalent network with tanh neurons

Source:  
Alexnet  
Krizhevsky et al 2012



Activation functions  
Green: sigmoid.  
Red: ReLU

Remark:  
In deep NNs ReLU is preferred since training is faster, in very shallow NNs sigmoids can be better to achieve smoothness.

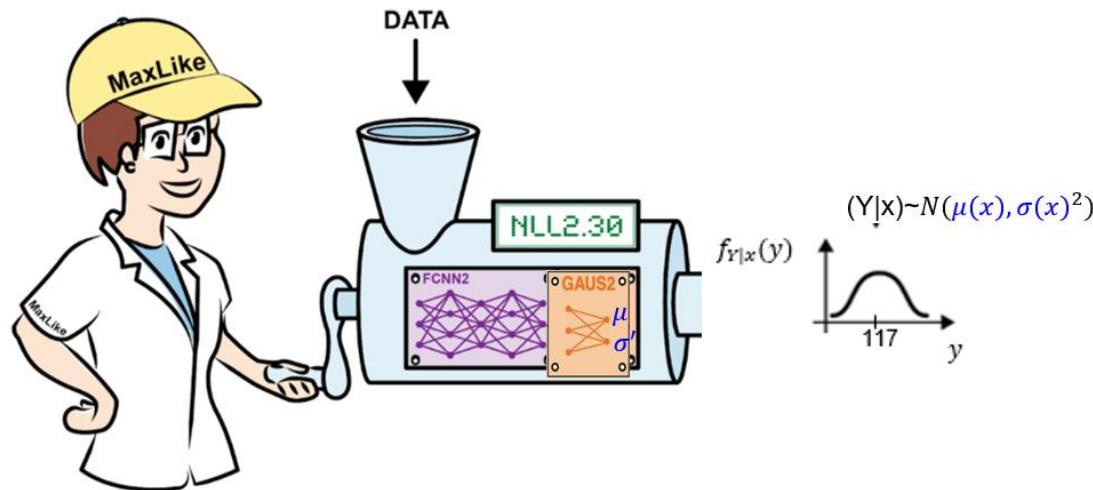
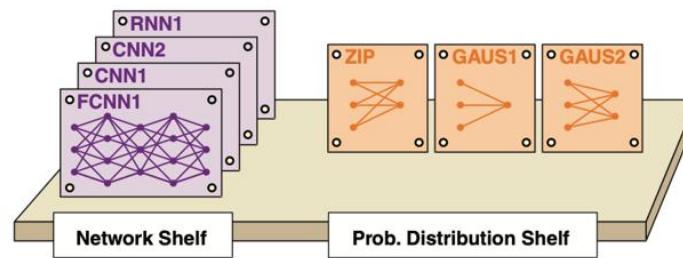
Epochs: "each training examples is used once"

# The miracle of gradient descent in DL

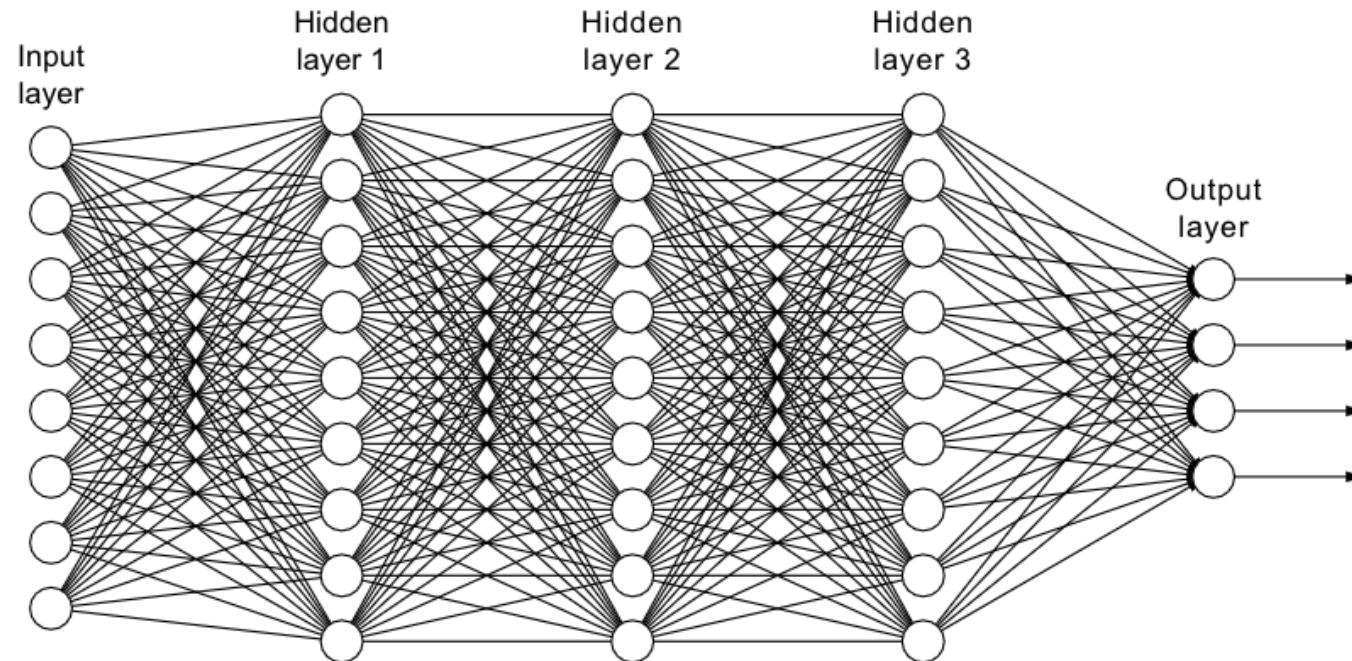


Loss surface in DL (is not convex) but SGD magically also works for non-convex problems.

# NN architectures

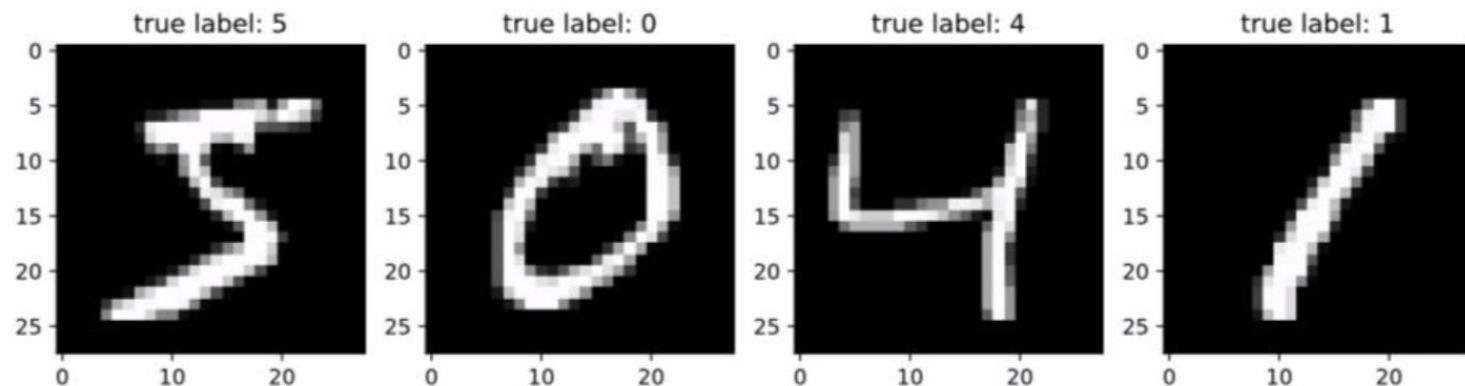


# A fully connected NN

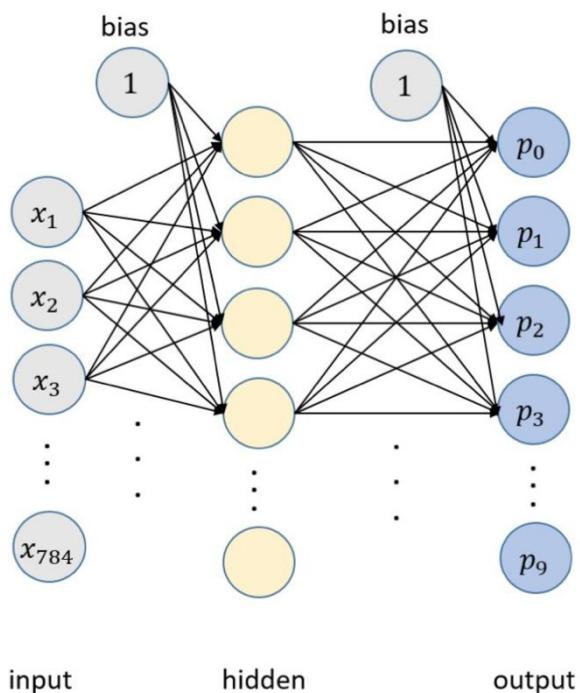


# Probabilistic classification

- Usually in DL the model predicts a probability for each possible class
- Example:
  - Banknote from exercise – classes: “real” or “fake”
  - Typical example Number from hand-written digit – classes: 0, 1, ..., 9



# Classification: Softmax Activation



$p_0, p_1 \dots p_9$  are probabilities for the classes 0 to 9.

Activation of last layer  $z_i$  incoming

$$p_i = \frac{e^{z_i}}{\sum_{j=0}^9 e^{z_j}}$$

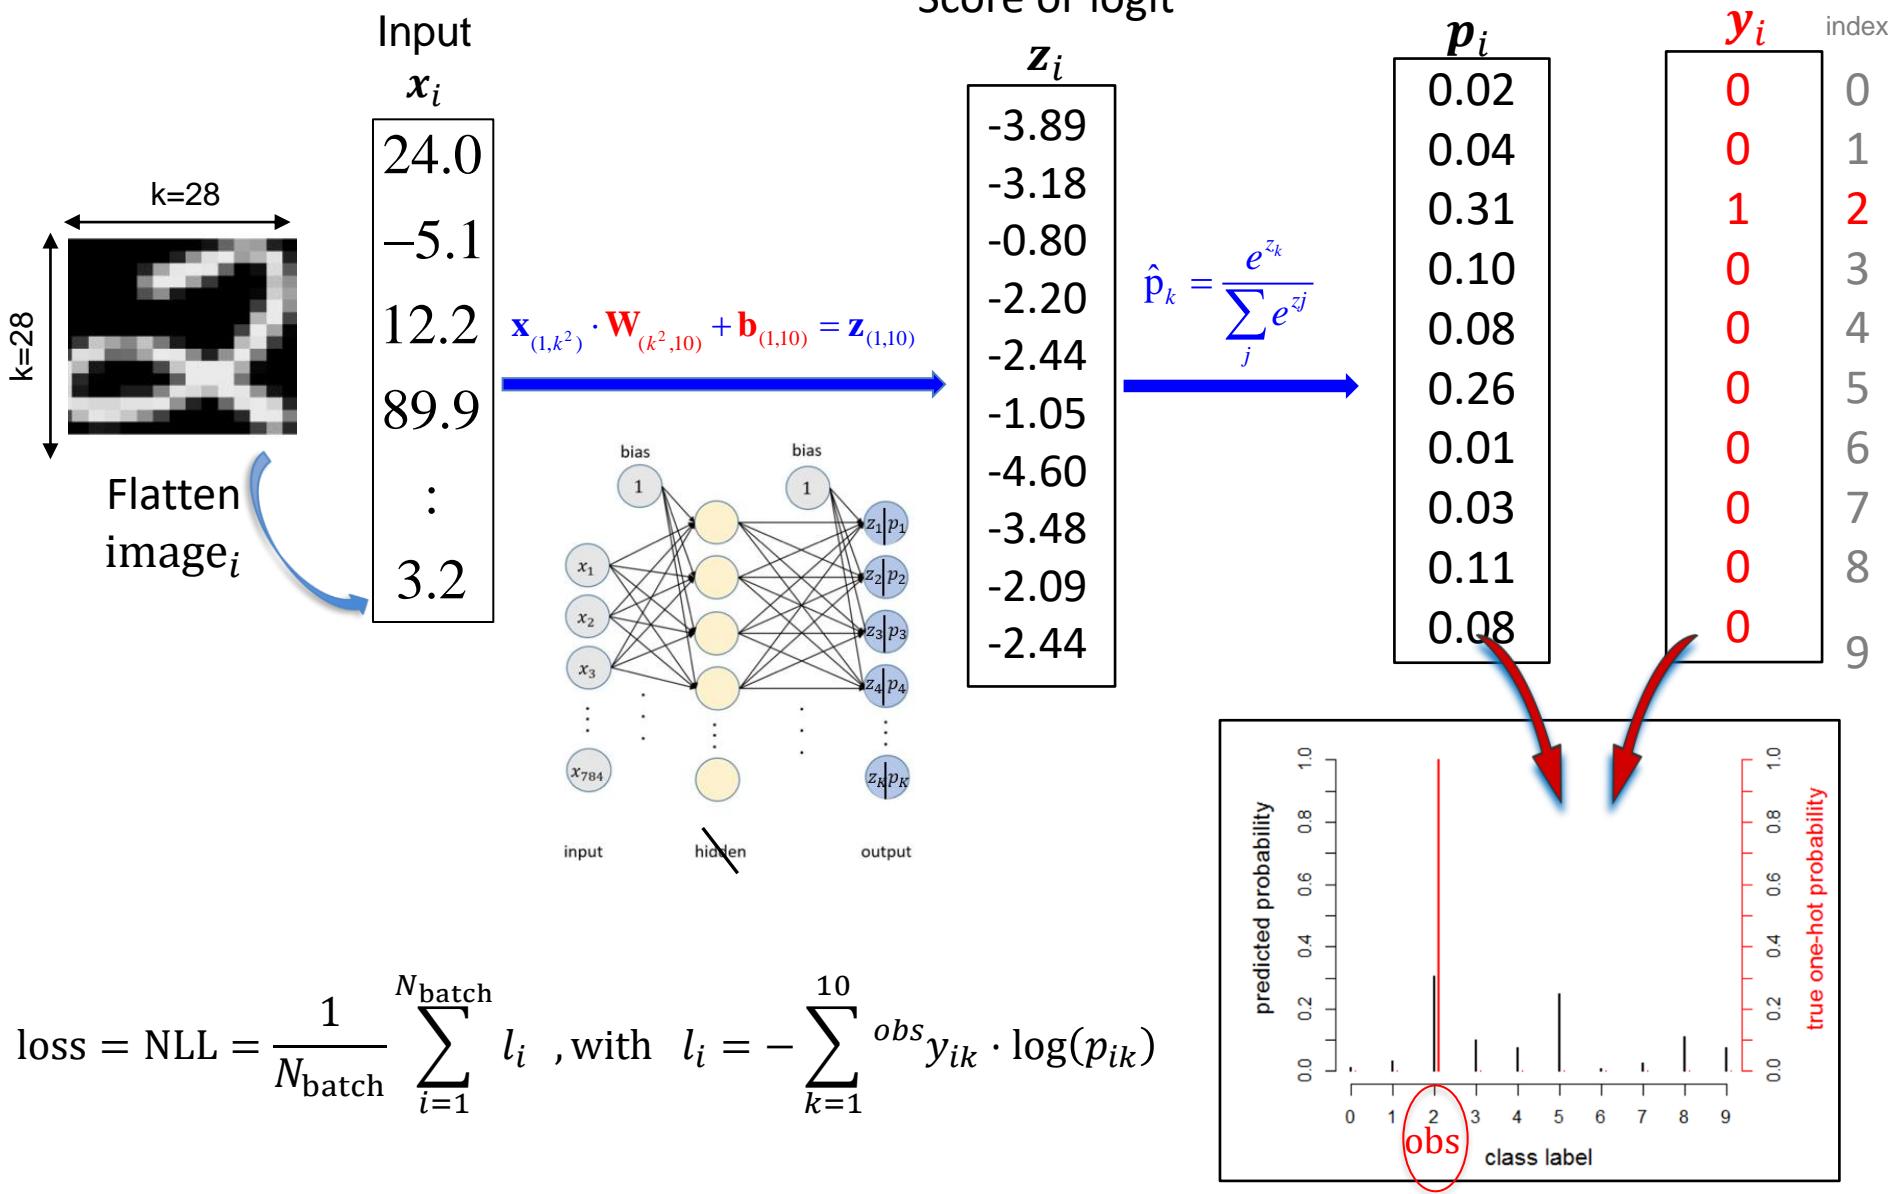
Makes outcome positive

Ensures that  $p_i$ 's sum up to one

This activation is called softmax

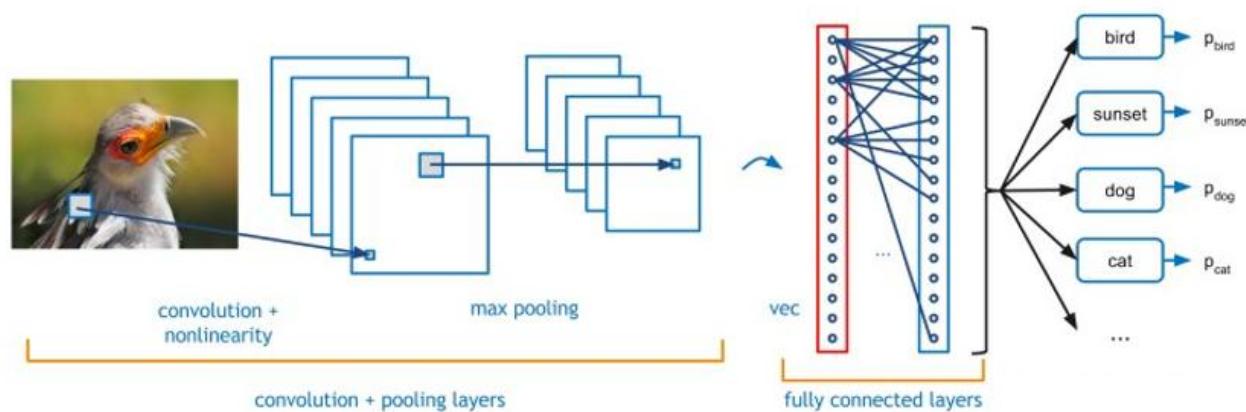
Figure 2.12: A fully connected NN with 2 hidden layers. For the MNIST example, the input layer has 784 values for the 28 x 28 pixels and the output layer out of 10 nodes for the 10 classes.

# NLL for multi-class = cross-entropy



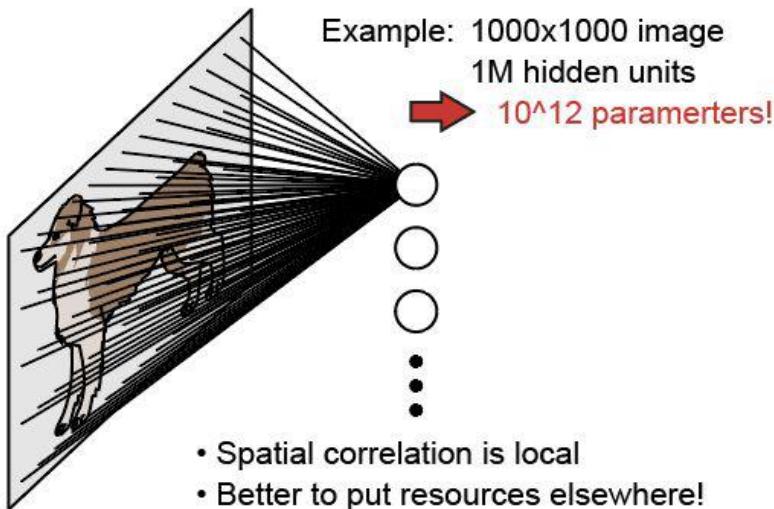
Loss = NLL = cross-entropy ( $-\sum p_i^{\text{obs}} \log p_i^{\text{pred}}$ ) averaged over all images in mini-batch

# Convolutional Neural Networks for image data



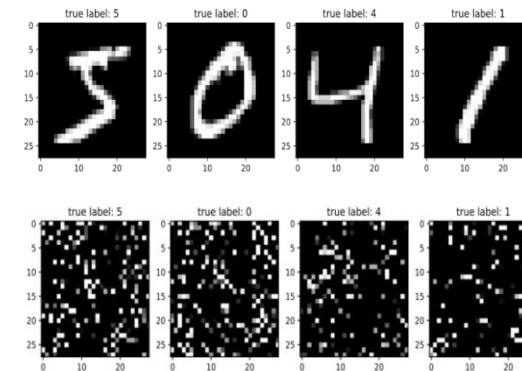
# Is a fully connected NN a good architecture for images?

## Fully connected neural net



$$z = b + \sum_i x_i w_i$$

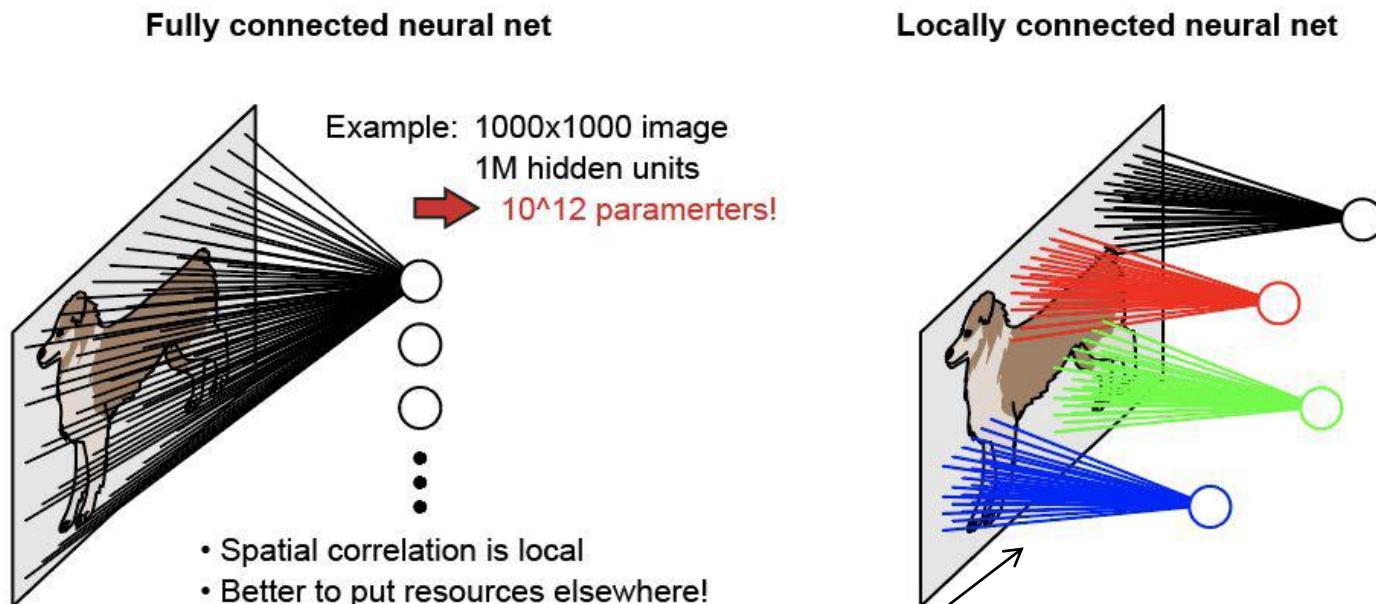
A fcNN does not take into account the spacial relationships of the pixels.



If you reshuffle the pixels of in all images in the same manner, the images look weird but the fcNN trains as with original images

[04\\_fcnn\\_mnist\\_shuffled.ipynb](#)

# Convolution extracts local information using few weights



## Shared weights:

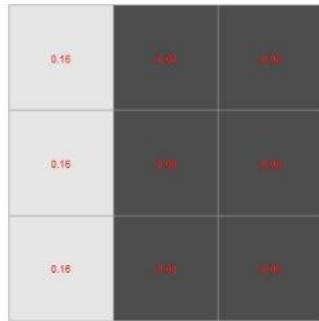
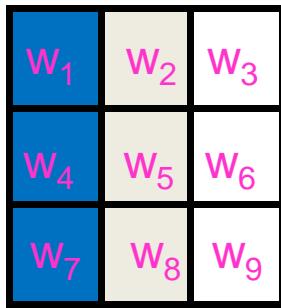
by using the **same weights** for each patch of the image we need much **less parameters** than in the fully connected NN and get from each patch the same kind of **local feature information** such as the presence of a edge.

# Convolutional networks use neighborhood information and replicated local feature extraction

In a locally connected network the calculation rule

$$z = b + \sum_i x_i w_i$$

Pixel values in a small image patch are element-wise multiplied with weights of a small filter/kernel:

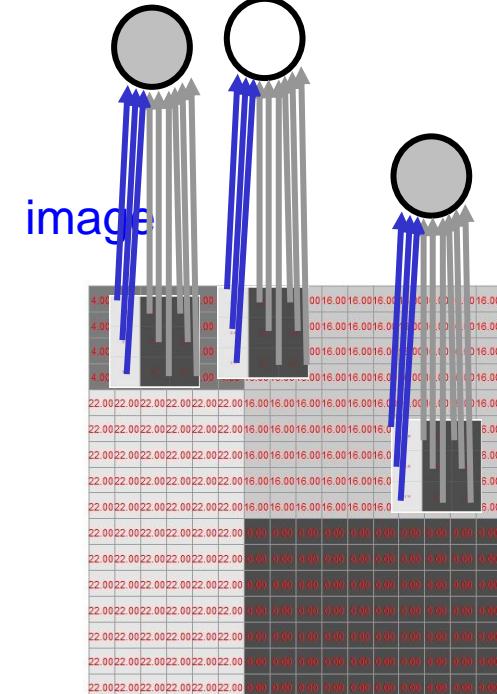


The filter is applied at each position of the image and it can be shown that the **result is maximal if the image pattern corresponds to the weight pattern.**

The results form again an image called **feature map** (=activation map) which shows at which position the feature is present.

feature/activation map

-0.00	0.00	0.00	0.00	-0.00	5.82	2.91	0.00	-0.00	-0.00	0.00	0.00	-0.00	0.00	-0.00	-0.00	-0.00	-0.00	-0.00
-0.00	0.00	0.00	0.00	-0.00	5.82	2.91	-0.00	0.00	-0.00	-0.00	0.00	0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00
0.00	0.00	0.00	0.00	-0.00	5.82	2.91	-0.00	0.00	0.00	0.00	0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00
-0.00	0.00	0.00	0.00	-0.00	-0.00	2.91	1.45	-0.00	0.00	0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00
0.00	0.00	0.00	0.00	-0.00	-0.00	0.00	-0.00	0.00	0.00	0.00	0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00
-0.00	0.00	0.00	0.00	-0.00	-0.00	-2.91	-1.45	-0.00	0.00	0.00	0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00
-0.00	0.00	0.00	0.00	-0.00	-0.00	-2.91	-1.45	-0.00	0.00	0.00	0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00
0.00	0.00	0.00	0.00	0.00	0.00	-2.91	-1.45	-0.00	0.00	0.00	0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00
0.00	0.00	0.00	0.00	0.00	0.00	0.00	-5.49	-2.75	-0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	0.00	0.00	0.00	0.00	0.00	0.00	-0.04	-0.00	-0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	0.00	0.00	0.00	0.00	0.00	0.00	-5.33	-0.00	-0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	0.00	0.00	0.00	0.00	0.00	0.00	-5.33	-0.00	-0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	-5.33	-0.00	-0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	-5.33	-0.00	-0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	-5.33	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00
0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	-5.33	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00	-0.00



## Exercise: Do one convolution step by hand

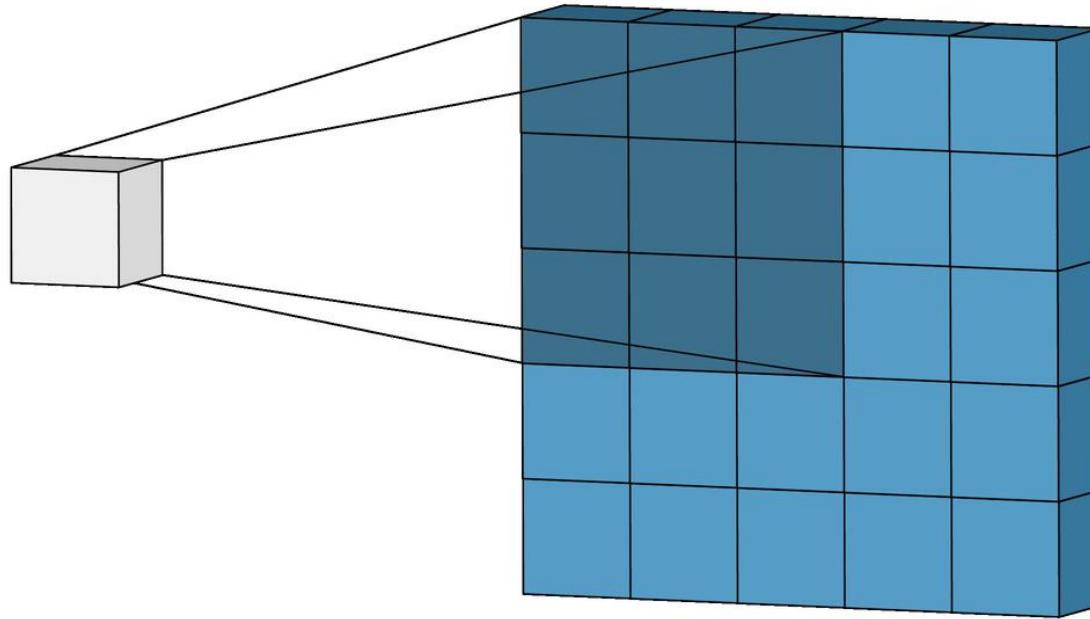


The kernel is 3x3 and is applied at each valid position  
– how large is the resulting activation map?

The small numbers in the shaded region are the kernel weights.  
Determine the position and the value within the resulting activation map.

3	3	2	1	0
0 <sub>0</sub>	0 <sub>1</sub>	1 <sub>2</sub>	3	1
3 <sub>2</sub>	1 <sub>2</sub>	2 <sub>0</sub>	2	3
2 <sub>0</sub>	0 <sub>1</sub>	0 <sub>2</sub>	2	2
2	0	0	0	1

## Applying the same 3x3 kernel at each image position



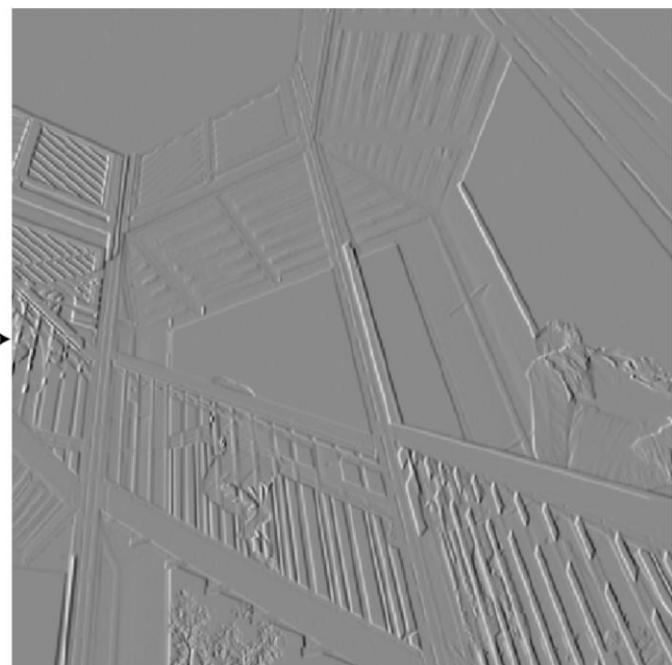
Applying the 3x3 kernel on a certain position of the image yields one pixel within the activation map where the position corresponds to the center of the image patch on which the kernel is applied.

# Example of designed Kernel / Filter



$$\begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix}$$

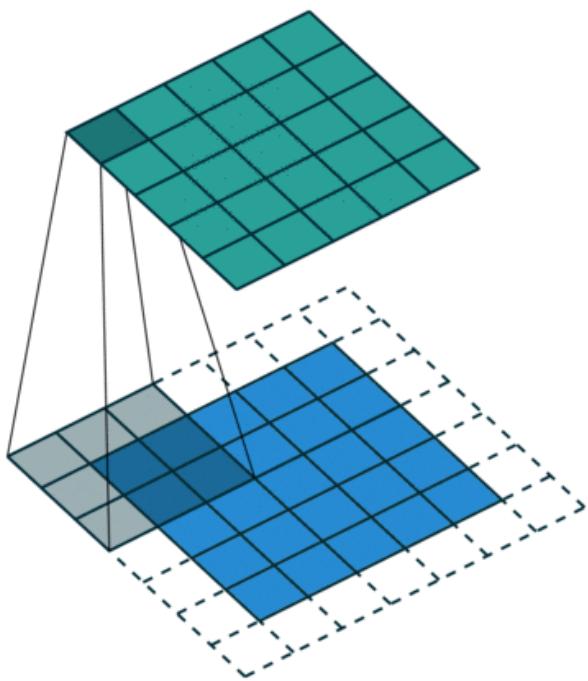
Horizontal Sobel kernel



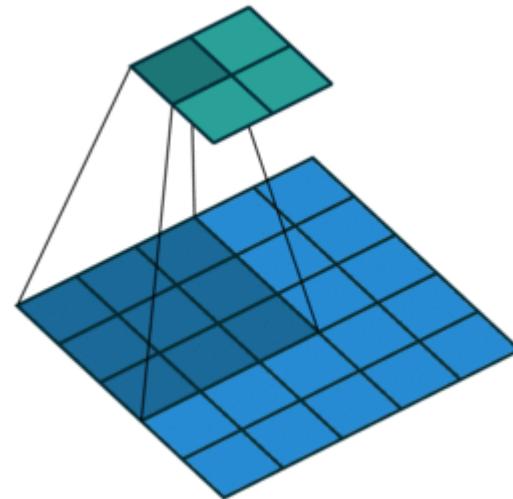
Applying a vertical edge detector kernel



# CNN Ingredient I: Convolution



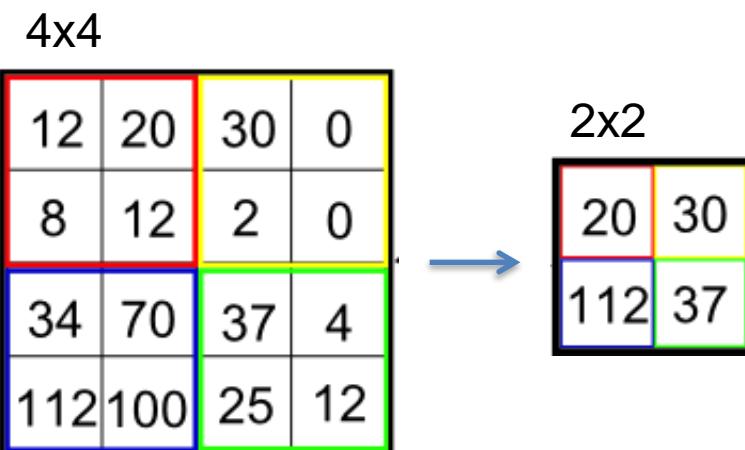
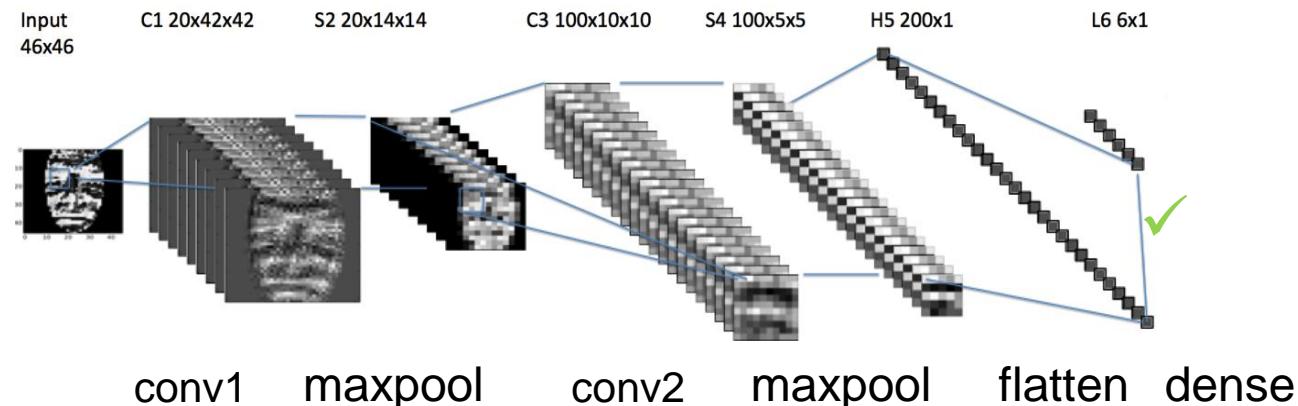
Zero-padding to achieve  
**same** size of feature and input



no padding to only use  
**valid** input information

The **same** weights are used at each position of the input image.

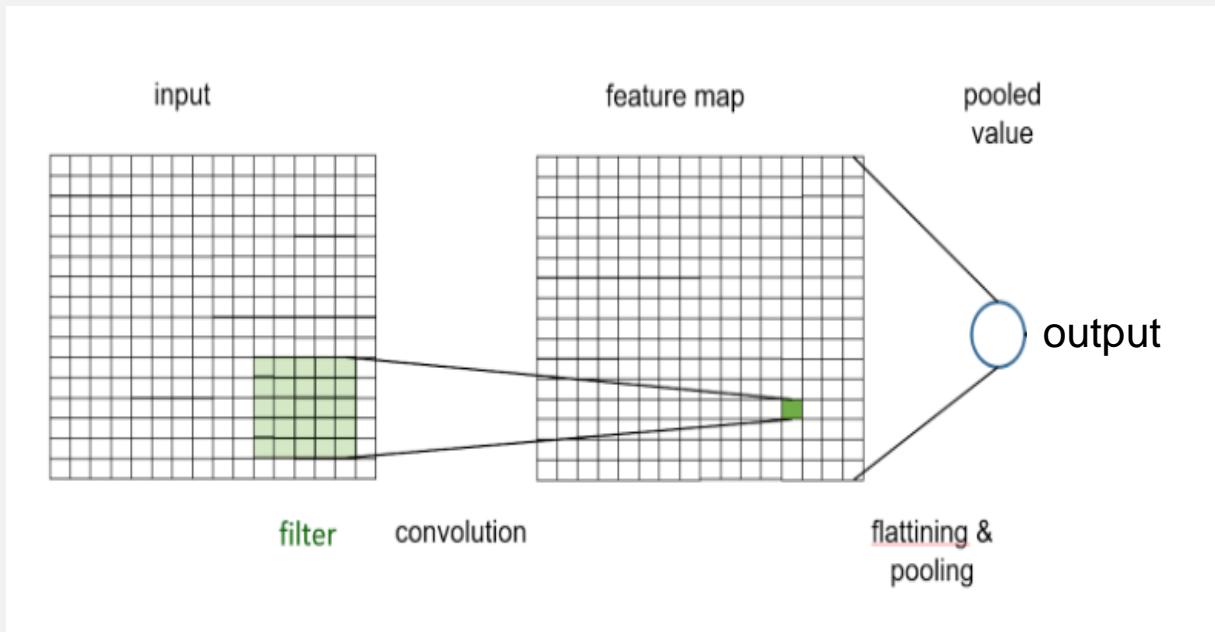
## CNN ingredient II: Maxpooling Building Blocks reduce size



Simply join e.g. 2x2 adjacent pixels in one by taking the max.  
→ less weights in model  
→ Less train data needed  
→ increased performance

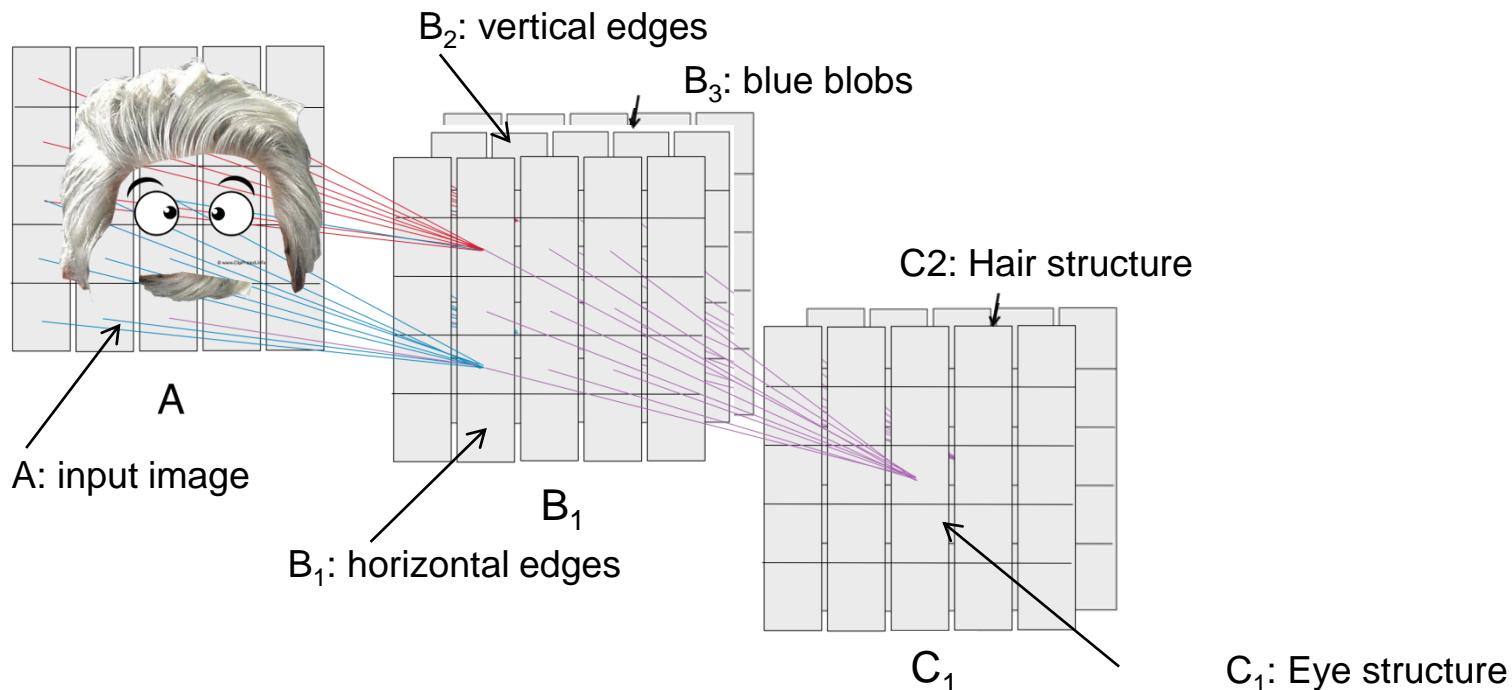
Hinton: „The pooling operation used in convolutional neural networks is a big mistake and the fact that it works so well is a disaster“

# Exercise: Artstyle Lover

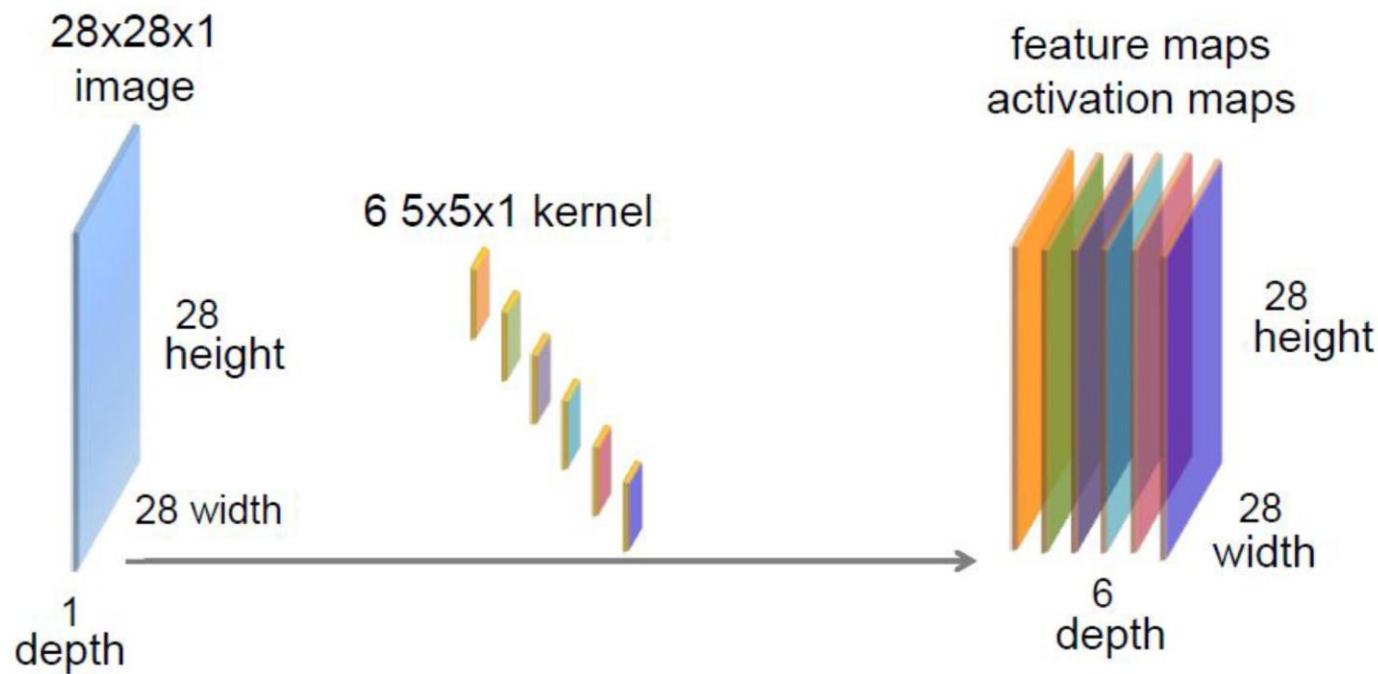


Open NB in: [https://github.com/tensorchiefs/dlwbl\\_eth25/blob/master/notebooks/02\\_cnn\\_edge\\_lover.ipynb](https://github.com/tensorchiefs/dlwbl_eth25/blob/master/notebooks/02_cnn_edge_lover.ipynb)

# More then one kernel (Motivatoion)

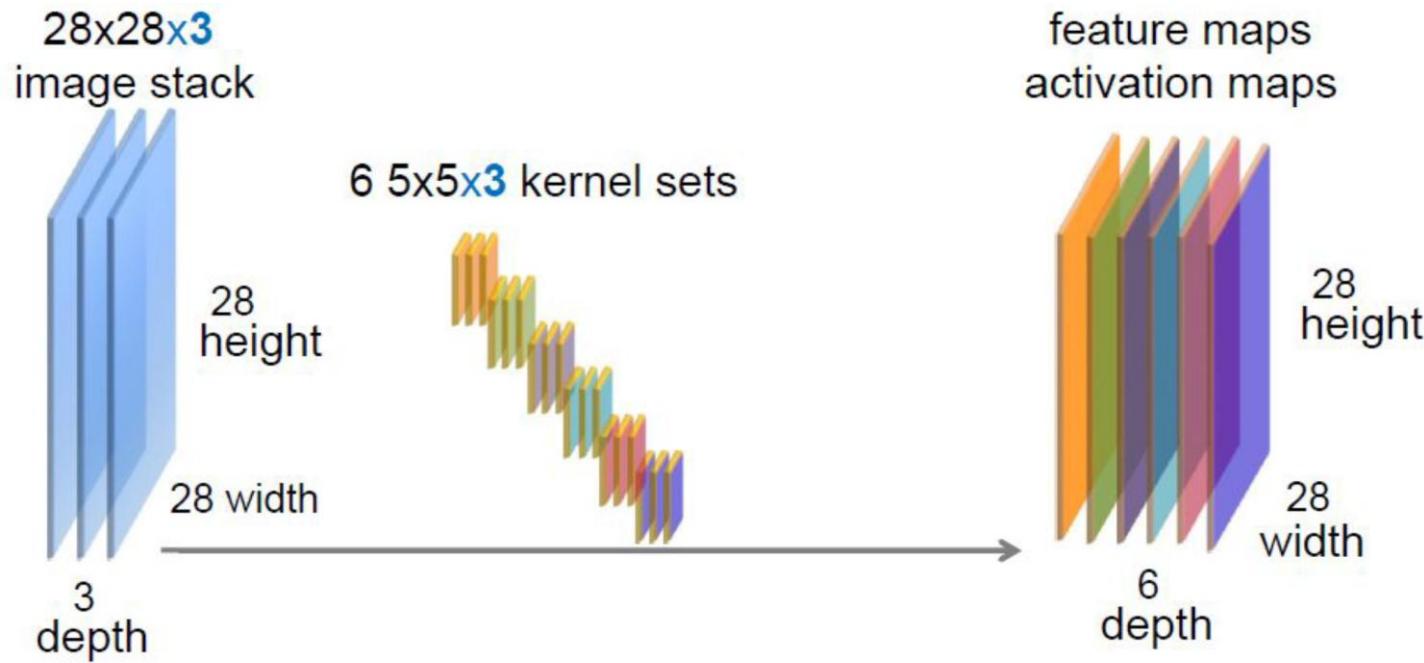


## Convolution layer with a 1-channel input and 6 kernels



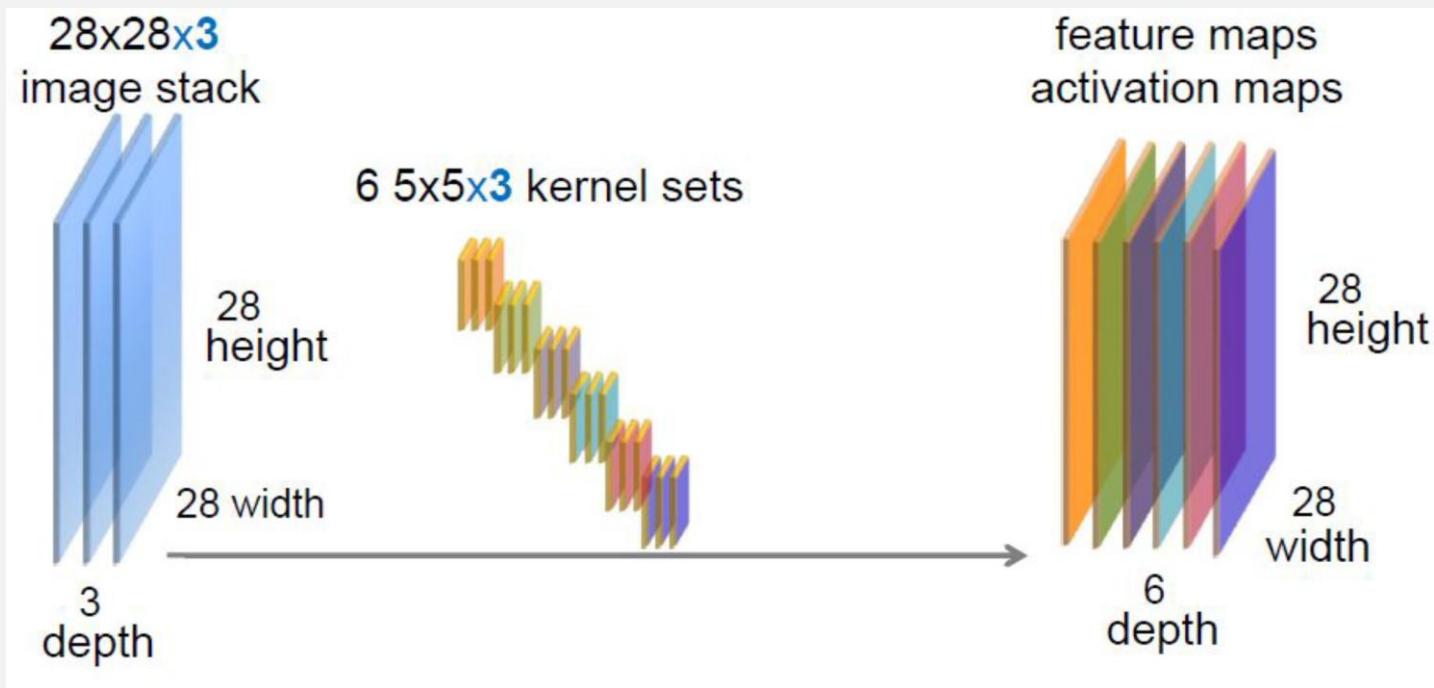
Convolution of the input image with 6 different kernels results in 6 activation maps.  
If the input image has only one channel, then each kernel has also only one channel.

## Convolution layer with a 3-channel input and 6 kernels



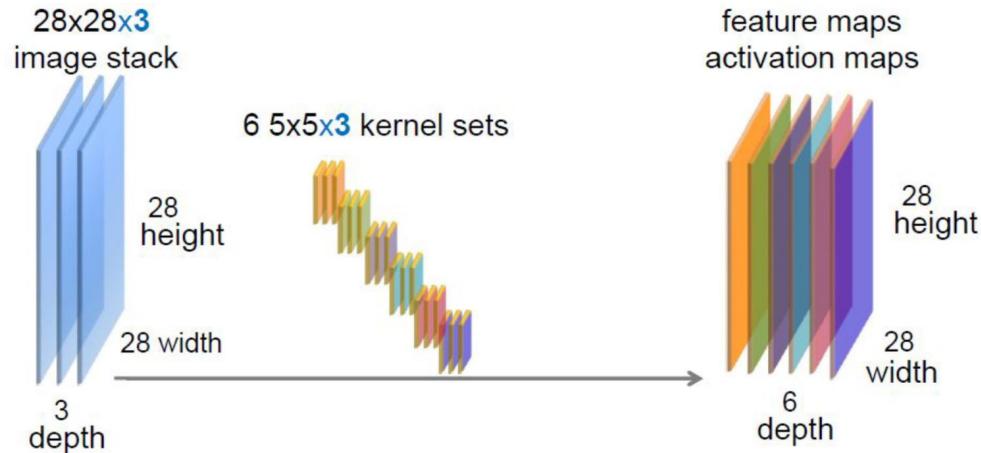
Convolution of the input image with 6 different kernels results in 6 activation maps.  
If the input image has 3 channels, then each filter has also 3 channels.

# Convolution layer with a 3-channel input and 6 kernels



How many weights?

# Solution



$$6 \times 5 \times 5 \times 3 + 6 = 456$$

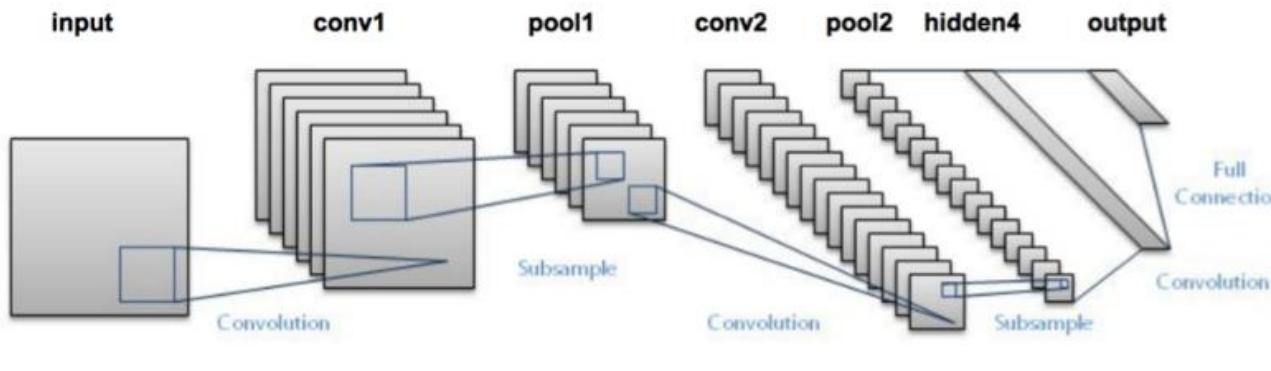
```
# Number of weights in Convolution
model = Sequential()
model.add(Convolution2D(6,kernel_size=(5,5),padding='same',input_shape=(28,28,3)))
model.summary()
```

Model: "sequential\_11"

Layer (type)	Output Shape	Param #
conv2d_44 (Conv2D)	(None, 28, 28, 6)	456

Total params: 456 (1.78 KB)  
Trainable params: 456 (1.78 KB)  
Non-trainable params: 0 (0.00 B)

# CNN for MNIST



```
model = Sequential()

model.add(Convolution2D(filters=8, kernel_size=(3,3),
                       padding='same', input_shape=(28,28,1)))
model.add(Activation('relu'))
model.add(Convolution2D(16, (3, 3), padding='same'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Flatten())
model.add(Dense(40))
model.add(Activation('relu'))
model.add(Dense(10))
model.add(Activation('softmax'))

# compile model and intitialize weights
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

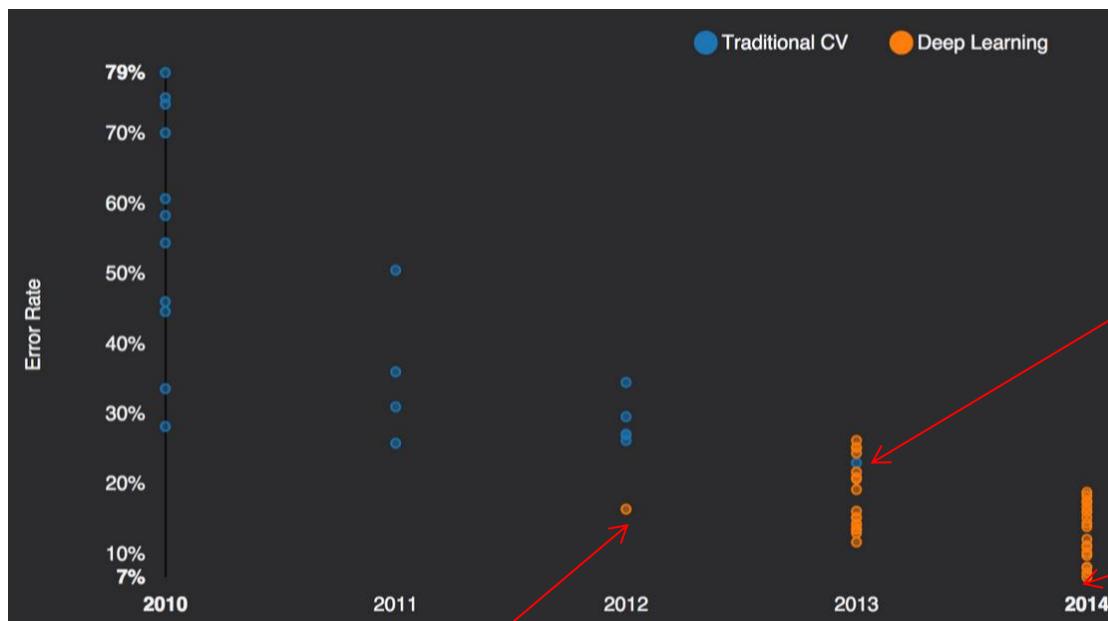
What is a good  
CNN architecture?

# CNN breakthrough in 2012: Imagenet challenge

1000 classes  
1 Mio samples



...



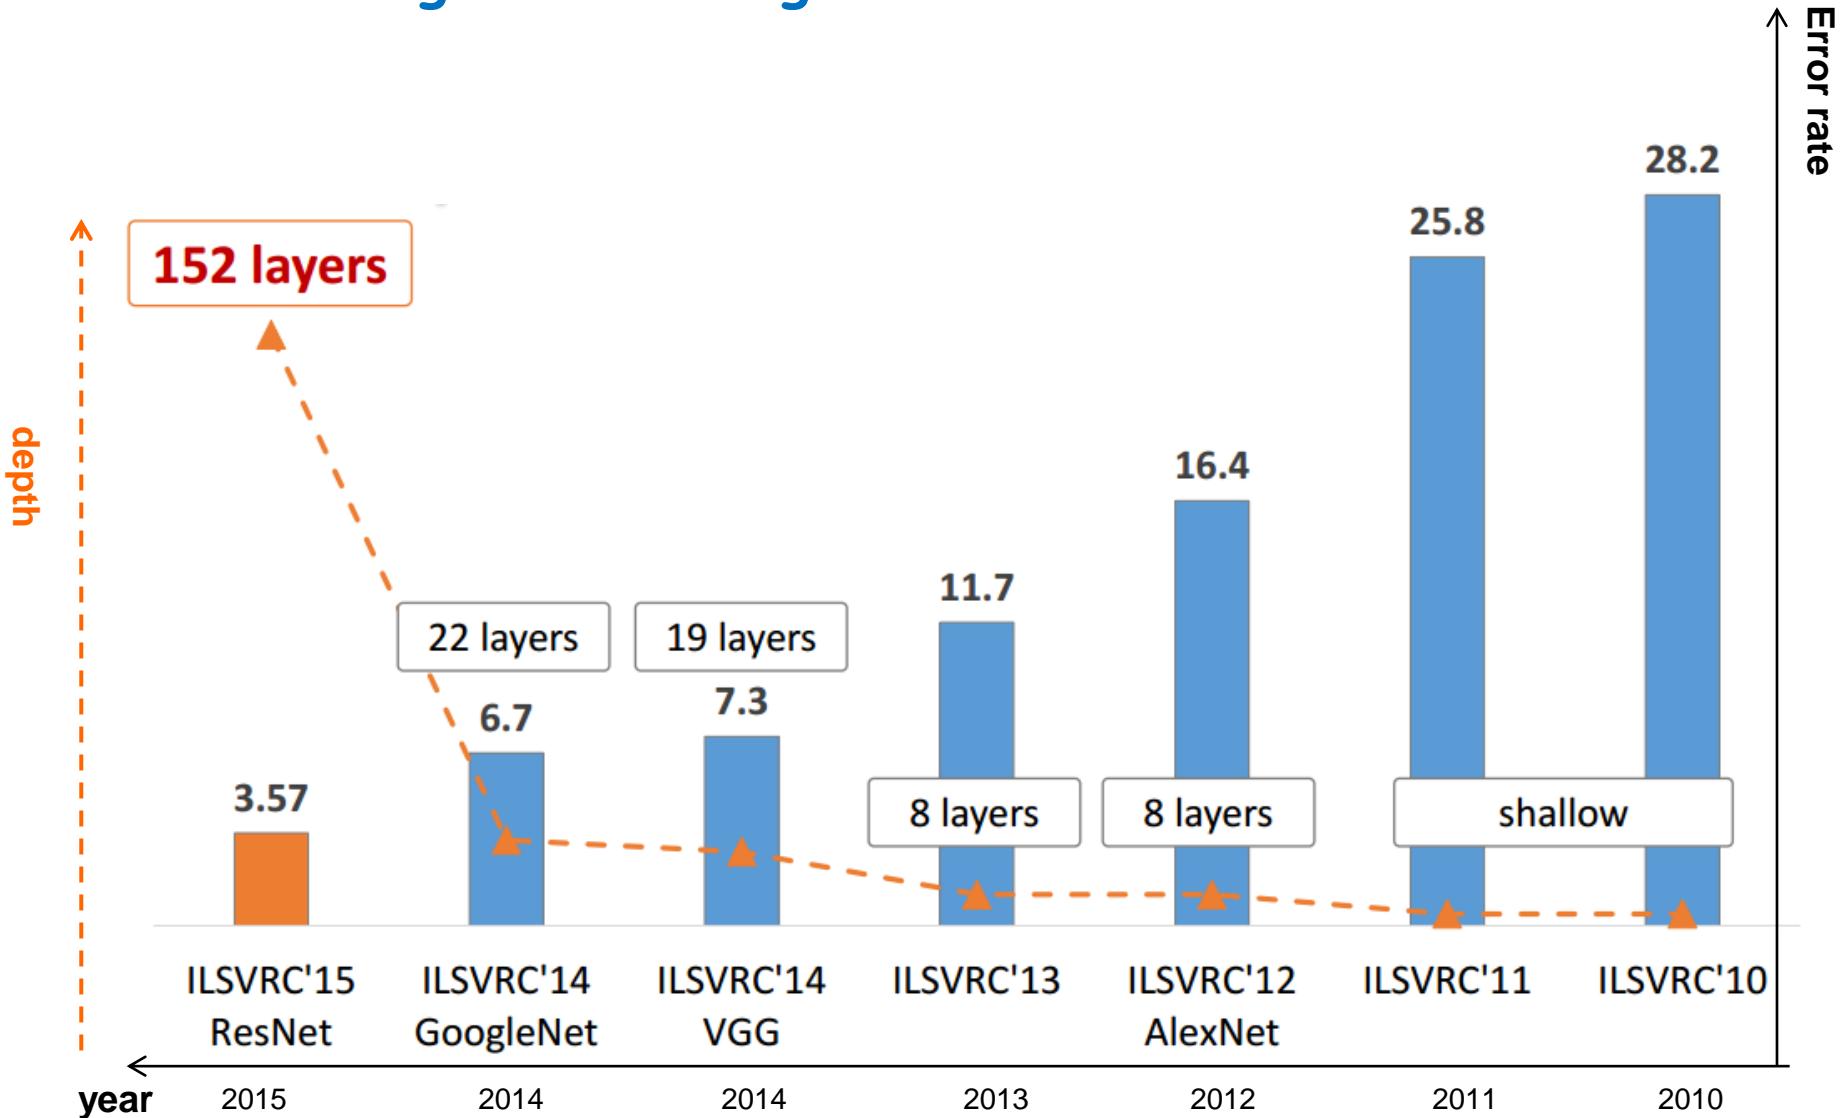
Human: 5% misclassification

Only one non-CNN approach in 2013

GoogLeNet 6.7%

A. Krizhevsky  
first CNN in 2012  
**Und es hat zoom gemacht**

# Review of ImageNet winning CNN architectures



# LeNet-5 1998: first CNN for ZIP code recognition

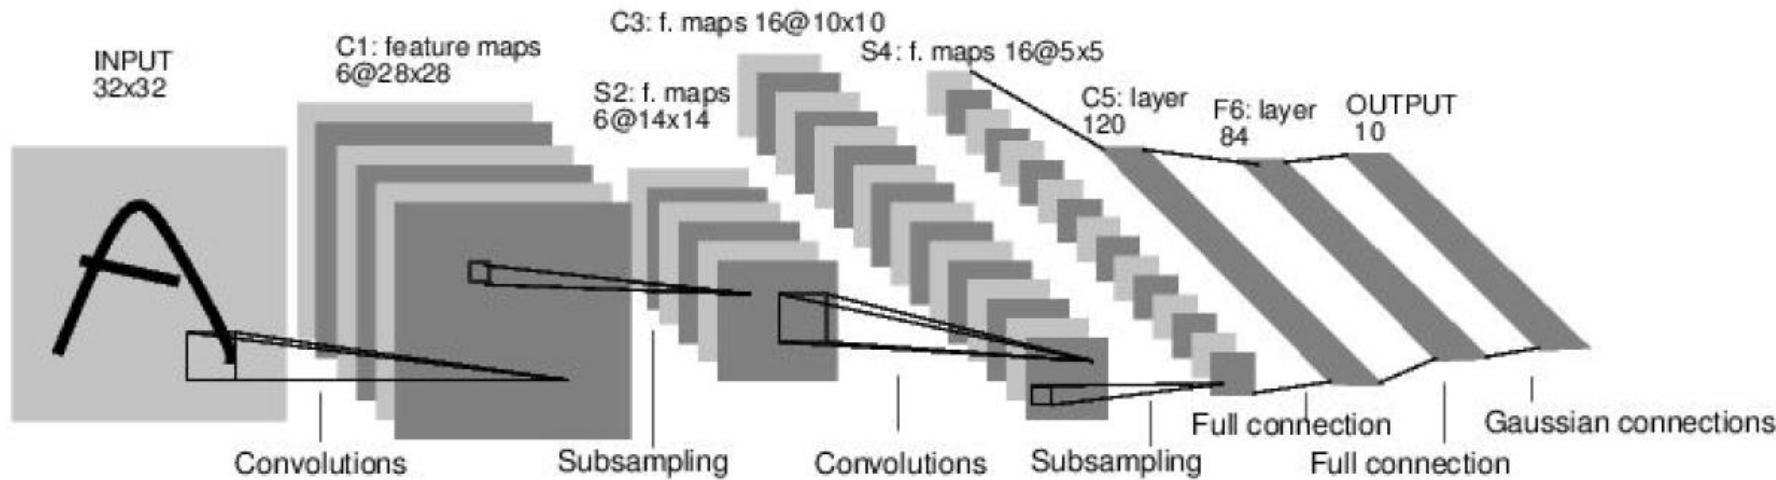


Image credits: <http://yann.lecun.com/exdb/publis/pdf/lecun-98.pdf>

Conv filters were **5x5**, applied at stride 1

Subsampling (Pooling) layers were **2x2** applied at stride 2  
i.e. architecture is [CONV-POOL-CONV-POOL-CONV-FC]

[Demo von 1993](#) Yann LeCun



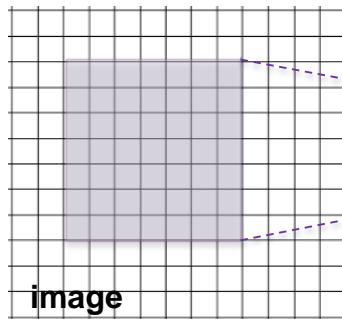
<http://yann.lecun.com/exdb/lenet/index.html>

# The trend in modern CNN architectures goes to small filters

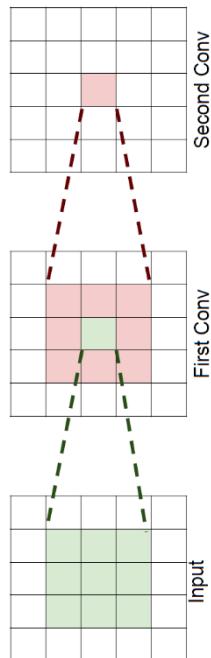
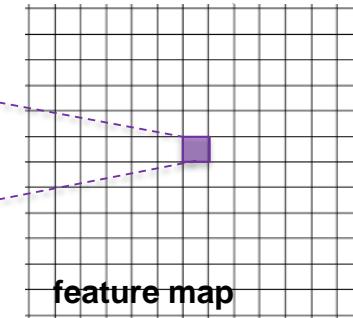
Why do modern architectures use very small filters?

Determine the receptive field in the following situation:

- 1) Suppose we have one  
7x7 conv layers (stride 1)  
49 weights

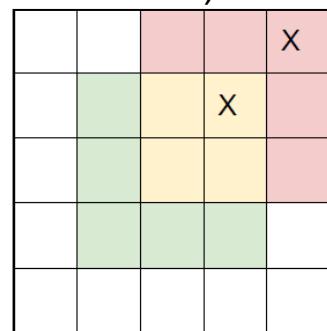


Answer1): 7x7



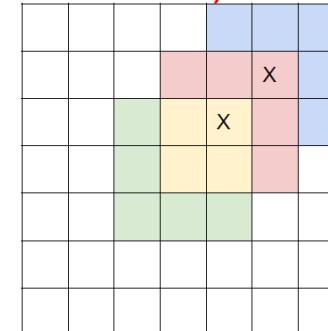
- 2) Suppose we stack two  
3x3 conv layers (stride 1)

Answer 2): 5x5



- 3) Suppose we stack three  
3x3 conv layers (stride 1)  
 $3 \times 9 = 27$  weights

Answer 3): 7x7



We need less weights for the same receptive field when stacking small filters!

# “Oxford Net” or “VGG Net” 2<sup>nd</sup> place

- 2<sup>nd</sup> place in the imageNet challenge
- More traditional, easier to train
- Small pooling
- Stacked 3x3 convolutions before maxpooling
  - > large receptive field
- no strides (stride 1)
- ReLU after conv. and FC (batchnorm was not introduced)
- Pre-trained model is available (see excercise)



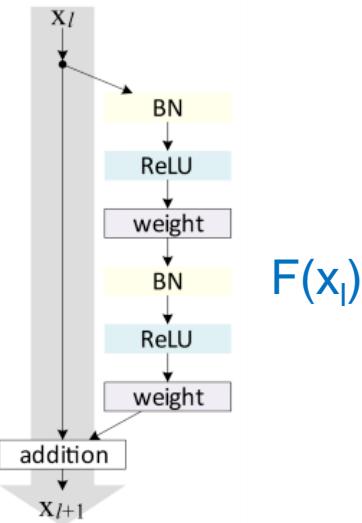
<http://arxiv.org/abs/1409.1556>

# "ResNet" from Microsoft 2015 winner of imageNet

152  
layers

ResNet basic design (VGG-style)

- add shortcut connections every two
- all 3x3 conv (almost)



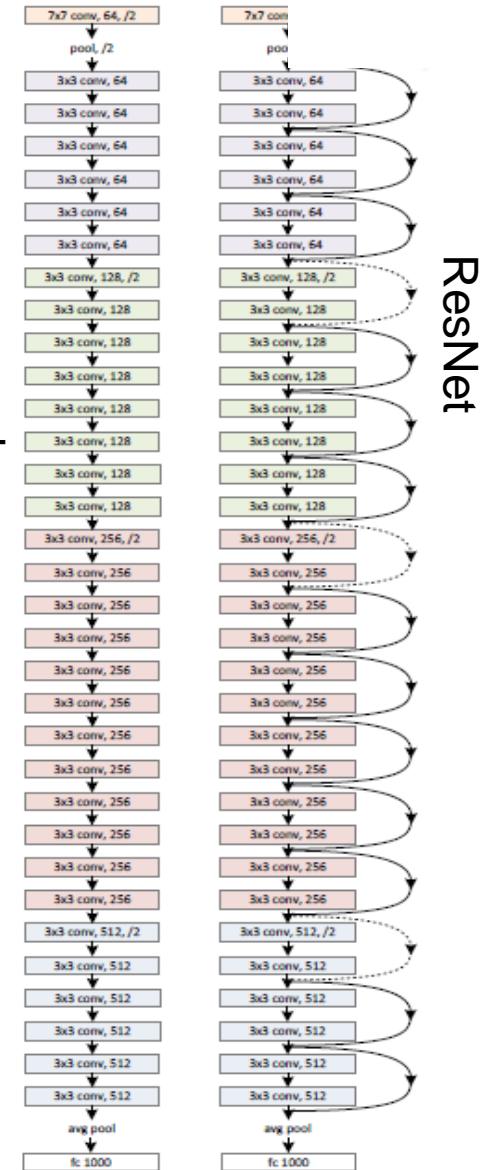
$$H(x_i) = x_{i+1} = x_i + F(x_i)$$

$F(x)$  is called "residual" since it only learns the "delta" which is needed to add to  $x$  to get  $H(x)$

152 layers:  
Why does this train at all?

This deep architecture  
could still be trained, since  
the gradients can skip  
layers which diminish the  
gradient!

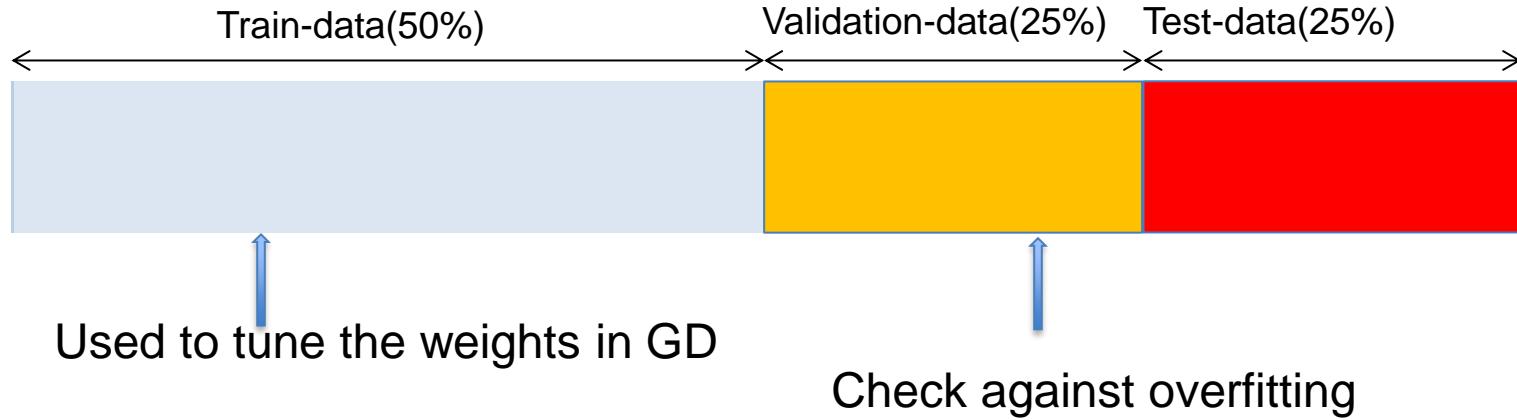
plain VGG



## Tricks of the trade

- Early stopping
- Input Standardization
- Batch Norm Layer
- Dropout
- Data augmentation

# Best practice: Split in Train, Validation, and Test Set

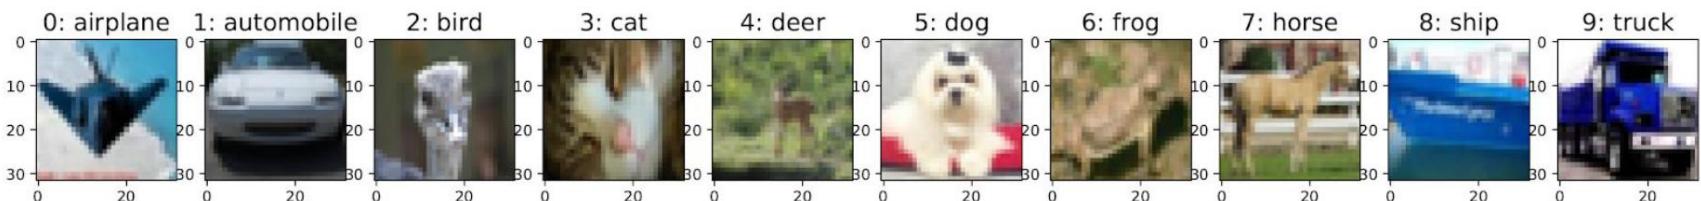


Best practice: Lock an extra **test data set** away, and use it only at the very end, to evaluate the chosen model, that performed best on your validation set.

Reason: **When trying many models, you probably overfit on the validation set.**

Determine performance metrics, such as MSE, to evaluate the predictions **on new validation or test data**

# CIFAR10 study with tensorflow notebook



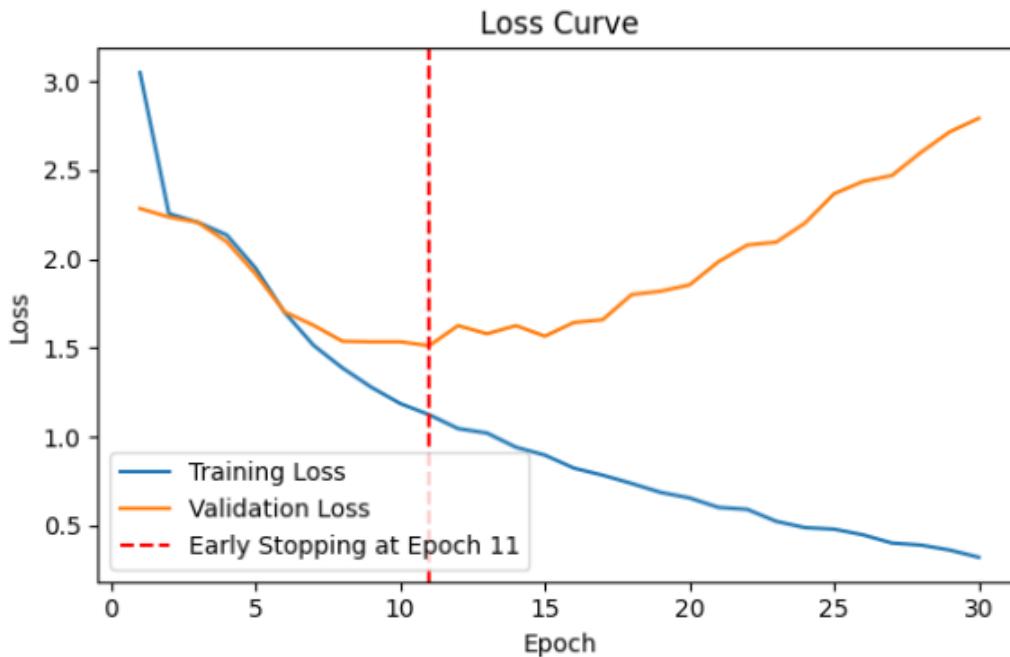
We define a CNN to classify cifar10 images (we have 10 classes)

Notebook:

[https://github.com/tensorchiefs/dl\\_course\\_2020/blob/master/notebooks/07\\_cifar10\\_norm\\_sol.ipynb](https://github.com/tensorchiefs/dl_course_2020/blob/master/notebooks/07_cifar10_norm_sol.ipynb)

# Loss curve and early stopping

Very common check: Plot loss in train and validation data vs epoch of training.

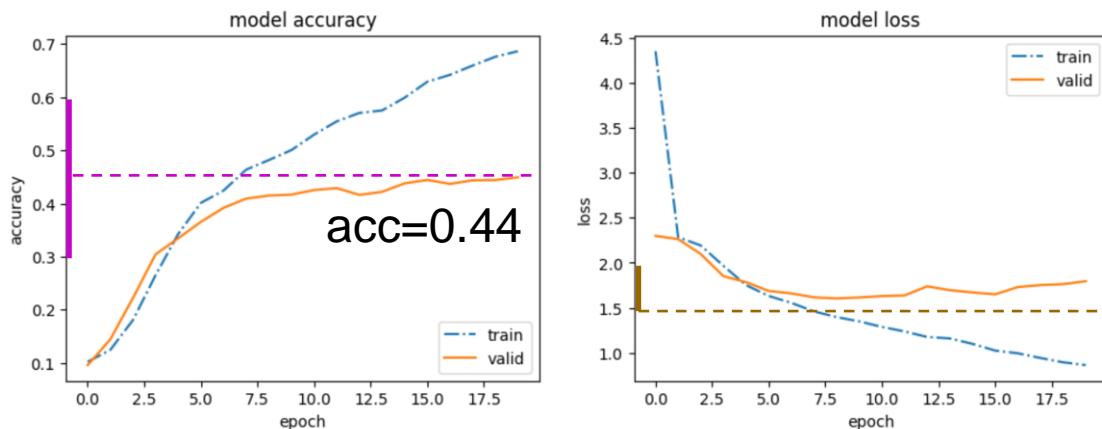


- If **training loss** does not go down to zero: model is not flexible enough
- Use weights @minimum of validation before overfitting
- Early stopping: stop to training if validation loss does not improve anymore

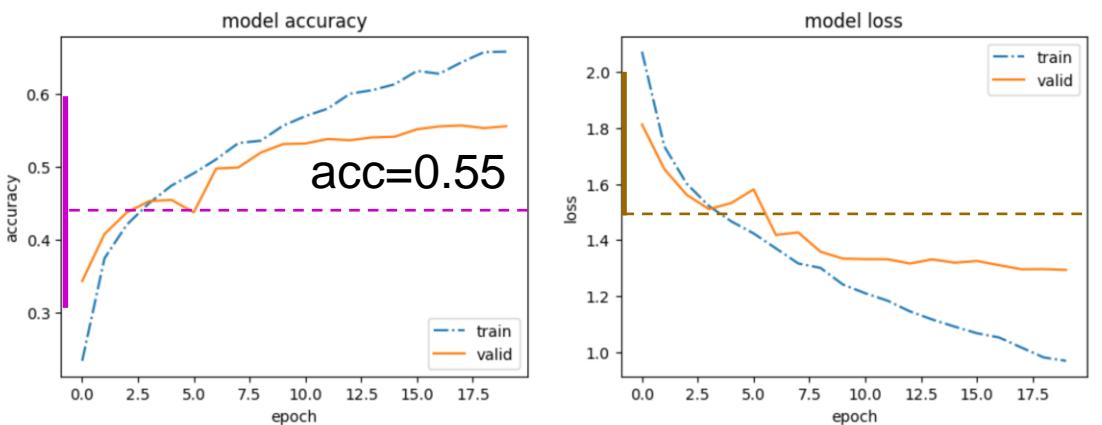
# Take-home messages from CIFAR10 CNN study

- DL does not need a lot of preprocessing, but working with standardized (small-valued) input data often helps.

Without normalizing  
the input to the CNN



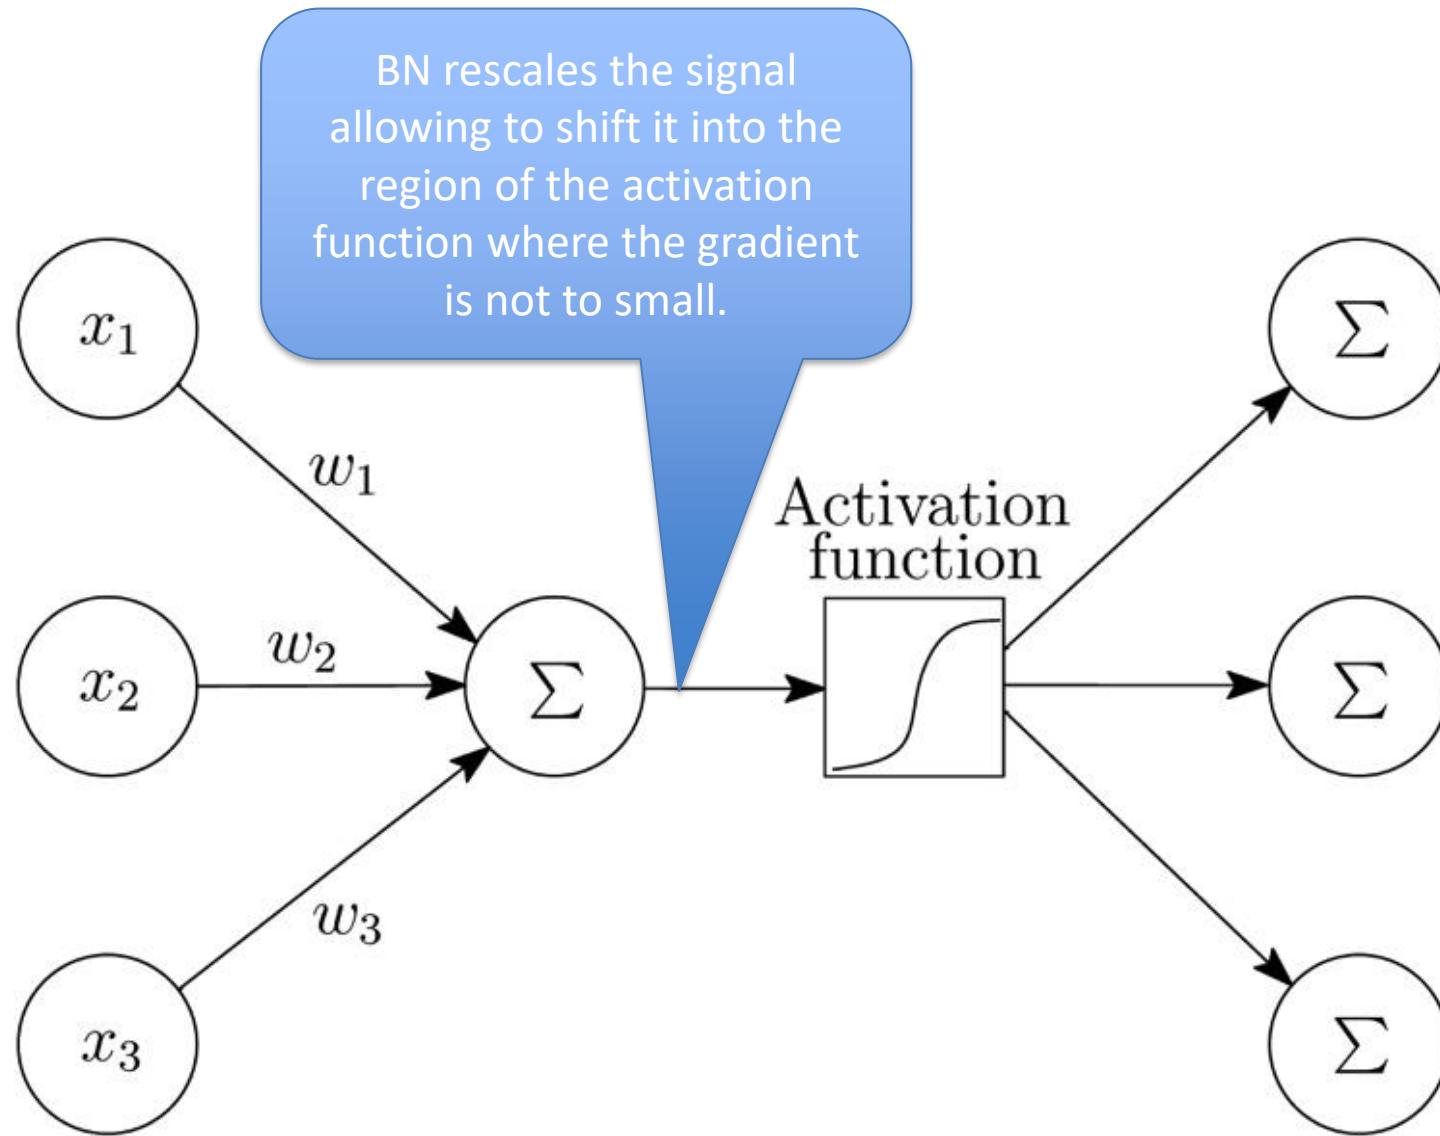
With normalizing  
(pixel-value/255)  
the input to the CNN



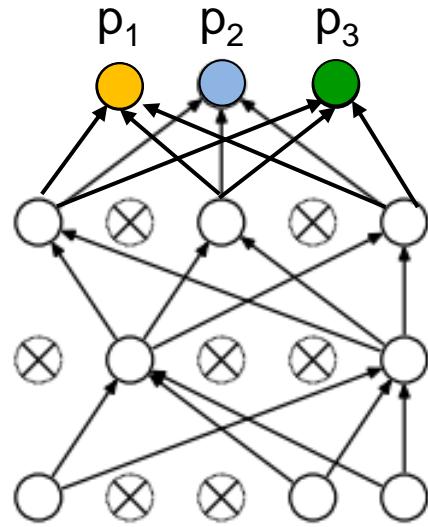
## Regularization to avoid overfitting

- Batchnorm layer
- Dropout layer

# What is the idea of Batch-Normalization (BN)



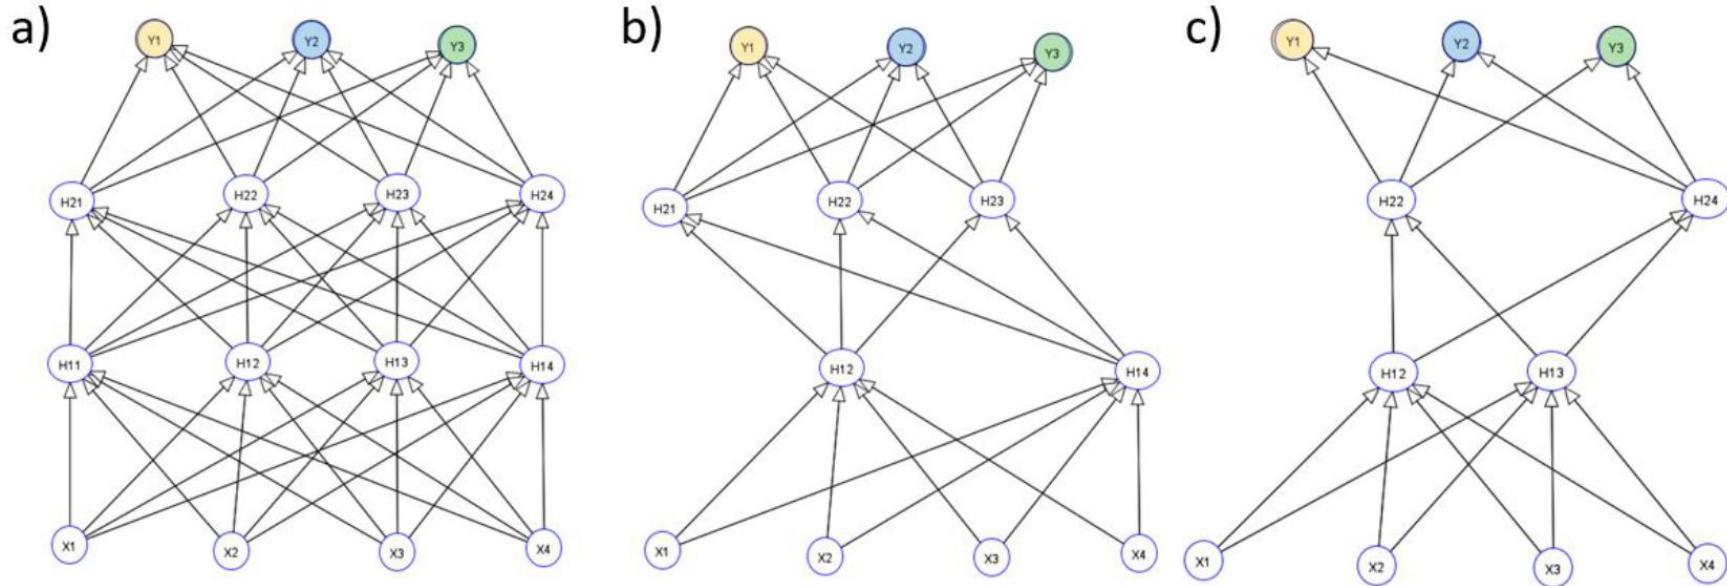
# Dropout helps to fight overfitting



Using **dropout during training implies:**

- In each training step only weights to not-dropped units are updated → we train a sparse sub-model NN
- For predictions with the trained NN we freeze the weights corresponding to averaging over the ensemble of trained models we should be able to “reduce noise”, “overfitting”
- JFI: To get same expected output in training (with dropout) and after training (test time - without dropout), the weights are multiplied after training by the dropout probability  $p=0.5$ .

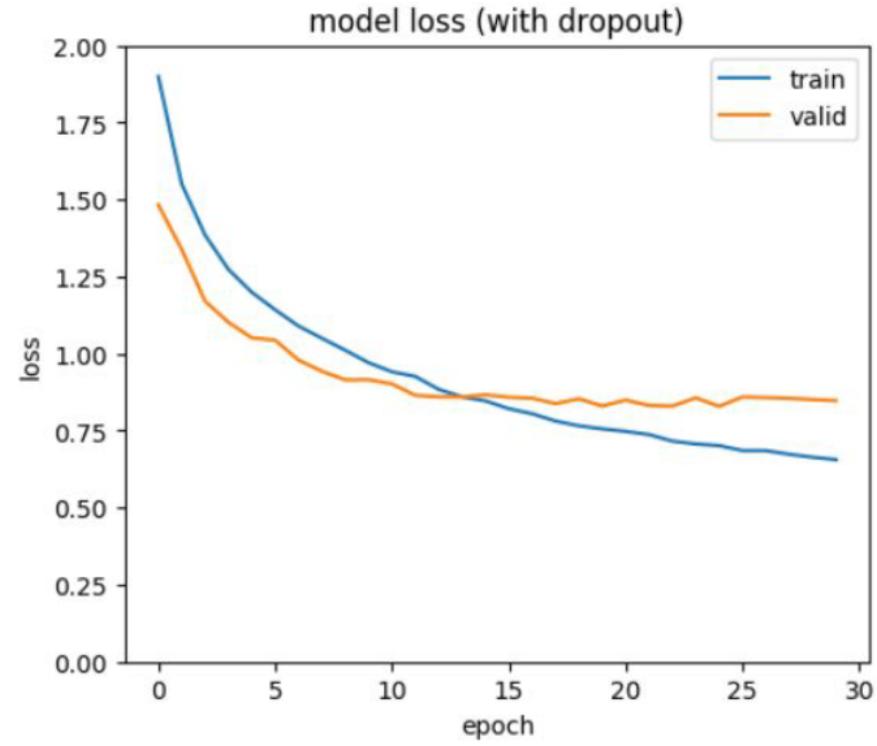
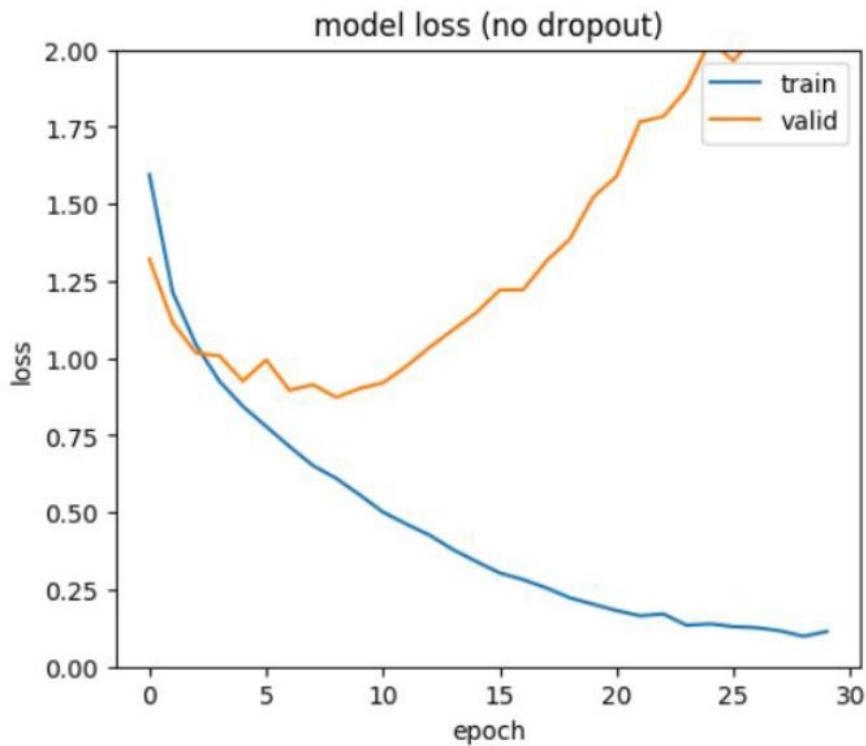
# Dropout



Three NNs:

a) shows the full NN with all neurons (as used when NN is trained),  
b) and c) show two versions of a thinned NN where some neurons are dropped (as done during training with dropout). Dropping neurons is the same as setting all connections that start from these neurons to zero.

# Dropout fights overfitting in a CIFAR10 CNN



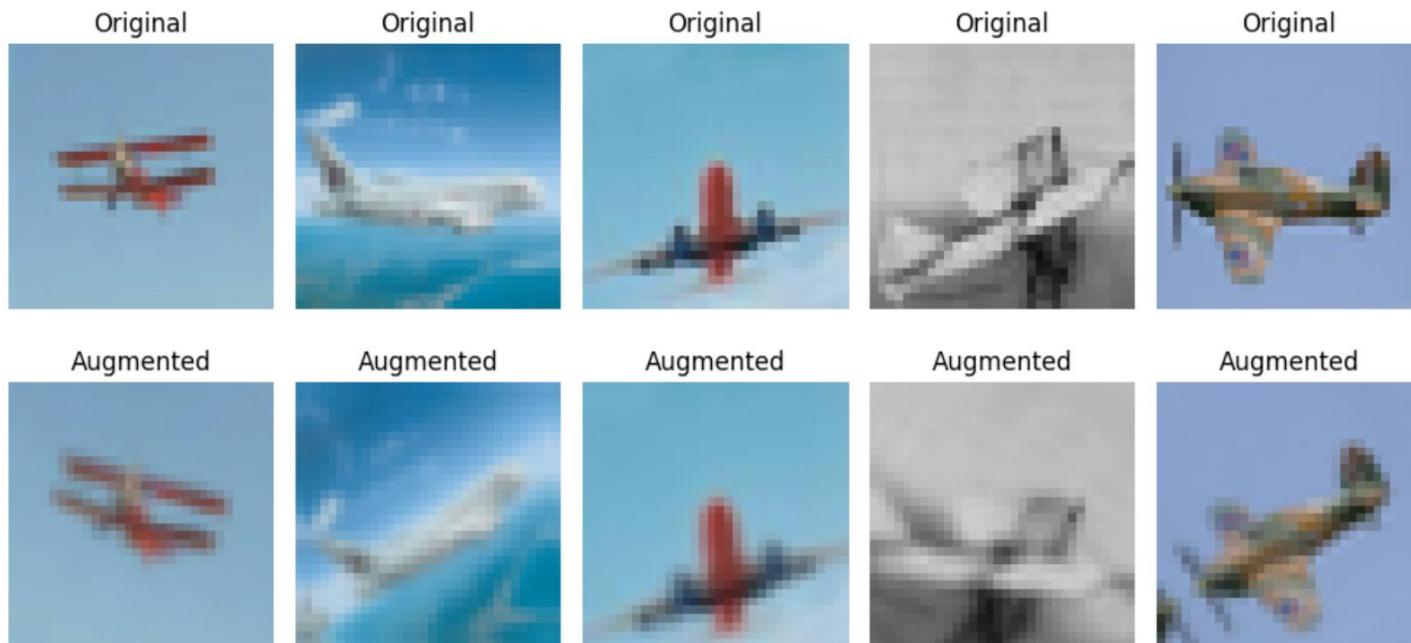
## What to do in case of limited data?

- Data augmentation
- Pretrained (foundation) models

# Fighting overfitting by Data augmentation:

During training random operations are done on the fly

- Rotate image within an angle range
- Flip image: left/right, up, down
- resize
- Take patches from images
- ....



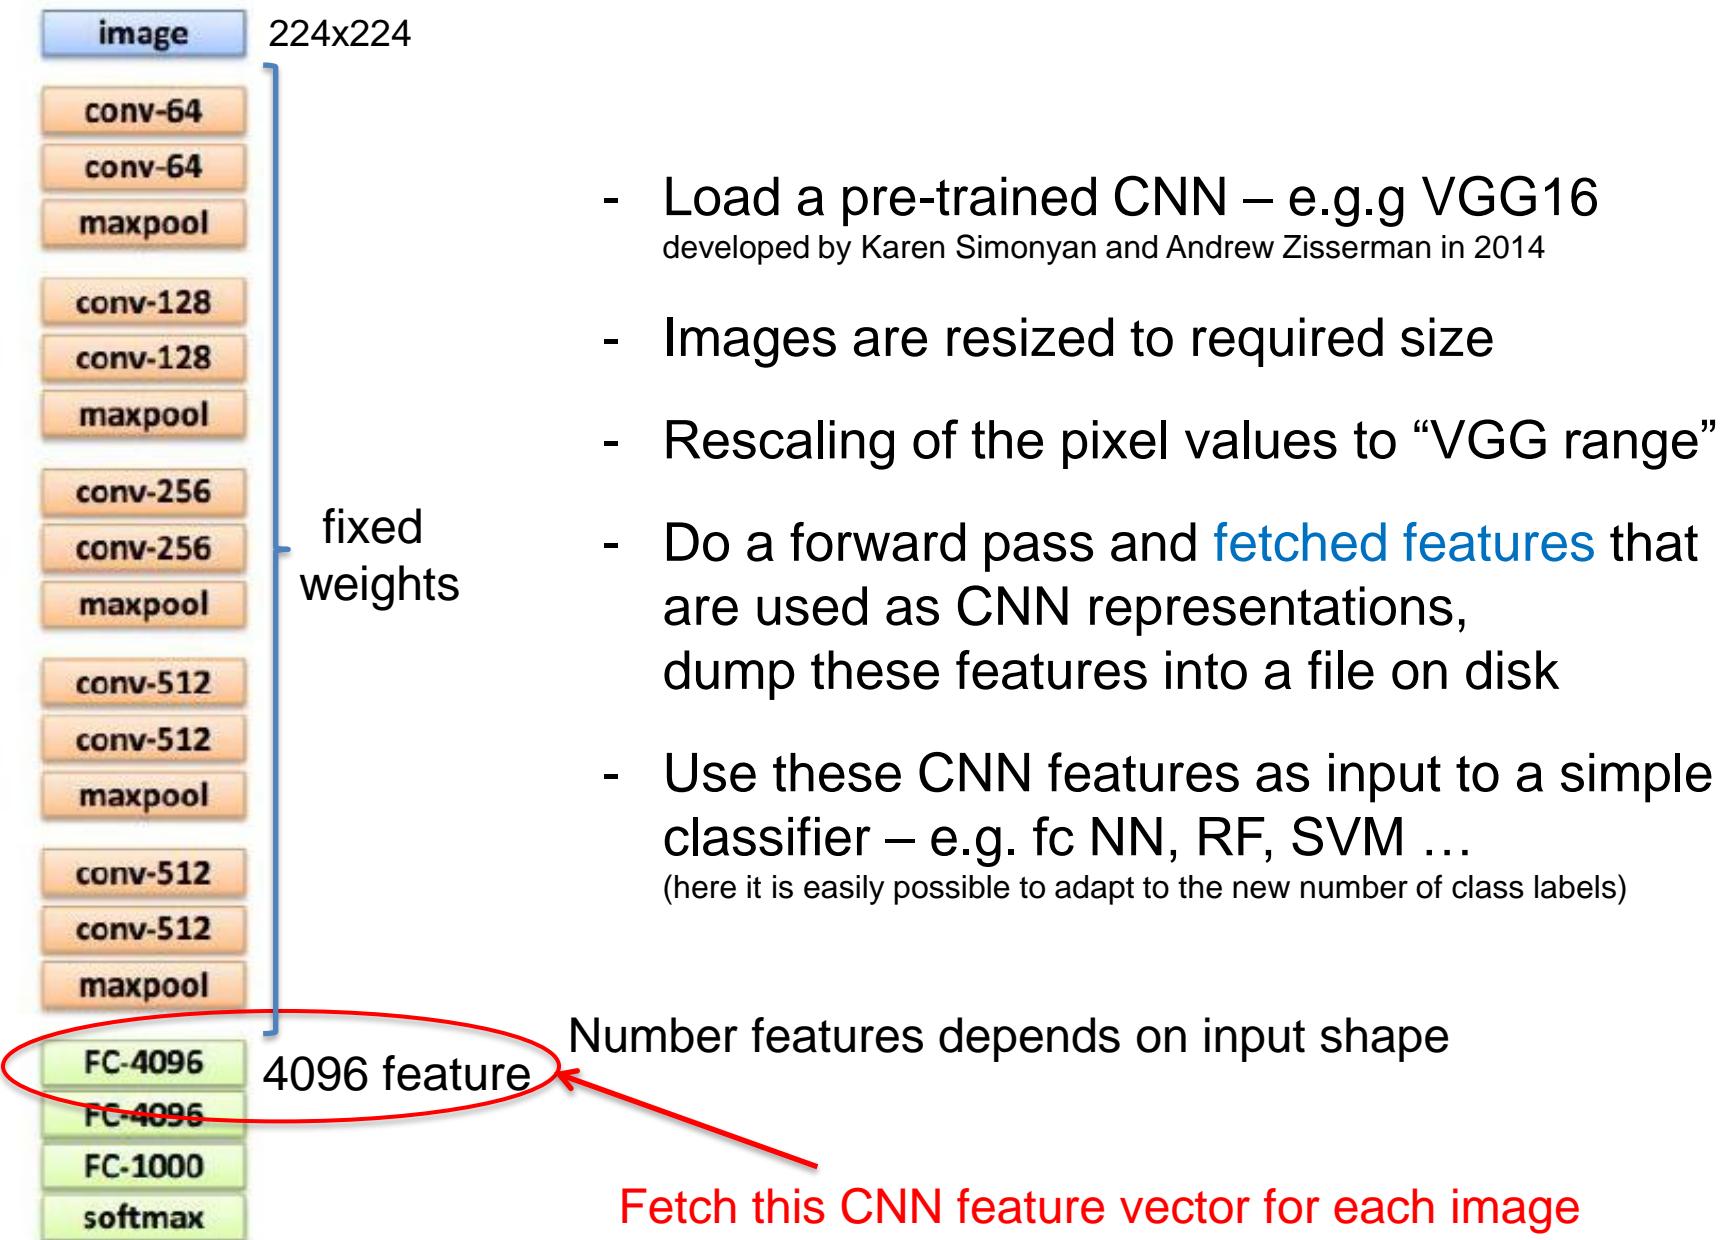
# In Code

```
1 # Define data augmentation pipeline
2 from keras.layers import RandomFlip, RandomRotation, RandomContrast, RandomZoom
3 data_augmentation = keras.Sequential([
4     RandomFlip("horizontal"),           # Randomly flip images horizontally
5     RandomRotation(0.1),              # Randomly rotate images by 10%
6     RandomZoom(0.1),                 # Randomly zoom in on images
7     RandomContrast(0.1),             # Adjust contrast randomly
8 ], name='Augmentation')
a
```

```
1 # Functional API Model Definition
2 x = Input(shape=input_shape)
3 x = data_augmentation(x)
4 # First Convolutional Block
5 x = Convolution2D(8, kernel_size, padding='same')(x)
6 ...
```

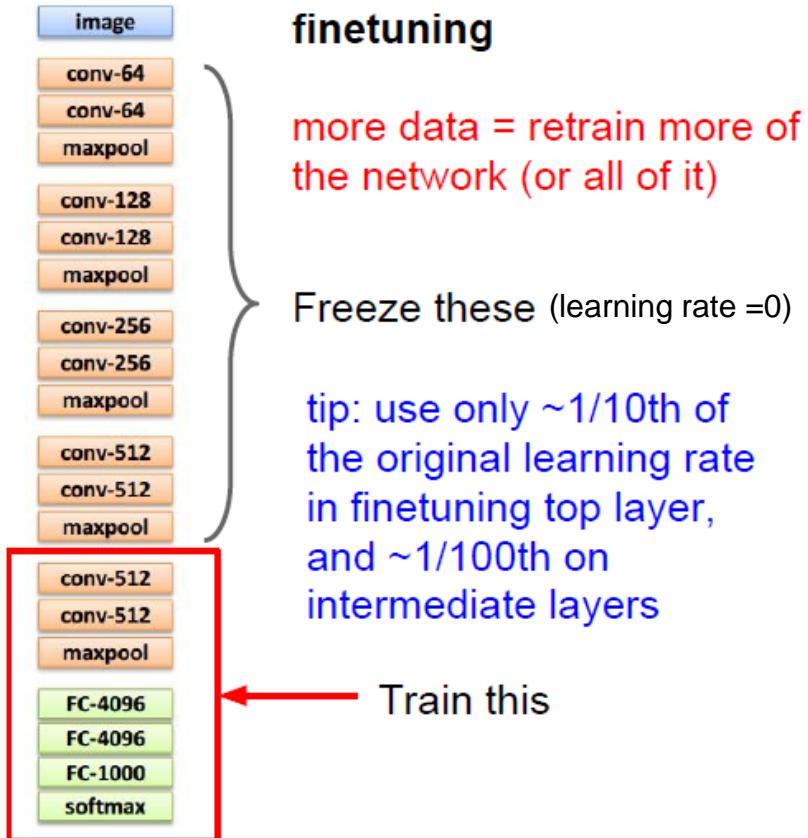
See: [https://github.com/tensorchiefs/dlwbl\\_eth25/blob/master/notebooks/02\\_transfer\\_learning.ipynb](https://github.com/tensorchiefs/dlwbl_eth25/blob/master/notebooks/02_transfer_learning.ipynb)

# Use pre-trained CNNs for feature generation



# Transfer learning beyond using off-shelf CNN feature

e.g. medium data set (<1M images)



The strategy for fine-tuning depends on the size of the data set and the type of images:

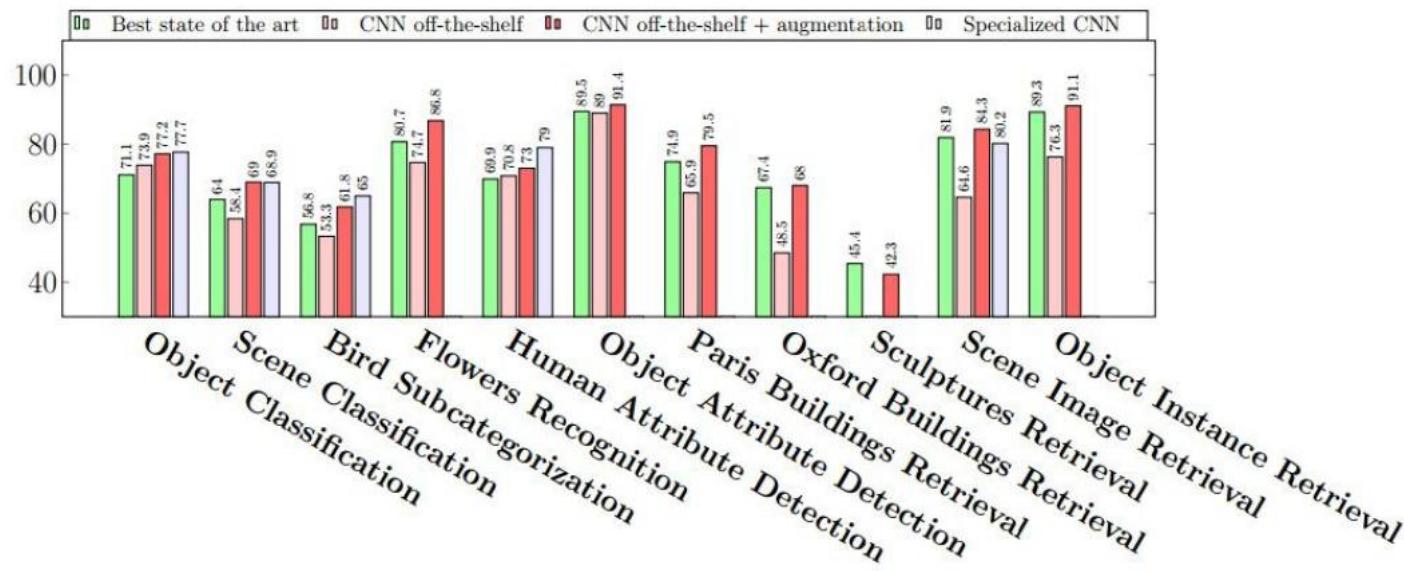
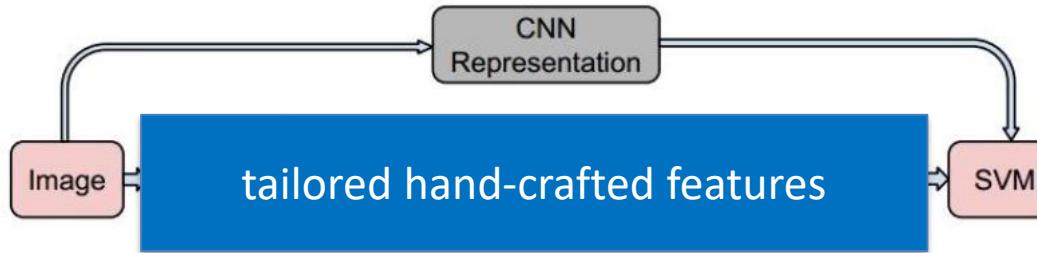
	<b>Similar task</b> (to imageNet challenge)	<b>Very different task</b> (to imageNet challenge)
<b>little data</b>	Extract CNN representation of one top fc layer and use these features to train an external classifier	You are in trouble - try to extract CNN representations from different stages and use them as input to new classifier
<b>lots of data</b>	Fine-tune a few layers including few convolutional layers	Fine-tune a large number of layers

Where to get pretrained CNNs like VGG16 or ResNet: <https://keras.io/api/applications/>

Hint: first retrain only fully connected layer, only then add convolutional layers for fine-tuning.

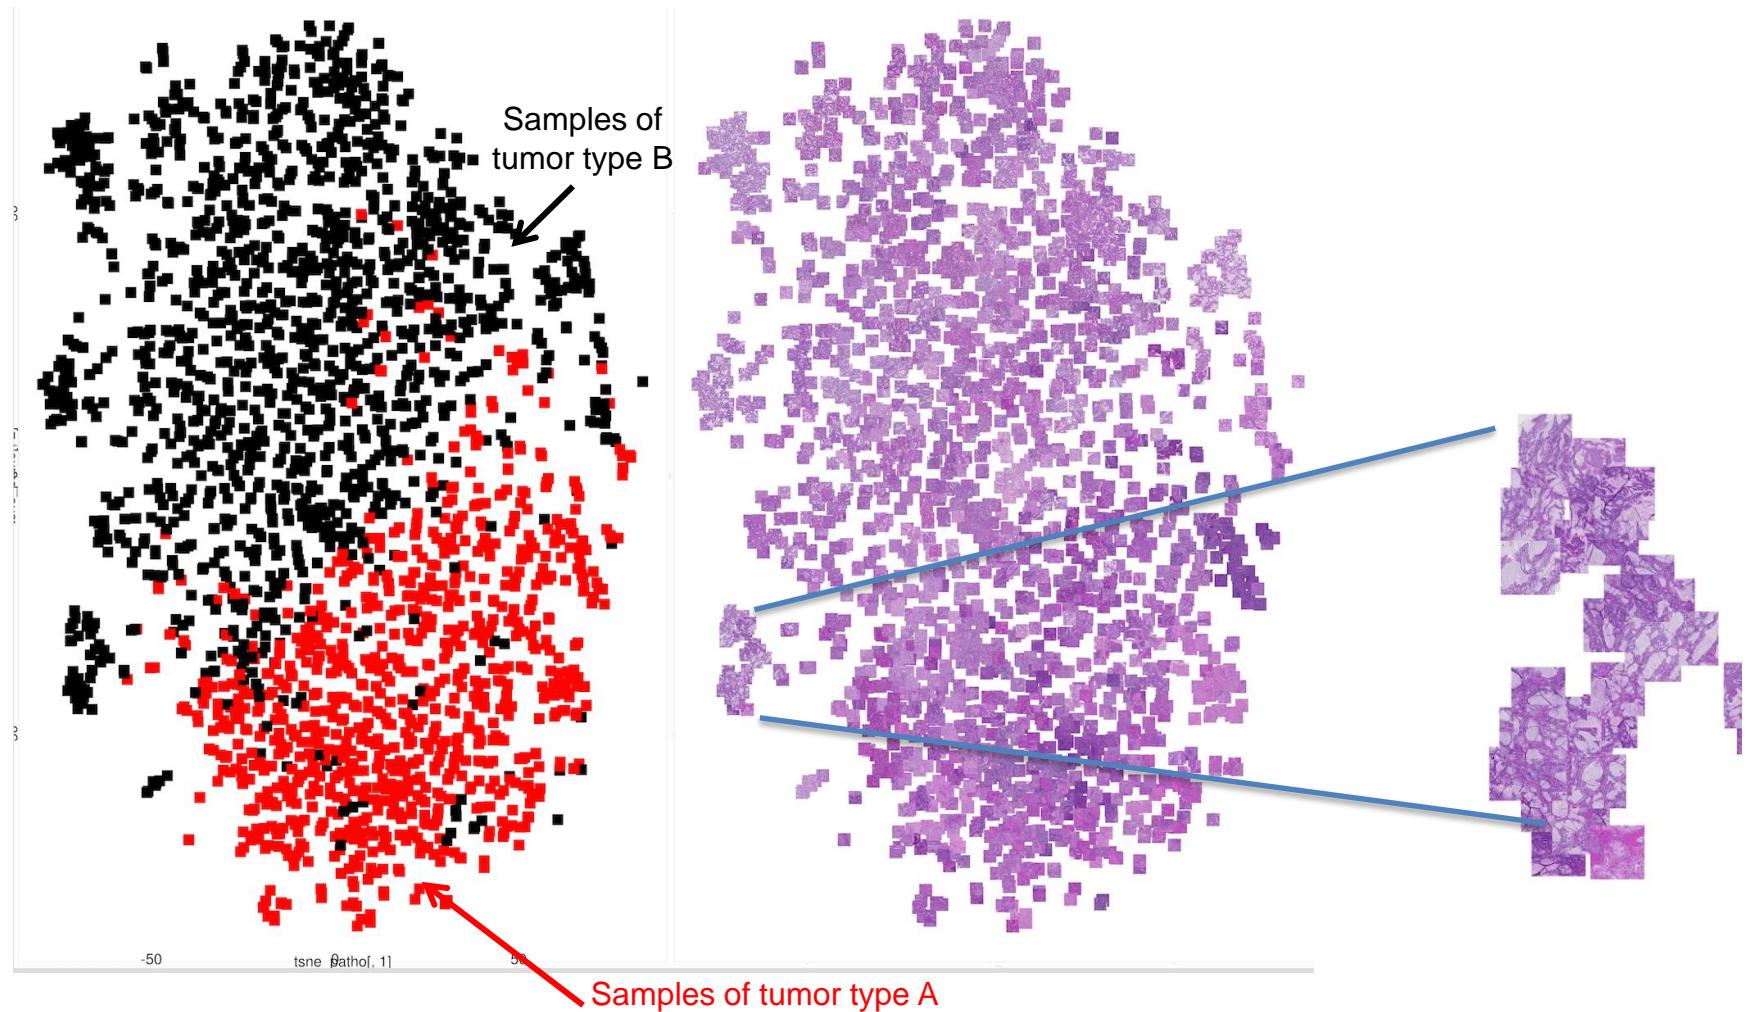
# Performance of off-the-shelf CNN features when compared to tailored hand-crafted features

CNN's distributed and compositional features generalize well to new tasks



“Astonishingly, we report consistent superior results compared to the highly tuned state-of-the-art systems in all the visual classification tasks on various datasets.”

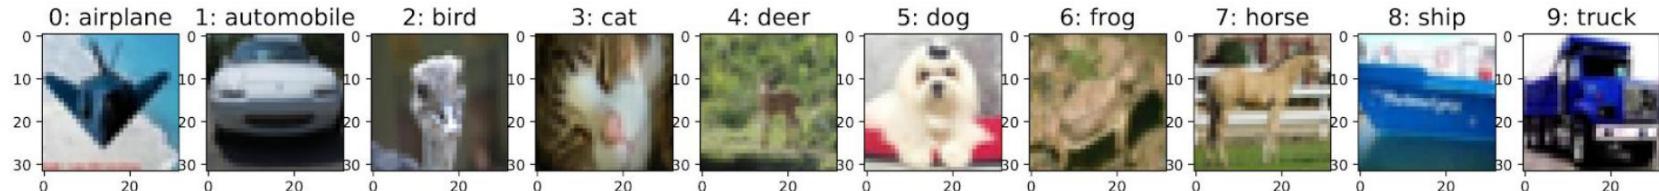
# Use features from a pretrained VGG as input to t-SNE



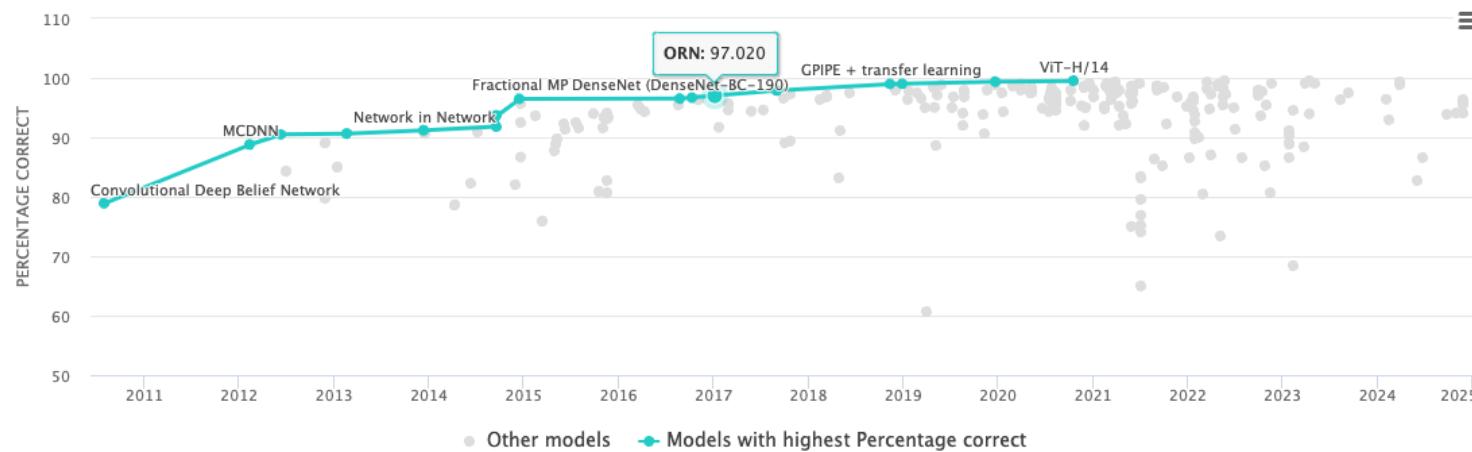
- Different tissue types cluster together in t-SNE: we could use knn as classifier
- VGG features even work on images that are far away from the 1000 imageNet classes

# Case Study: CIFAR10 Dataset

# Cifar10 Data



10 Classes, with 5'000 images for training



Some examples:

VGG-19 with GradInit	2021	94.71
ResNet-18	2022	95.55

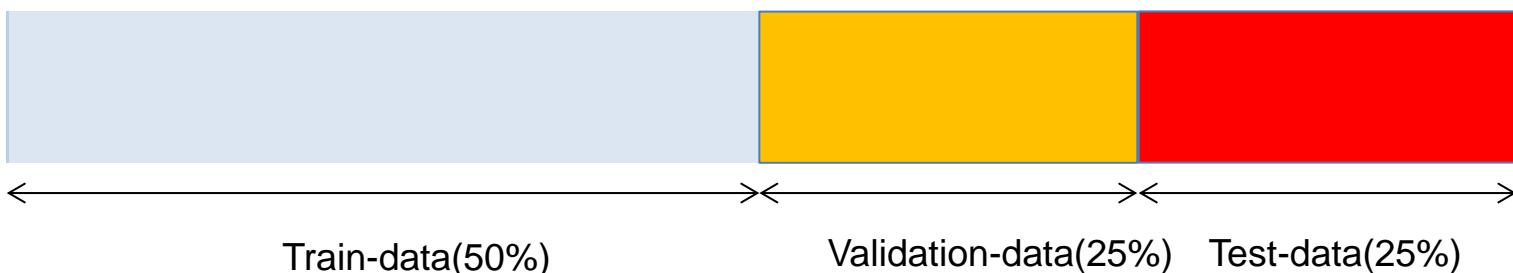
# Dataset Splitting in Machine Learning Competitions

- Training Set
    - The set you get for training, you can do what you want with it.
  - Test Set:
    - Final evaluation set, never seen by the model during training (e.g., 10,000 images in CIFAR-10).
- Note: CIFAR-10 and many other datasets do NOT provide a separate validation set.
- Typical Workflow
    - Hyperparameter Tuning
      - Create validation set from training set
        - Helps in tuning hyperparameters and preventing overfitting.
        - Common practice: split 80% training / 20% validation (e.g., 40,000 train / 10,000 validation).
      - Final Training: After hyperparameter tuning, models are often retrained on the entire training set
      - Evaluation on test set
        - Sometimes test set is secret and you have to upload your predictions (private test set)

```
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

x, y\_train

x, y\_test



# Exercise: Transfer learning with CIFAR10 data

